

AI and Deep Learning

Deep Learning


Jeju National University

Yungcheol Byun

Agenda

- Merging gates in a computation graph
- Vanishing gradient and ReLU
- MNIST application
- Overfitting and drop-out
- Deep Learning


Influence of w change on E




is multiplication of all the local gradients in the graph
(chain rule)

Influence of w change on E


Merging gates



Influence of w change on E




Influence of w change on E



is multiplication of **all the local gradients** in the graph
(chain rule)


Influence of w change on E



is multiplication of all the local gradients in the graph
(chain rule)

Influence of w change on E

Merging gates



Therefore, the local gradient is derivative of the E .

Influence of w change on E


$$E = (w \cdot 1 - 1)^2$$

Derivative of E with respect to w

$$\frac{\partial E}{\partial w} = \frac{\partial}{\partial w} (w \cdot 1 - 1)^2 = 2(w - 1)$$

Cost/Error function

for logistic regression



$$hypo = \frac{1}{1 + e^{-wx}}$$



$$E = -y \log(hypo) - (1 - y) \log(1 - hypo)$$

Binary Cross Entropy

Computational Graph

Binary Cross Entropy 오류 함수(loss function)

$$E = -y \log(\text{hypo}) - (1 - y) \log(1 - \text{hypo})$$

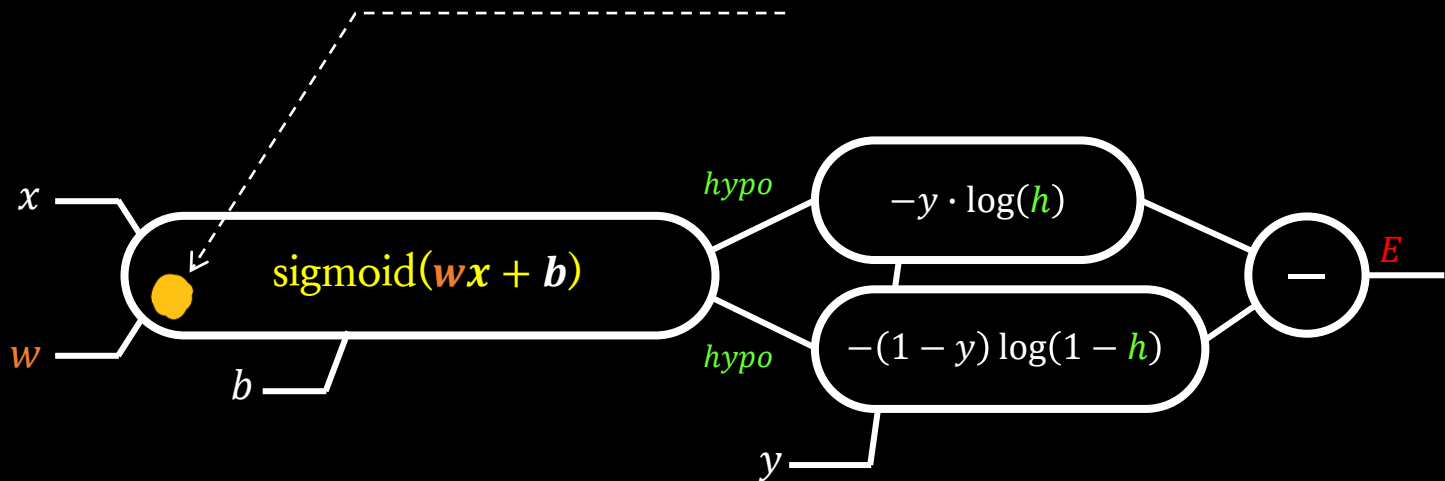



$$\frac{\partial E}{\partial w} =$$

Computational Graph

Merging gates


How we get **the local gradient** of the merged gate(sigmoid)?






$$\frac{\partial \text{sigmoid}}{\partial w} = \text{sigmoid}(1 - \text{sigmoid})$$

1 Neuron(2-layer)




3-layer NN




Why? what for?


3-layer NN (simplified)



Influence of w change on E



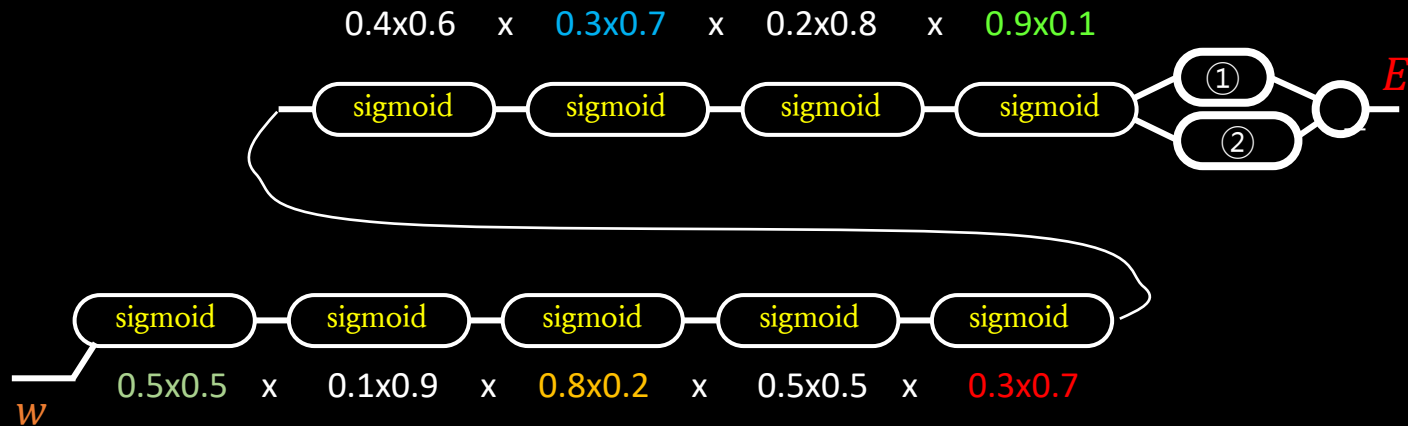
3-layer NN \rightarrow 2 sigmoid



10-layer Neural Network

The giant monster, computational graph!

Influence of w change on E



$$\frac{\partial E}{\partial w} = ?$$

Hint: chain rule!

Multiplication of the 9 derivative of
sigmoid

Vanishing Gradient

- The **derivative** of a **sigmoid** function is **sigmoid x (1-sigmoid)**
- 2 multiplication of sigmoid for a single neuron, 18 multiplications for 9 connected neurons
- Each **sigmoid** gives us the value between 0 and 1.

Ex) $0.5 \times 0.5 \times 0.1 \times 0.9 \times 0.8 \times 0.2 \times 0.5 \times 0.5 \times 0.3 \times 0.7 \times 0.4 \times 0.6 \times 0.3 \times 0.7 \times 0.2 \times 0.8 \times 0.9 \times 0.1$
= 0.00000010886 x ①

Vanishing Gradient

- The influence of w change on E is calculated through *many* multiplications of the values between 0 and 1, which gives us almost 0.
- **Vanishing Gradient**
- $w = w - \alpha \cdot \text{Slope} (\approx 0)$
- $b = b - \alpha \cdot \text{Slope} (\approx 0)$
- Therefore, **no updates** in w and b

(Lab) 18.py

- XOR problem using 4-layer neural networks
- Failed owing to vanishing gradient



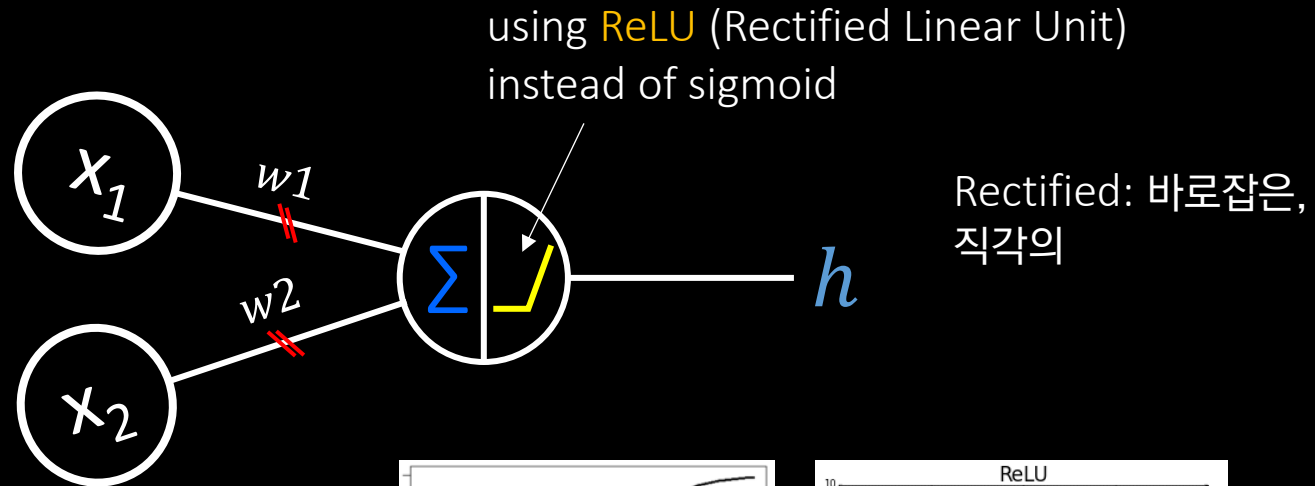
<https://github.com/yungbyun/neuralnetworks>

The **Dark Age** in AI (~2006)

since **back-propagation** by **Hinton** in 1986

ReLU

proposed by [Hahnloser](#) in 2000 and
demonstrated for deep networks
in 2011



(Lab) 19.py

- Solving vanishing gradient problem using ReLU activation function
- Back-propagation is working by using ReLU.

<https://github.com/yungbyun/neuralnetworks>

So, now can go deeper.

MNIST







Modified National Institute of
Standards and Technology
(USA)



MNIST

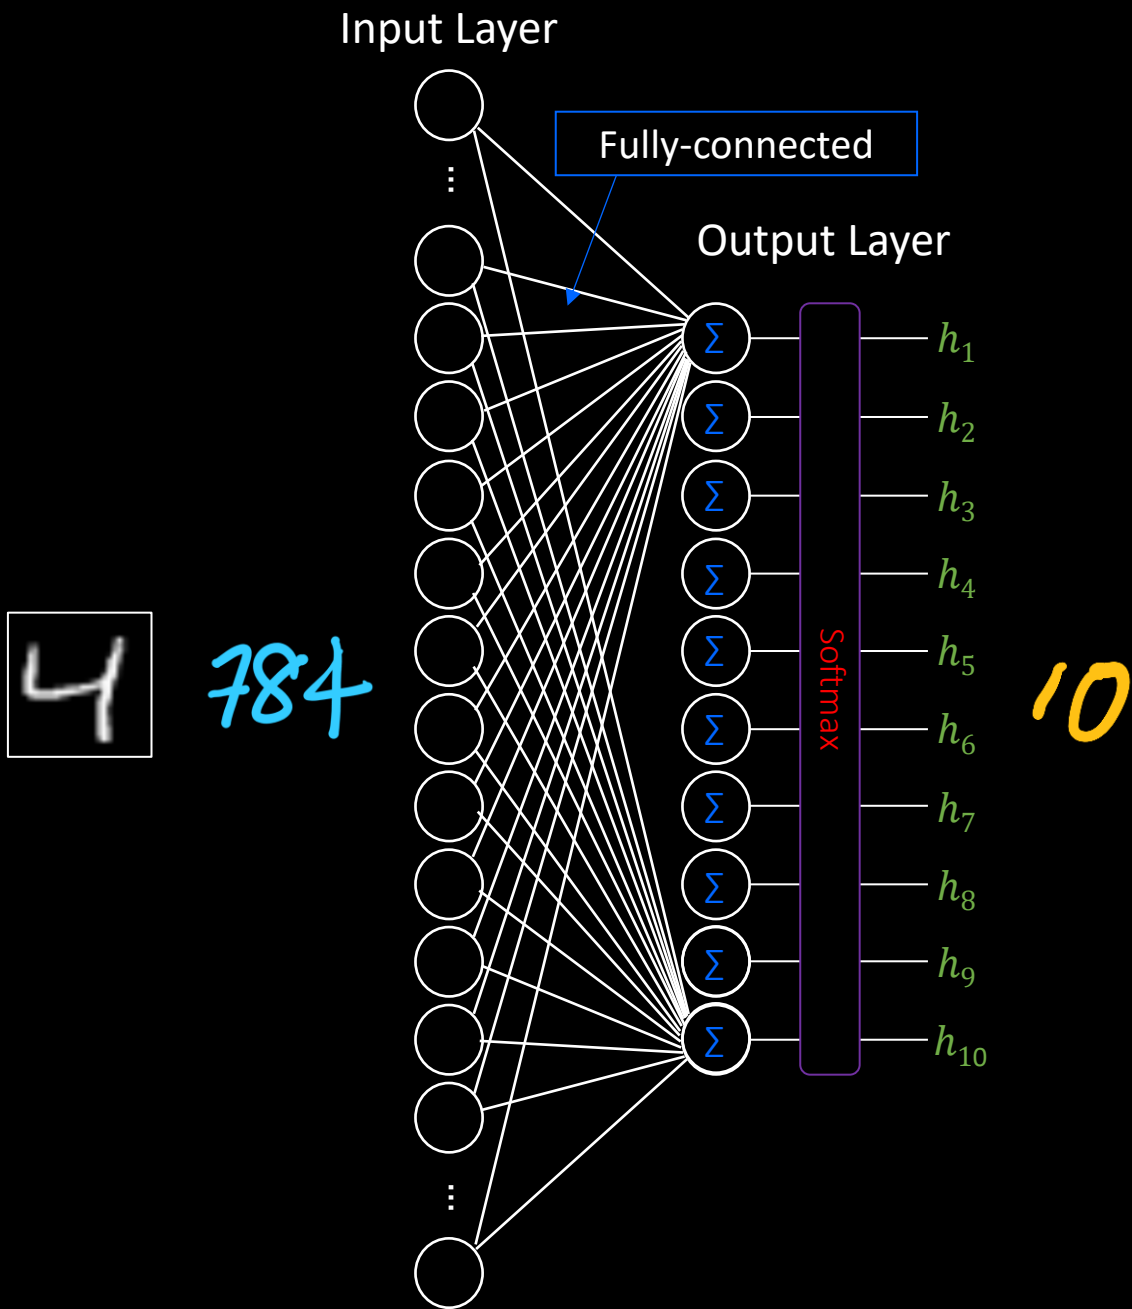





(Lab) 20.py

- 60,000 training images + 10,000 testing images
- Input image : $28 * 28$ pixels \rightarrow 784 pixels
- 784 input dimension
- 10 classes (output: 0 ~ 9)
- Softmax
- 90.23% of recognition rate


<https://github.com/yungbyun/neuralnetworks>





(Lab) 21.py

- Deep Neural Network (4-layer)
- ReLU
- 94.55% accuracy




<https://github.com/yungbyun/neuralnetworks>

(Lab) 22.py

- Applying **initialization** method for w and b , not randomly
- **97.23%** of accuracy


<https://github.com/yungbyun/neuralnetworks>

(Lab) 23.py



- Applying initialization
- 6-layer deep neural networks
- 97.83% of accuracy


Decision Boundry



“

Lots of neurons (connections, synapses) give us
so **complicated** decision boundary.

Which do you think is desirable decision boundary?



*While the black line fits the data well,
the green line is overfit.*


<https://elitedatascience.com>

Overfitting and drop-out


- The deeper the network is, the more the decision boundary is complex.
- Doing good at training data but errors for test data → overfitted to the training data
- Making it less complex by drop-out some neurons while learning.

Regularization: **Dropout**

“randomly set some neurons to zero in the forward pass”



(a) Standard Neural Net



(b) After applying dropout.

[Srivastava et al., 2014]


(Lab) 24.py

- Applying dropout
- 98.13% (\leftarrow 97.83%) of recognition accuracy


How to Prevent Overfitting

- Train with more data
- Reduce features
- Early stopping before overfitting
- Ensemble
- Regularization


Early stopping

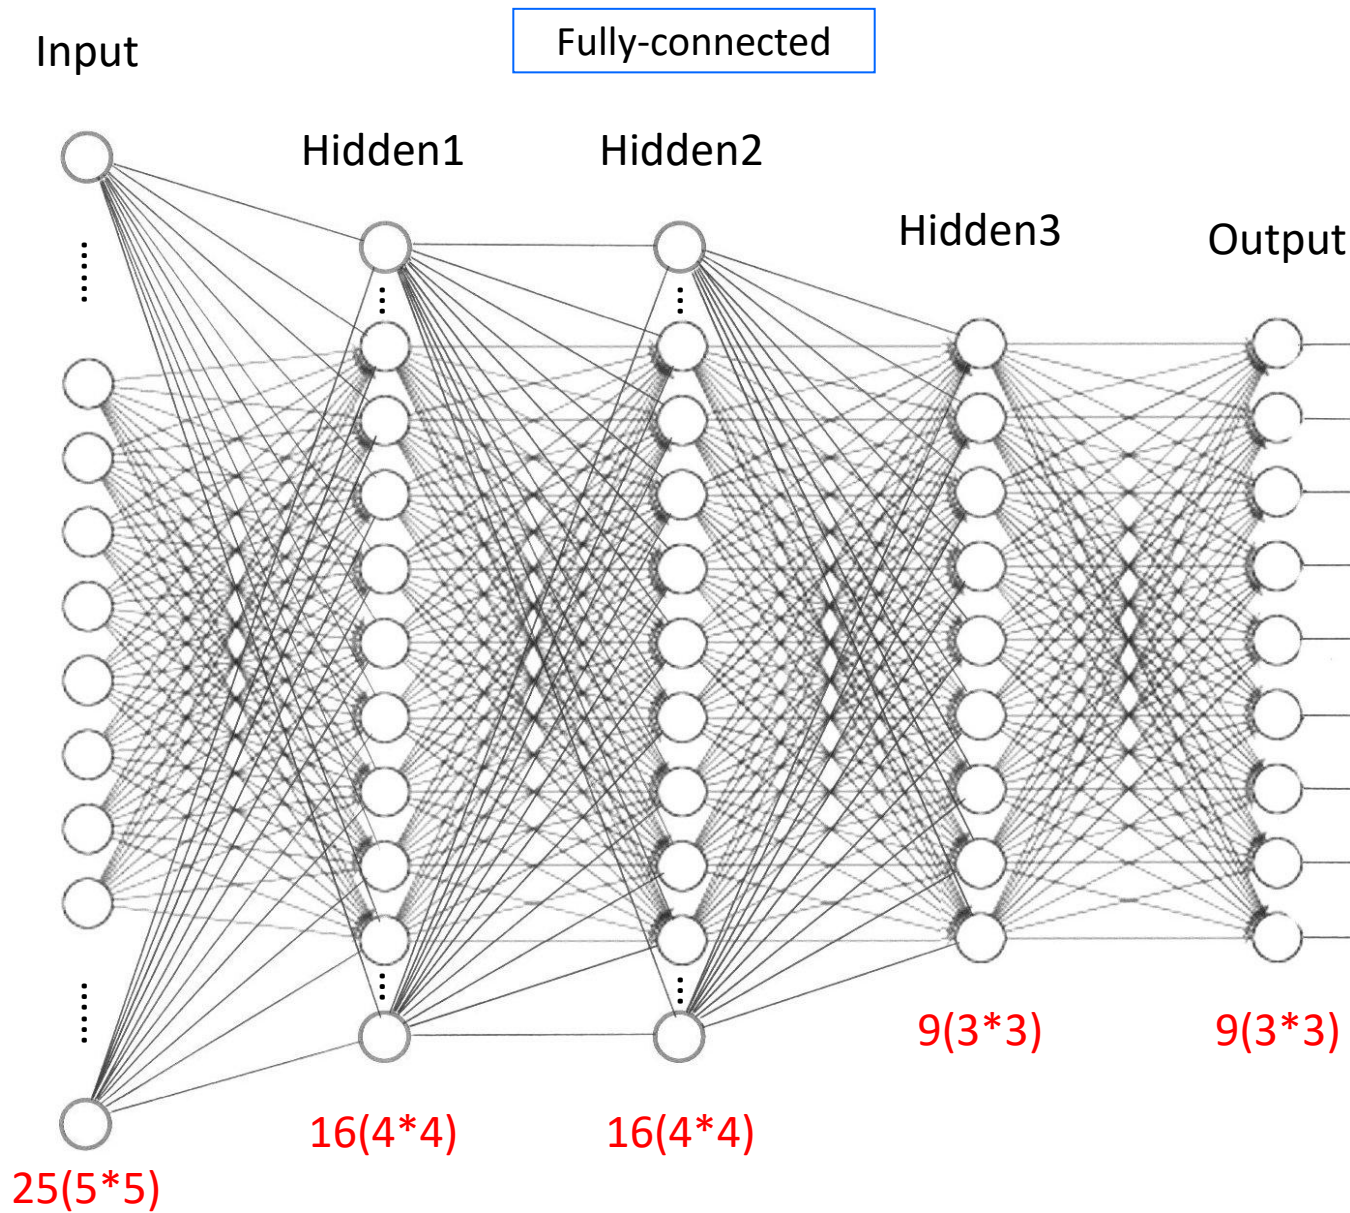


Fully-connected



Fully-connected






Fully connected, then how many connections(synapses, parameters) are there?


$$25 * 16 + 16 * 16 + 16 * 9 + 9 * 9 = 881$$



Fully-connected



Fully-connected

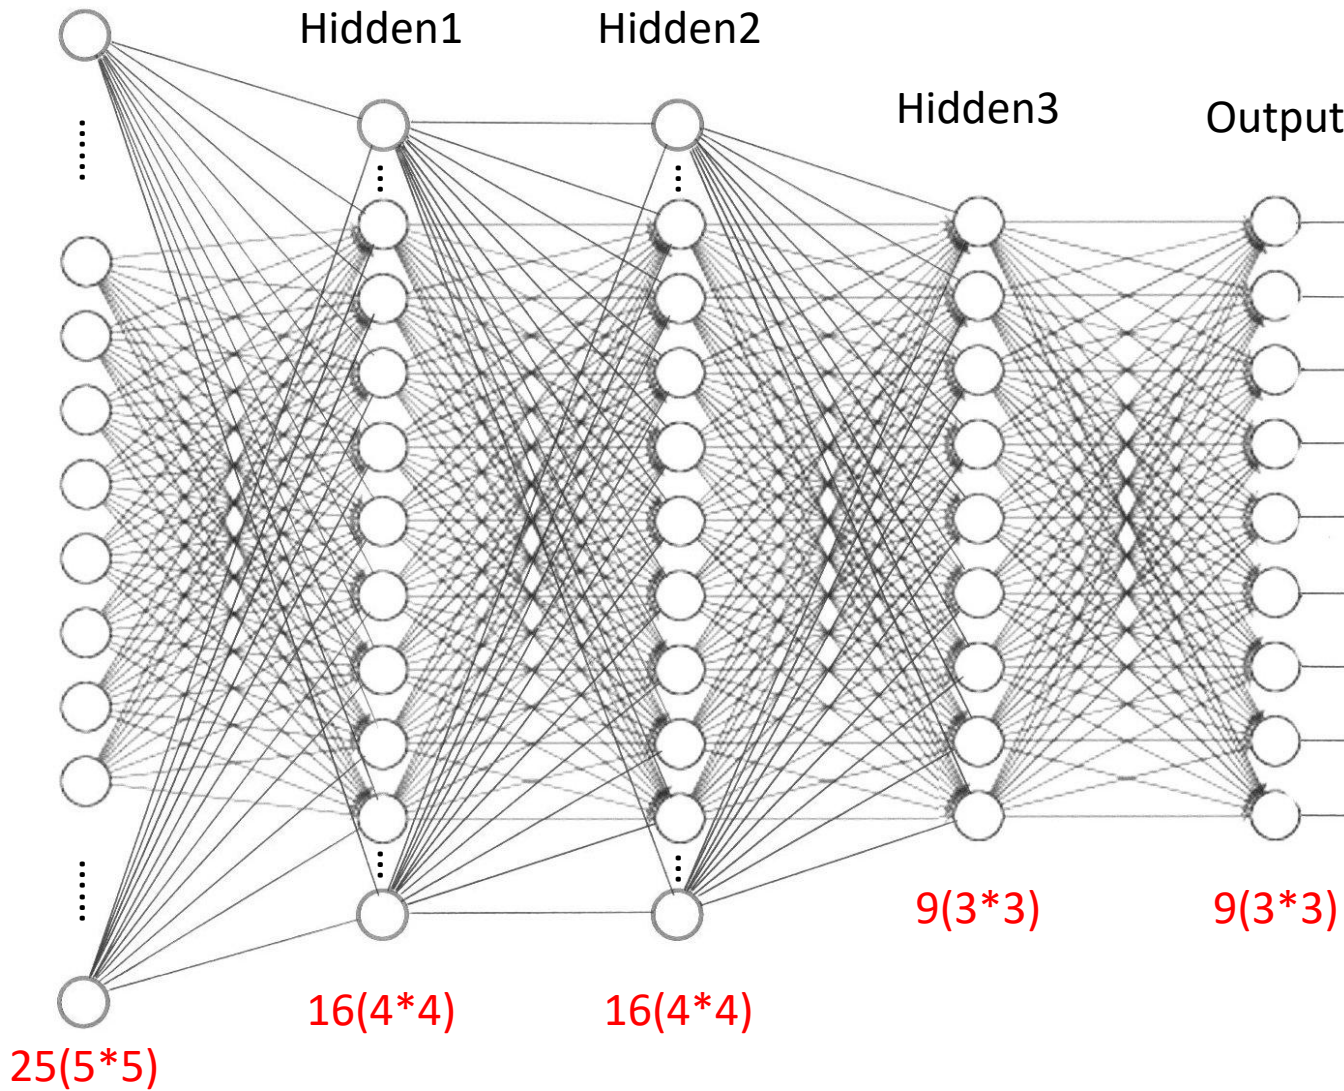


Fully connected, so how many connections are there?

$$25 * 16 + 16 * 16 + 16 * 9 + 9 * 9 = 881$$

Input

Fully-connected



Fully connected, then how many connections(synapses, parameters) are there?

$$25 * 16 + 16 * 16 + 16 * 9 + 9 * 9 = 881$$




Geoffrey Hinton, Yann LeCun, Yoshua Bengio, Andrew Ng



Deep Learning

- in early 2000s (2006, 2010, 2012)
- Deep Neural Networks
- Activation functions (ReLU)
- Weight initialization methods
- Dropout (2014)
- Big data
- GPU

Fully-connected



FCNN

Fully-connected Neural Networks

Any problem? Yes, so **CNN**.