

AI and Deep Learning

CNN, Convolutional Neural Network

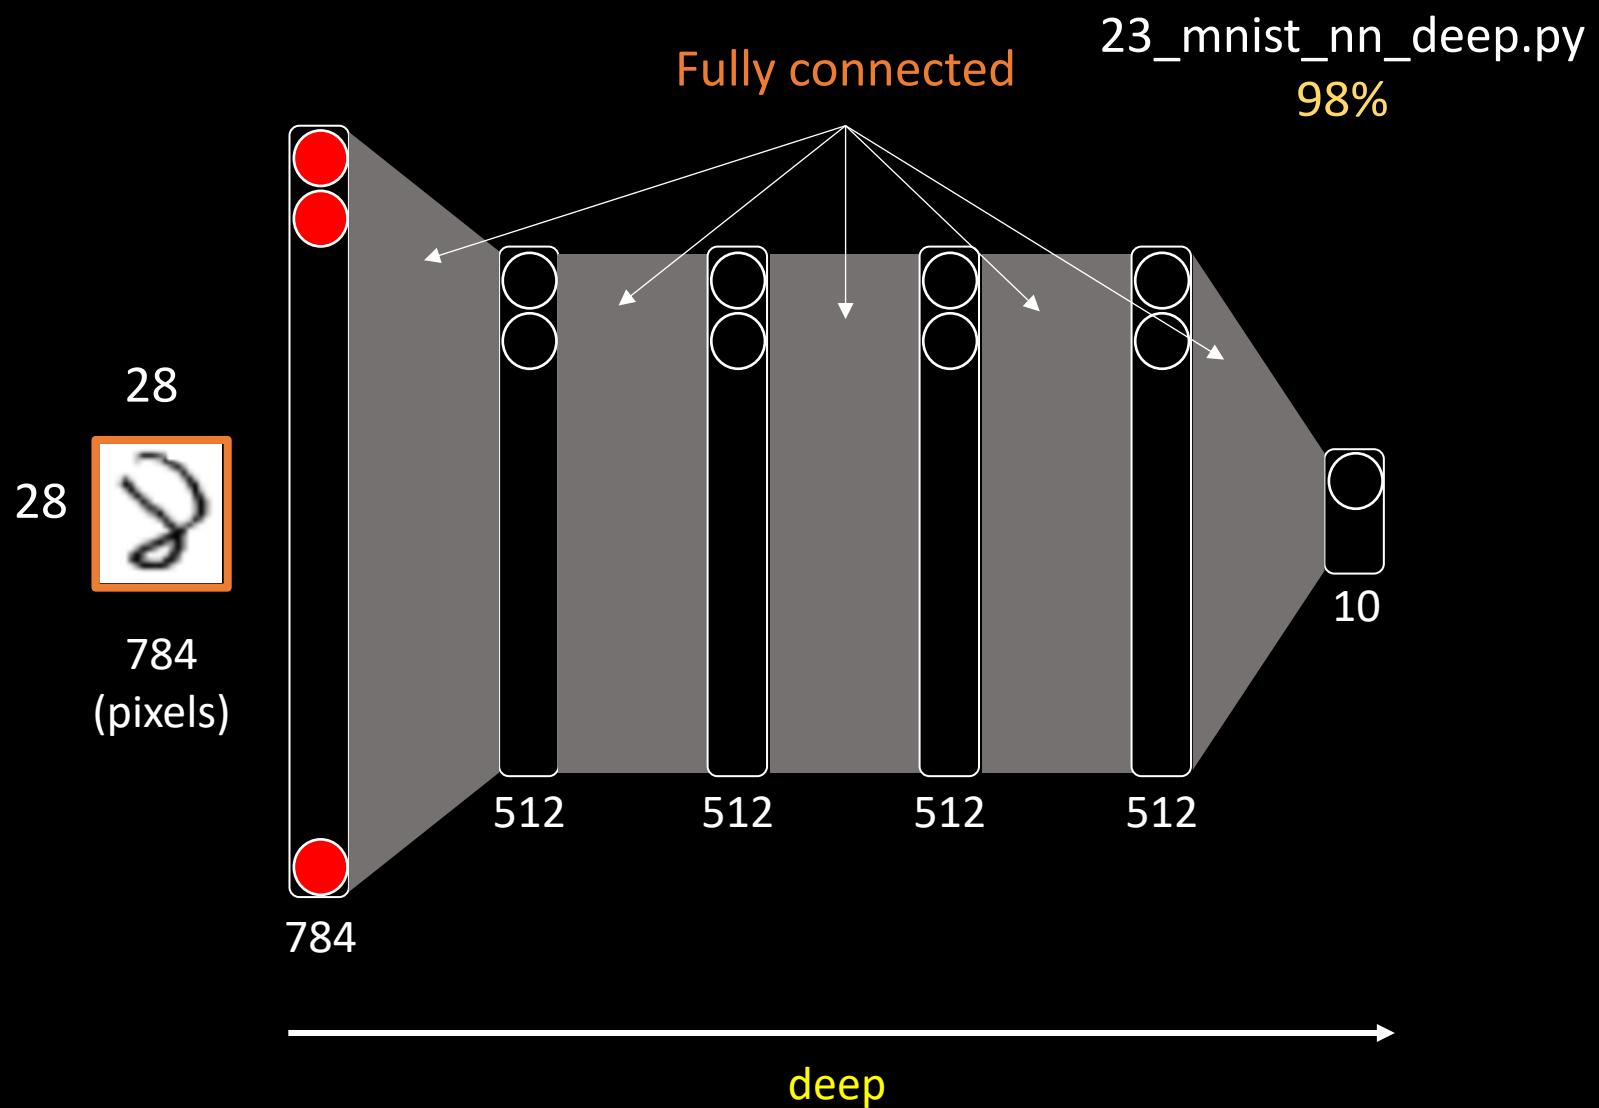
Jeju National University

Yungcheol Byun

Agenda

- Drawbacks of FCNN
- Feature and Filter
- Activation Maps
- Pooling
- Theoretical Background
- Case Study

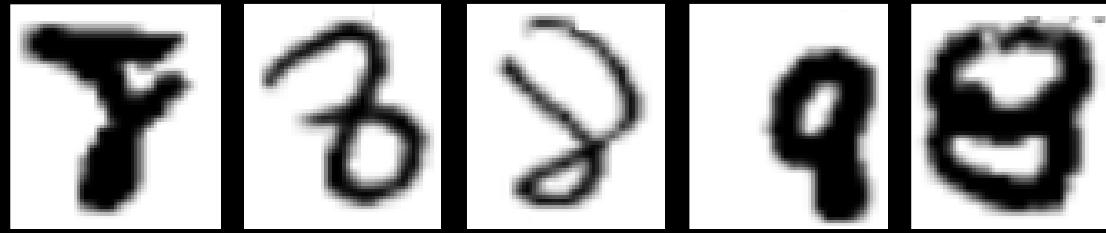
Drawbacks of Fully Connected NN





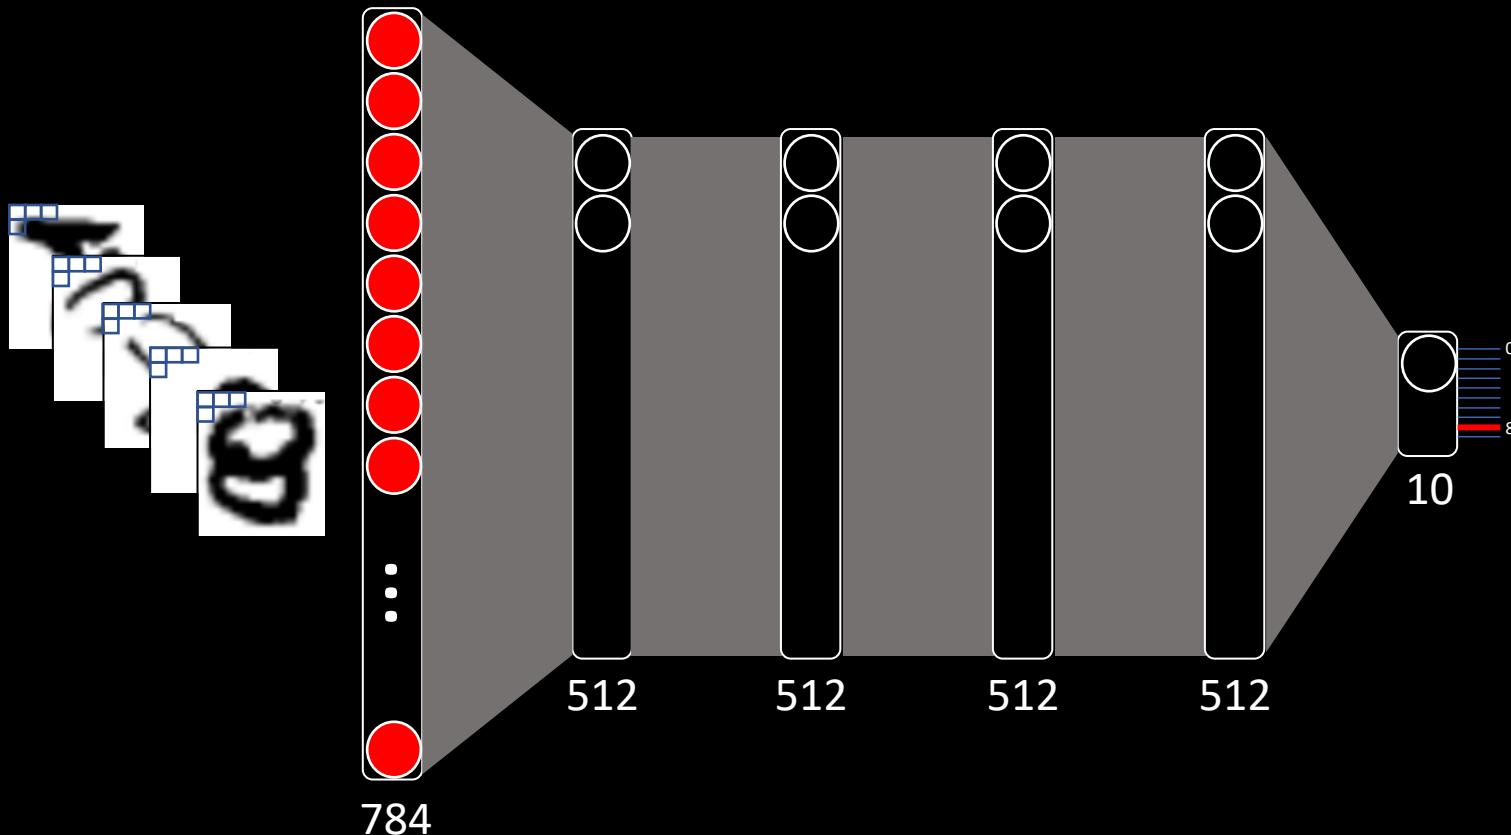
- Various shapes - size, location, skewness, distortion



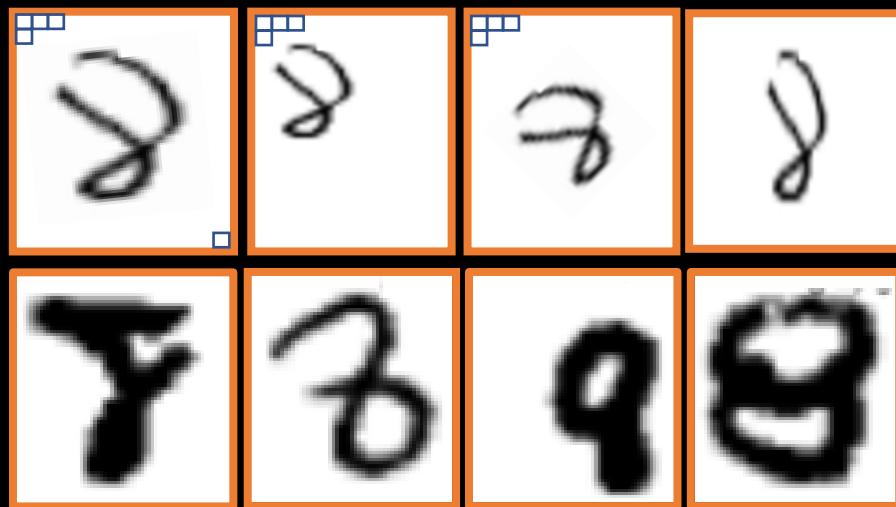


Many variations of eight(8)
from **writing styles**

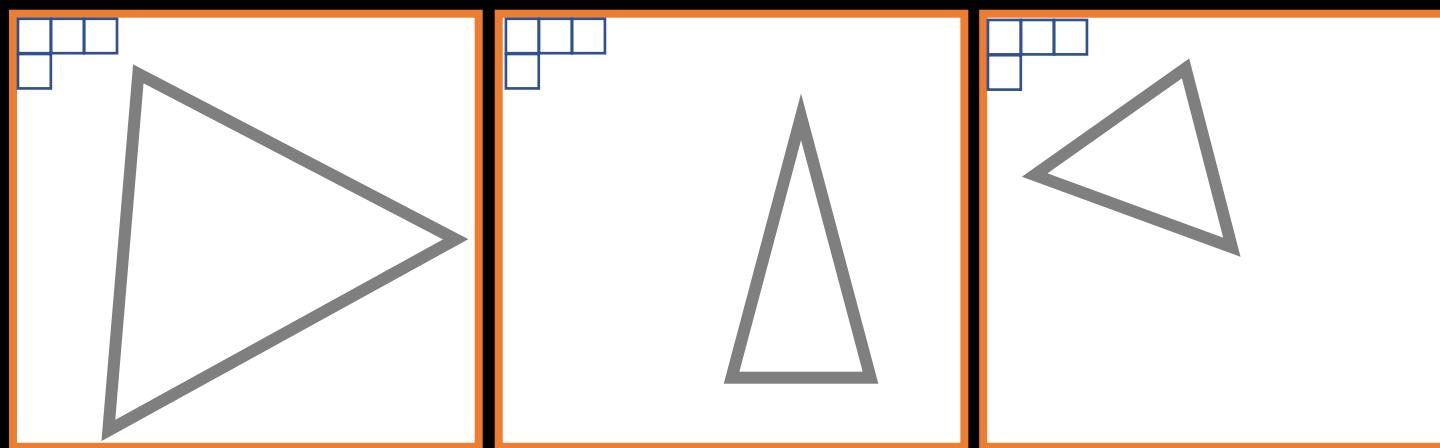
Same digits but
different sets of values
as inputs



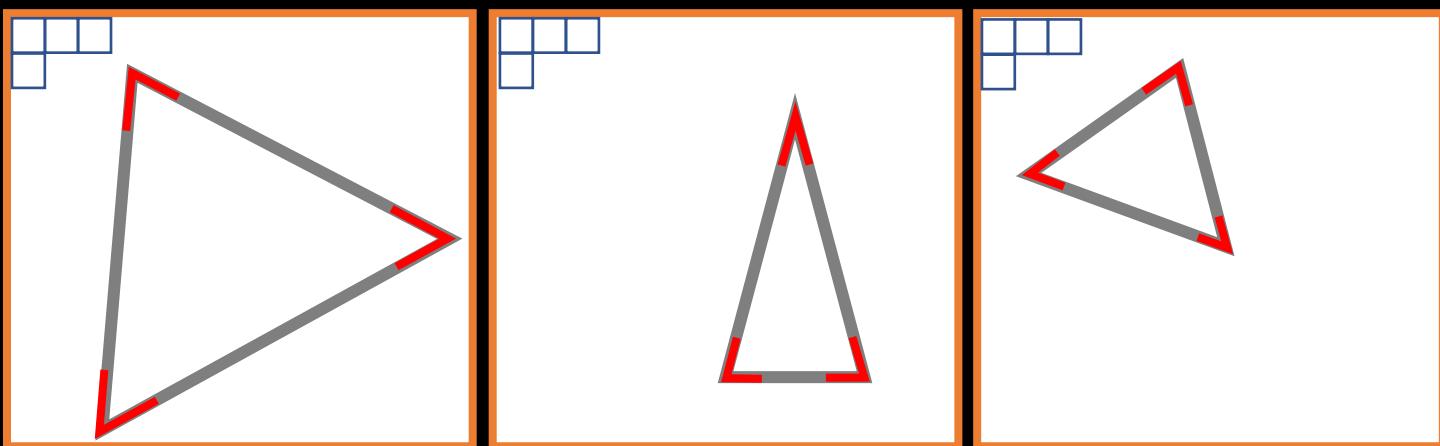
Any common features?



Any common features?

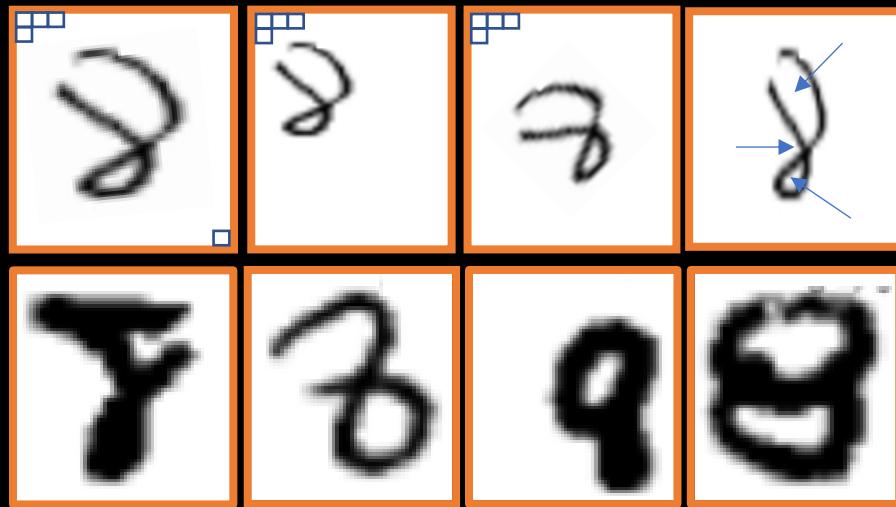


Any common features?

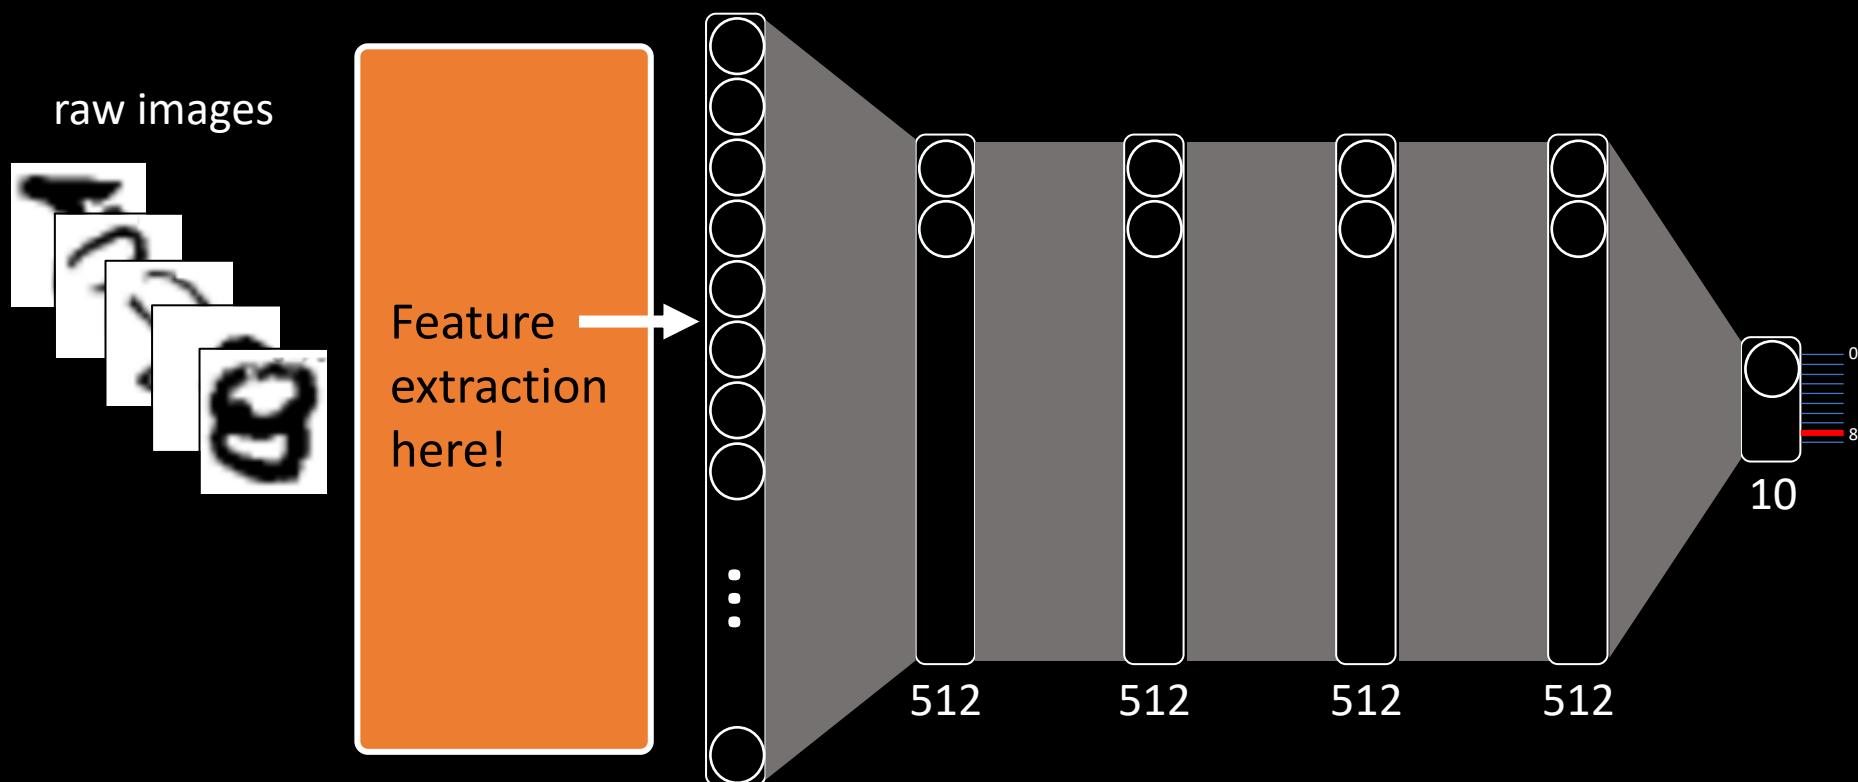


Redundant information

Any common features?



Feature extraction for the FCNN

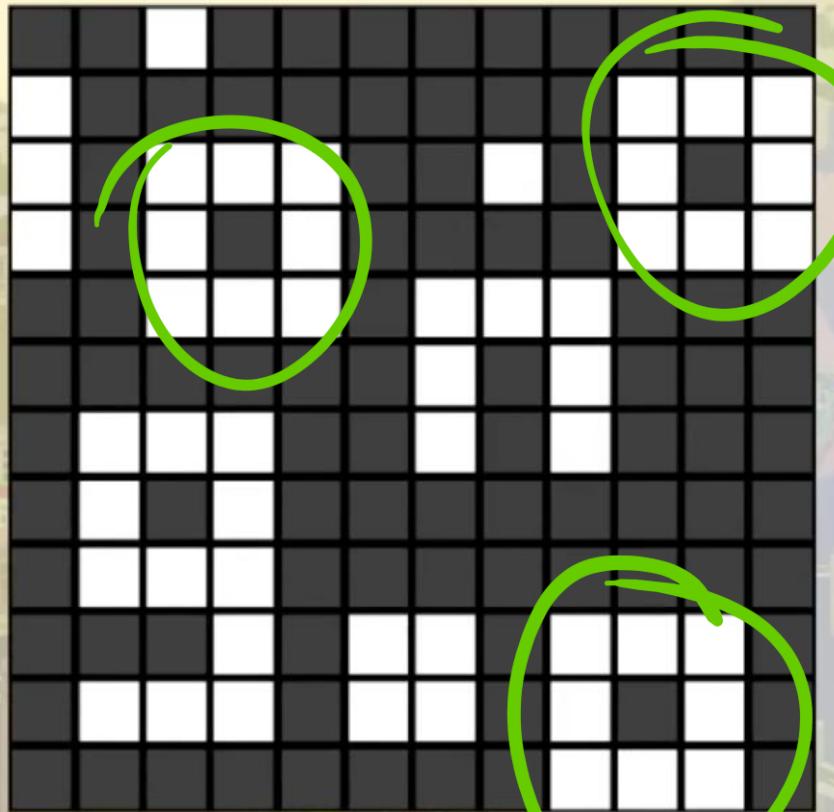


Instead of using the raw image as input,
we extract features first and use them as input for the FCNN.

How to extract **features**

Doughnut features

도넛 특징을 찾아라!



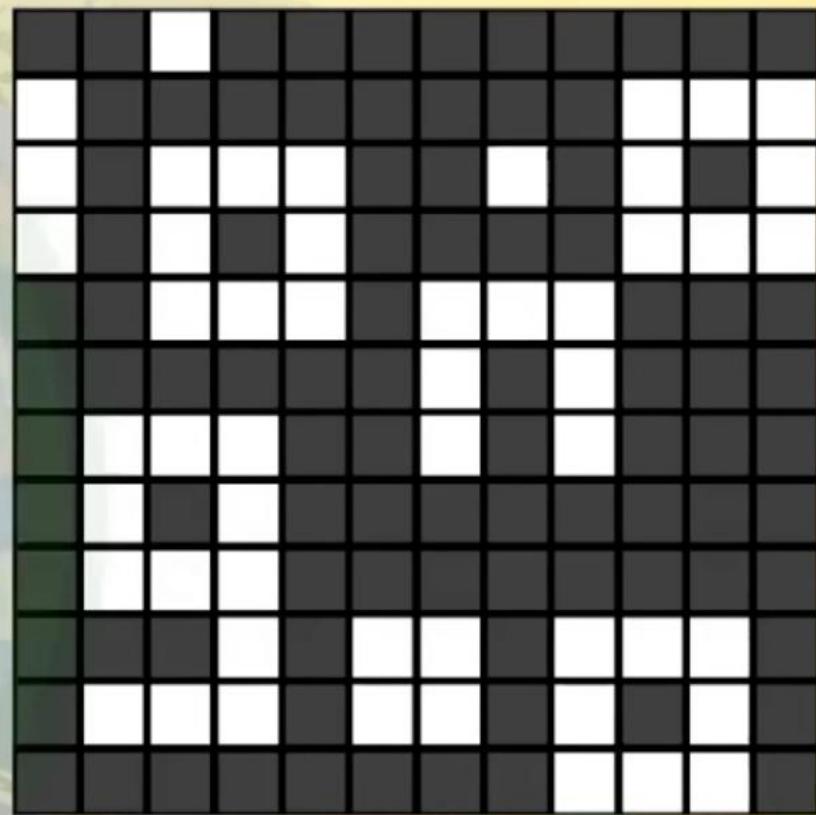
IMAGE



Filter to extract the doughnut feature

"DONUT FILTER"

DOESN'T MATTER	MUST BE BLACK	MUST BE BLACK	MUST BE BLACK	DOESN'T MATTER
MUST BE BLACK	MUST BE WHITE	MUST BE WHITE	MUST BE WHITE	MUST BE BLACK
MUST BE BLACK	MUST BE WHITE	MUST BE BLACK	MUST BE WHITE	MUST BE BLACK
MUST BE BLACK	MUST BE WHITE	MUST BE WHITE	MUST BE WHITE	MUST BE BLACK
DOESN'T MATTER	MUST BE BLACK	MUST BE BLACK	MUST BE BLACK	DOESN'T MATTER



How to extract doughnut features



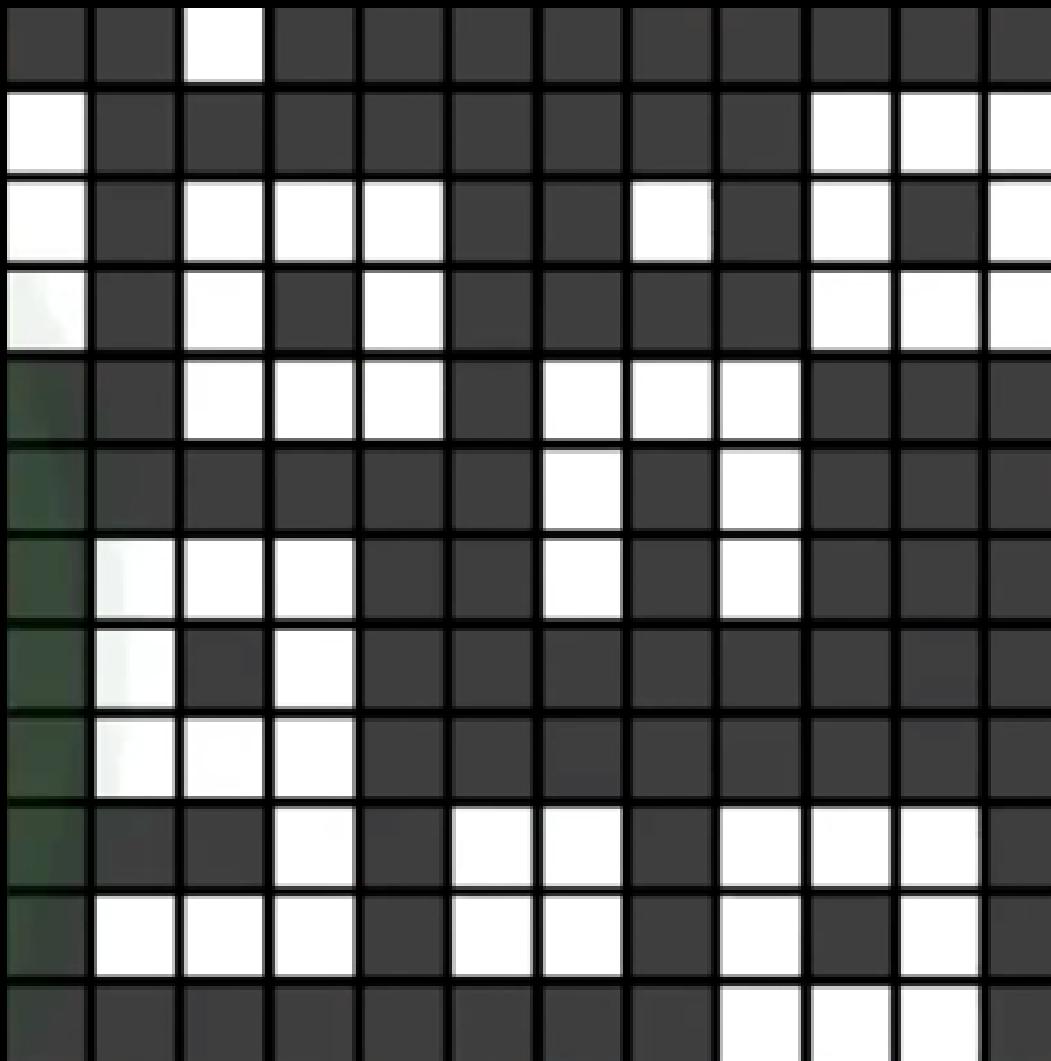
Out-of-bounds pixels are considered as black pixels.

"Is every condition of the filter satisfied?"

It's not a donut.

DOESN'T MATTER ✓ Yes	MUST BE BLACK ✓ Yes	MUST BE BLACK ✓ Yes	MUST BE BLACK ✓ Yes	DOESN'T MATTER ✓ Yes
MUST BE BLACK ✓ Yes	MUST BE WHITE ✗ No	MUST BE WHITE ✗ No	MUST BE WHITE ✗ No	MUST BE BLACK ✓ Yes
MUST BE BLACK ✓ Yes	MUST BE WHITE ✗ No	MUST BE BLACK ✓ Yes	MUST BE WHITE ✗ No	MUST BE BLACK ✗ No
MUST BE BLACK ✓ Yes	MUST BE WHITE ✗ No	MUST BE WHITE ✓ Yes	MUST BE WHITE ✗ No	MUST BE BLACK ✓ Yes
DOESN'T MATTER ✓ Yes	MUST BE BLACK ✓ Yes	MUST BE BLACK ✗ No	MUST BE BLACK ✓ Yes	DOESN'T MATTER ✓ Yes

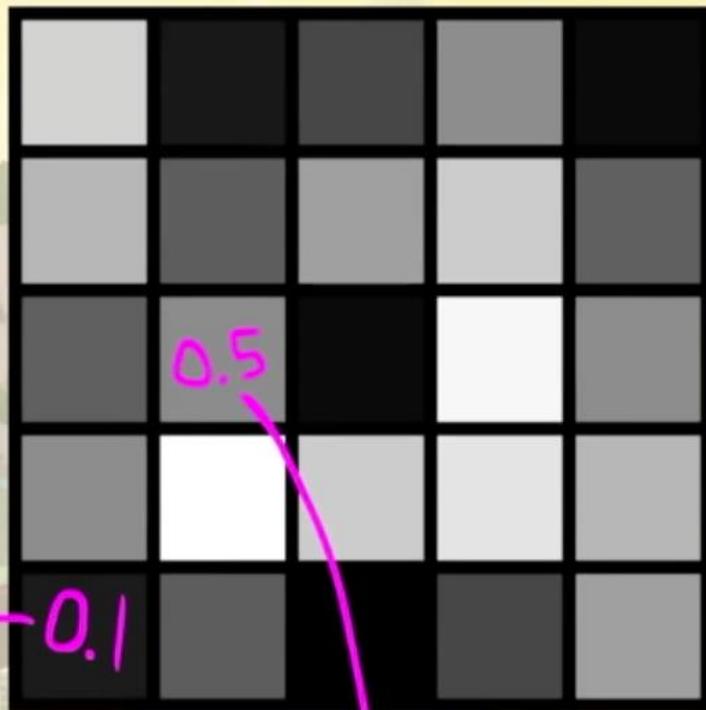
Black & White



Real images are color or gray



An example of a gray image



0

0.0

255

1.0

~~DONUT FILTER~~

DOESN'T MATTER	MUST BE BLACK	MUST BE BLACK	MUST BE BLACK	DOESN'T MATTER
MUST BE BLACK	MUST BE WHITE	MUST BE WHITE	MUST BE WHITE	MUST BE BLACK
MUST BE BLACK	MUST BE WHITE	MUST BE BLACK	MUST BE WHITE	MUST BE BLACK
MUST BE BLACK	MUST BE WHITE	MUST BE WHITE	MUST BE WHITE	MUST BE BLACK
DOESN'T MATTER	MUST BE BLACK	MUST BE BLACK	MUST BE BLACK	DOESN'T MATTER

New improved filter

x0.0	x-0.2	x-0.1	x-0.2	x0.0
x-0.2	x0.7	x1.0	x0.7	x-0.2
x-0.1	x1.0	x-2.5	x1.0	x-0.1
x-0.2	x0.7	x1.0	x0.7	x-0.2
x0.0	x-0.1	x-0.1	x-0.1	x0.0

"set of multipliers"

New improved filter

"If you want to
be considered
donutty..."



x-0.2	x-0.1	x-0.2	x0.0
x0.7	x1.0	x0.7	x-0.2
0.1	x1.0	x-2.5	x-0.1
x-0.2	x0.7	x1.0	x0.7
x0.0	x-0.1	x-0.1	x0.0

...you'd better
have a high
value for this
pixel."

Input Image (입력 이미지)

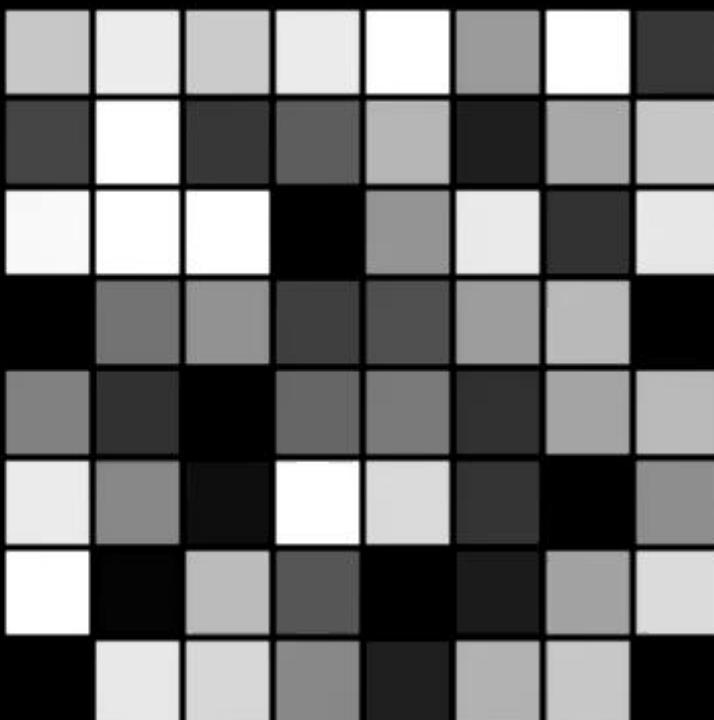
THE FILTER ↴

x0.0	x-0.2	x-0.1	x-0.2	x0.0
x-0.2	x0.7	x1.0	x0.7	x-0.2
x-0.1	x1.0	x-2.5	x1.0	x-0.1
x-0.2	x0.7	x1.0	x0.7	x-0.2
x0.0	x-0.1	x-0.1	x-0.1	x0.0

0.73	0.86	0.74	0.85	0.99	0.58	0.96	0.25
0.29	0.92	0.25	0.38	0.67	0.16	0.63	0.73
0.90	0.99	0.95	0.06	0.56	0.85	0.23	0.84
0.00	0.45	0.55	0.28	0.33	0.59	0.69	0.01
0.50	0.23	0.03	0.11	0.48	0.23	0.62	0.69
0.85	0.52	0.11	0.95	0.80	0.24	0.03	0.54
0.99	0.07	0.69	0.35	0.05	0.15	0.61	0.80
0.03	0.84	0.79	0.52	0.17	0.66	0.73	0.01

0.00	x-0.20	x-0.10	x-0.20	x0.00					
0.20	x0.70	x1.00	x0.70	x-0.20					
0.10	x1.00	0.73 x-2.50	0.86 x1.00	0.74 x-0.10	0.85	0.99	0.58	0.96	0.25
0.20	x0.70	0.29 x1.00	0.92 x0.70	0.25 x-0.20	0.38	0.67	0.16	0.63	0.73
0.00	x-0.20	0.90 x-0.10	0.99 x-0.20	0.95 x0.00	0.06	0.56	0.85	0.23	0.84
	0.00	0.45	0.55	0.28	0.33	0.59	0.69	0.01	
	0.50	0.23	0.03	0.41	0.48	0.23	0.62	0.69	
	0.85	0.52	0.11	0.95	0.80	0.24	0.03	0.54	
	0.99	0.07	0.69	0.35	0.05	0.15	0.61	0.80	
	0.03	0.84	0.79	0.52	0.17	0.66	0.73	0.01	

How donutty is each pixel?



Convolution?

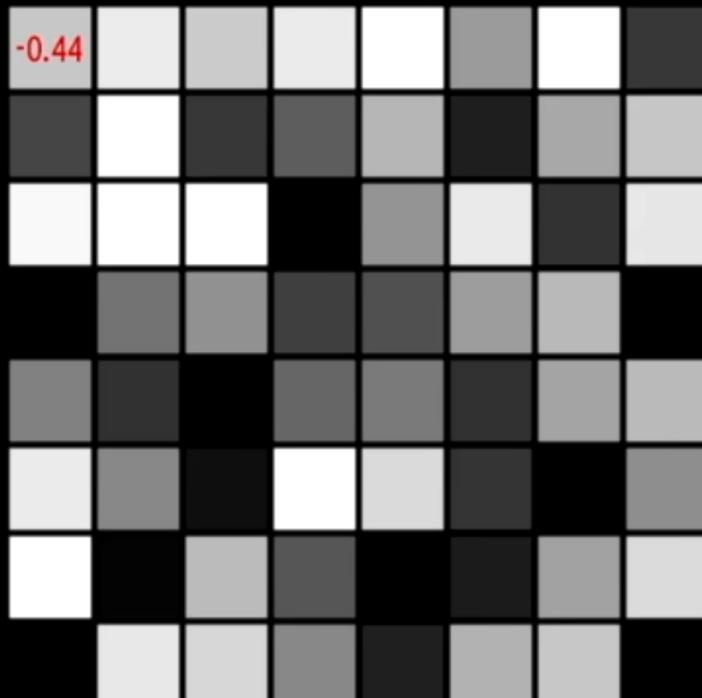
dot product with the image
and slide(move) the filter spatially to the next right pixel.

Filter

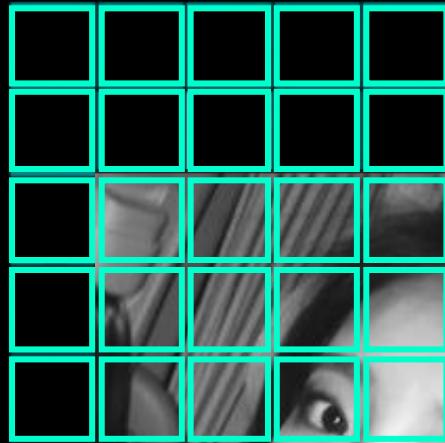
x0.00 0.00	x-0.20 -0.00	x-0.10 -0.00	x-0.20 -0.00	x0.00 0.00				
x-0.20 -0.00	x0.70 0.00	x1.00 0.00	x0.70 0.00	x-0.20 -0.00				
x-0.10 -0.00	0.73 x1.00 0.73	0.86 x-2.50 -2.14	0.74 x1.00 0.74	0.85 x-0.10 -0.09	0.99	0.58	0.96	0.25
x-0.20 -0.00	0.29 x0.70 0.92	0.92 x1.00 0.92	0.25 x0.70 0.00	0.38 x-0.20 -0.08	0.67	0.16	0.63	0.73
x0.00 0.00	0.90 x-0.20 -0.18	0.99 x-0.10 -0.10	0.95 x-0.20 -0.19	0.06 x0.00 0.00	0.56	0.85	0.23	0.84
0.00	0.45	0.55	0.28	0.33	0.59	0.69	0.01	
0.50	0.23	0.03	0.41	0.48	0.23	0.62	0.69	
0.85	0.52	0.11	0.95	0.80	0.24	0.03	0.54	
0.99	0.07	0.69	0.35	0.05	0.15	0.61	0.80	
0.03	0.84	0.79	0.52	0.17	0.66	0.73	0.01	

Image

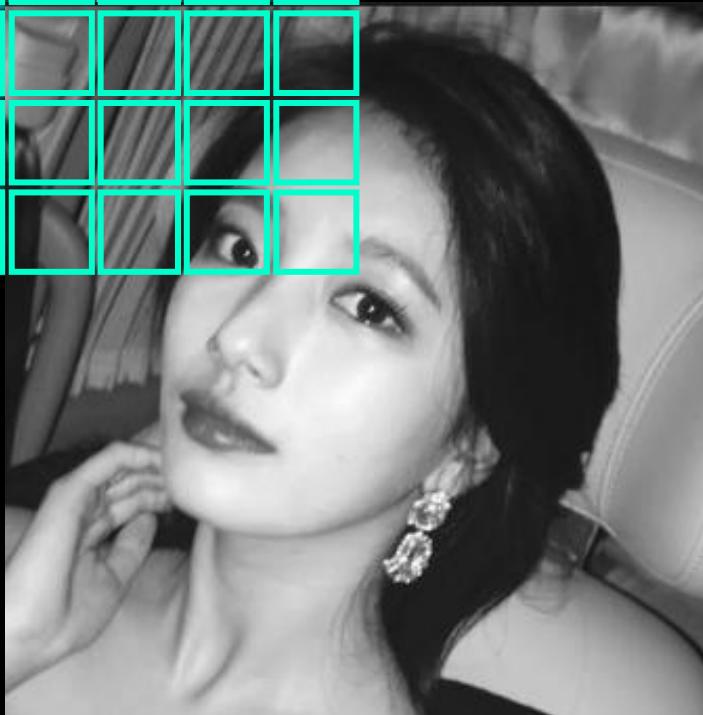
How donutty is each pixel?



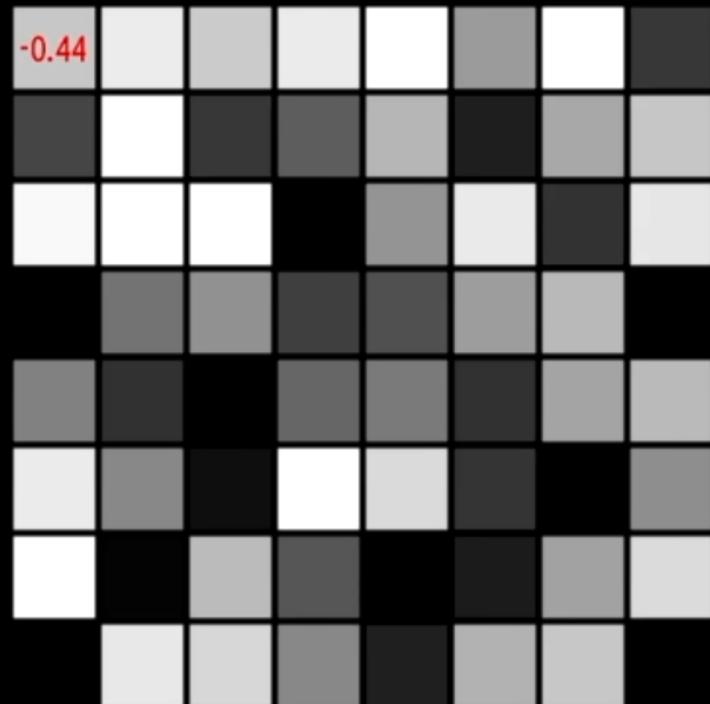
Filter



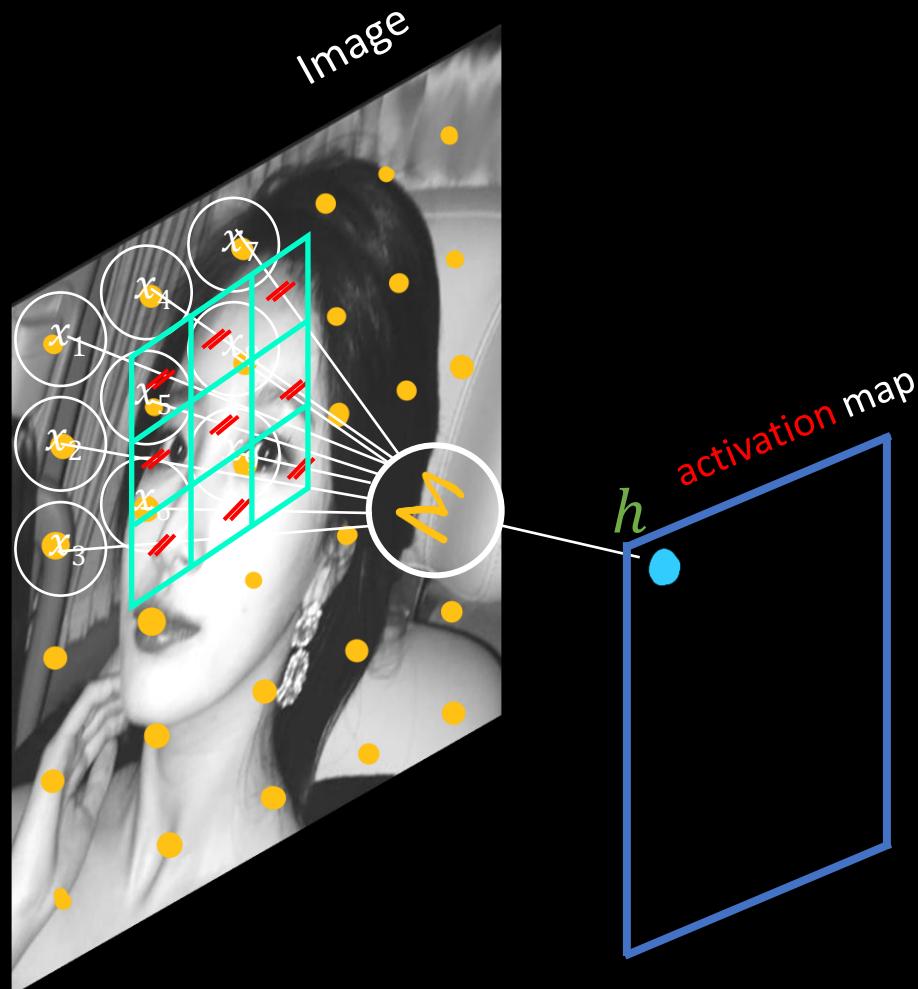
Image



activation map



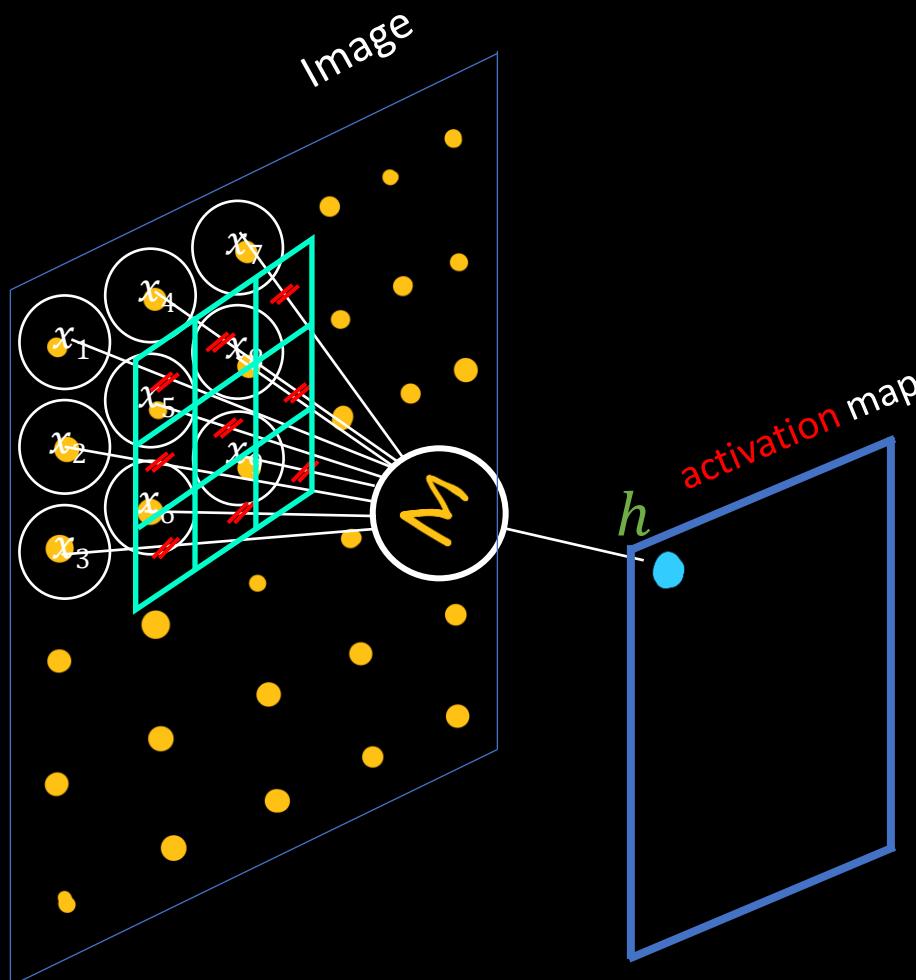
Convolution in 3D view



9 connections = 9 synaps (parameters)

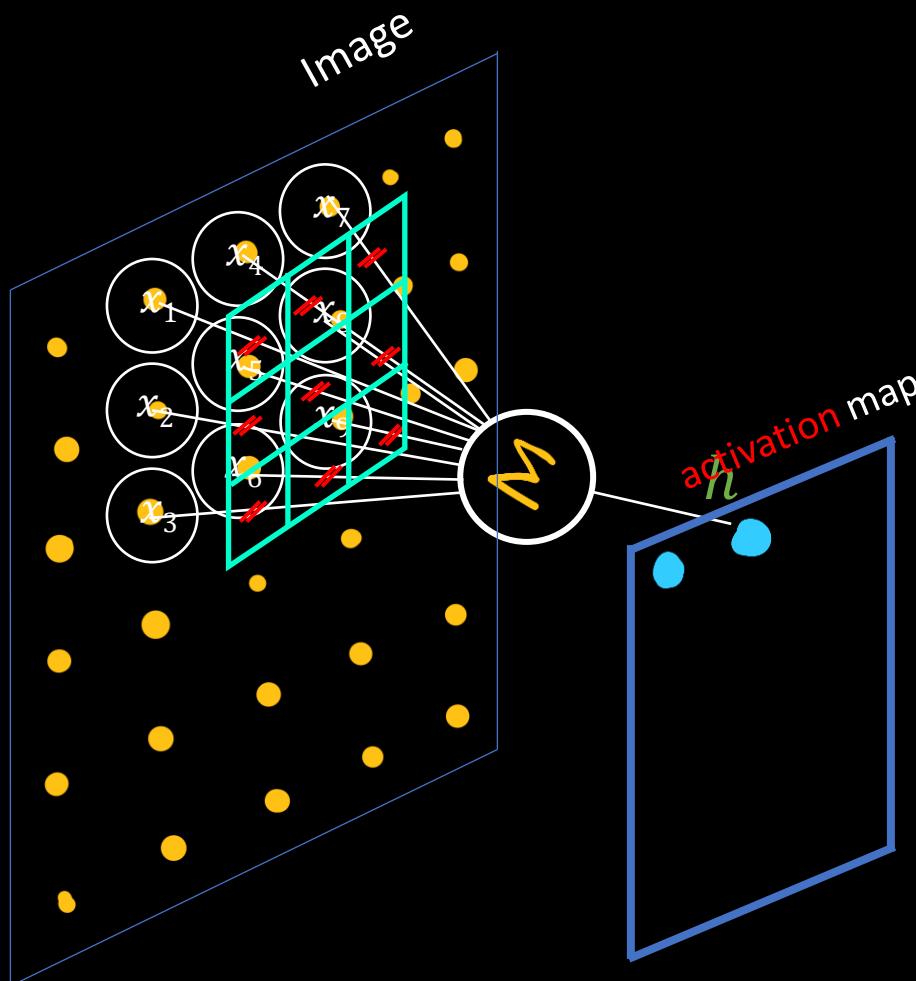
So, what is a **filter**?

Convolution in 3D view



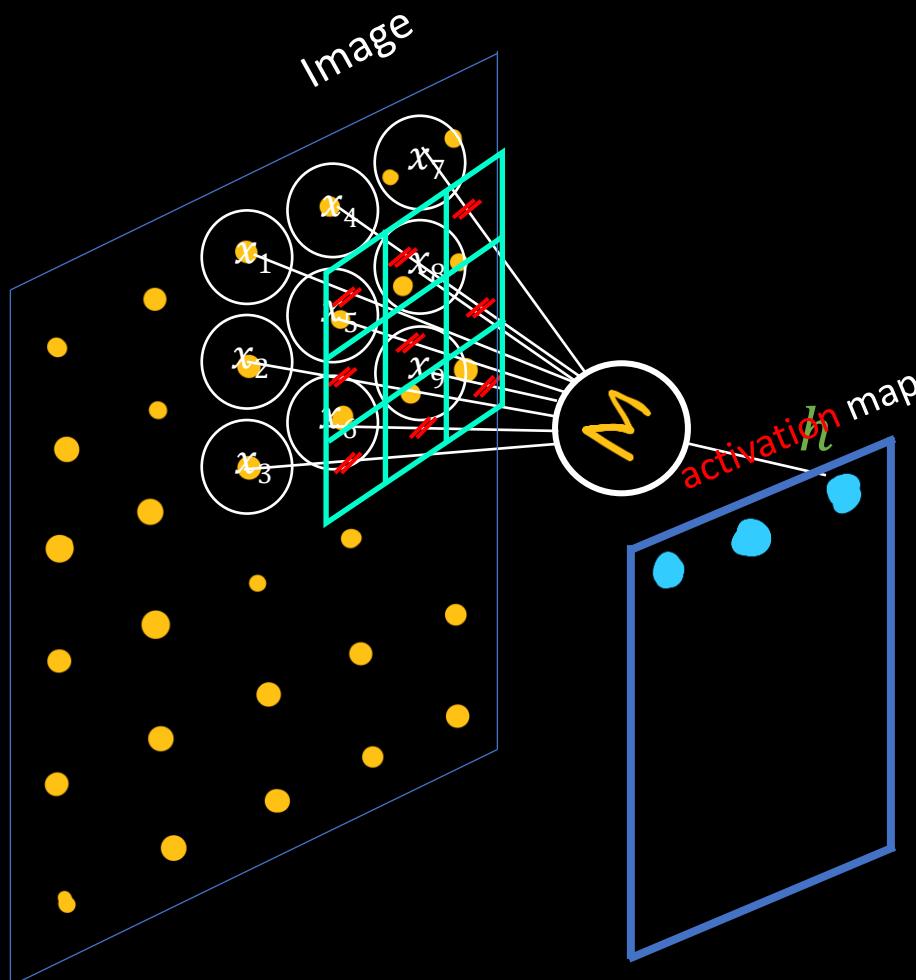
9 connections = 9 synaps (parameters)

Convolution in 3D view



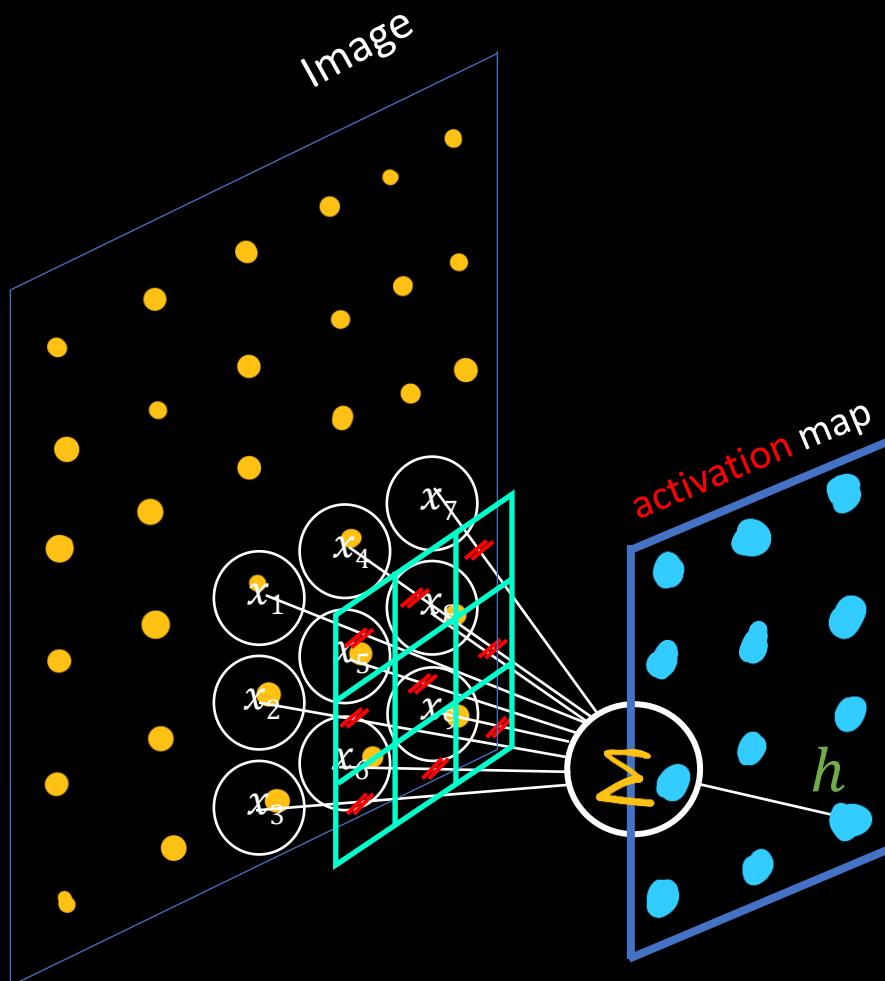
9 connections = 9 synaps (parameters)

Convolution in 3D view



9 connections = 9 synaps (parameters)

Convolution in 3D view

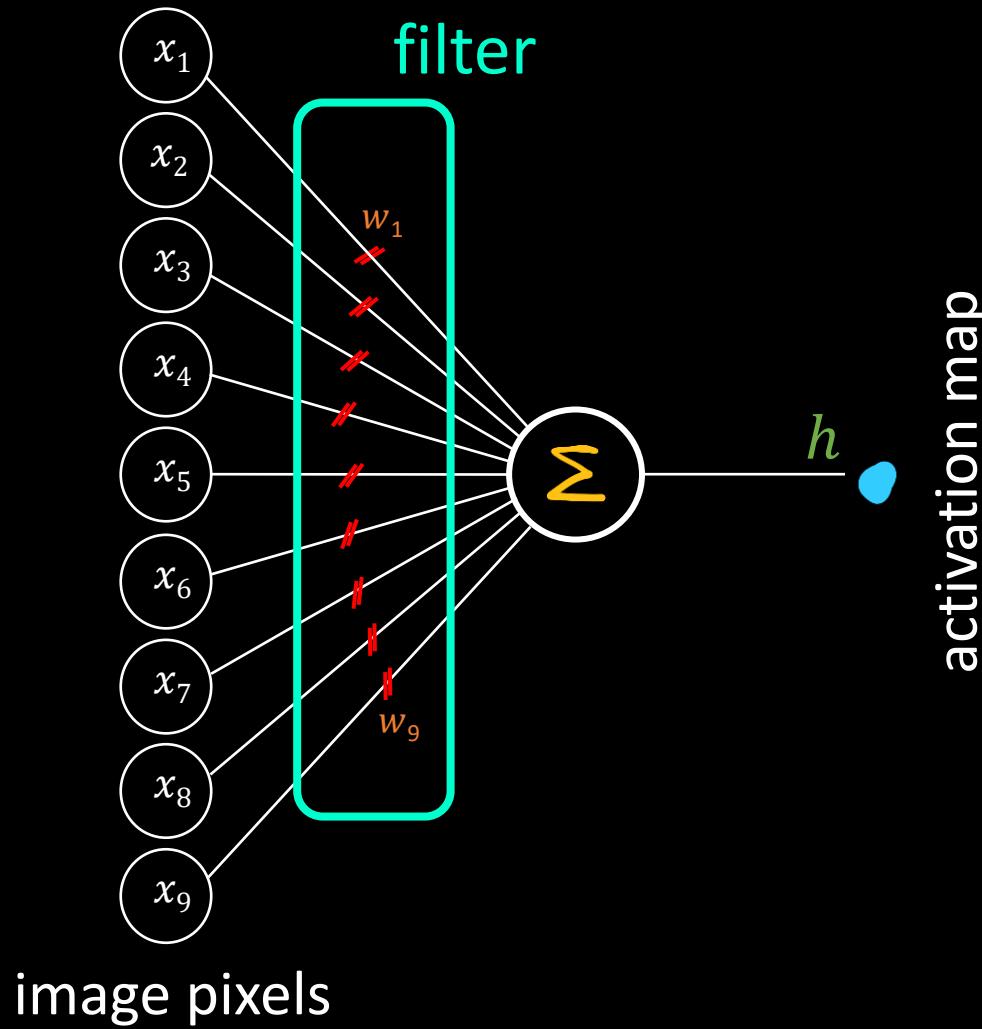


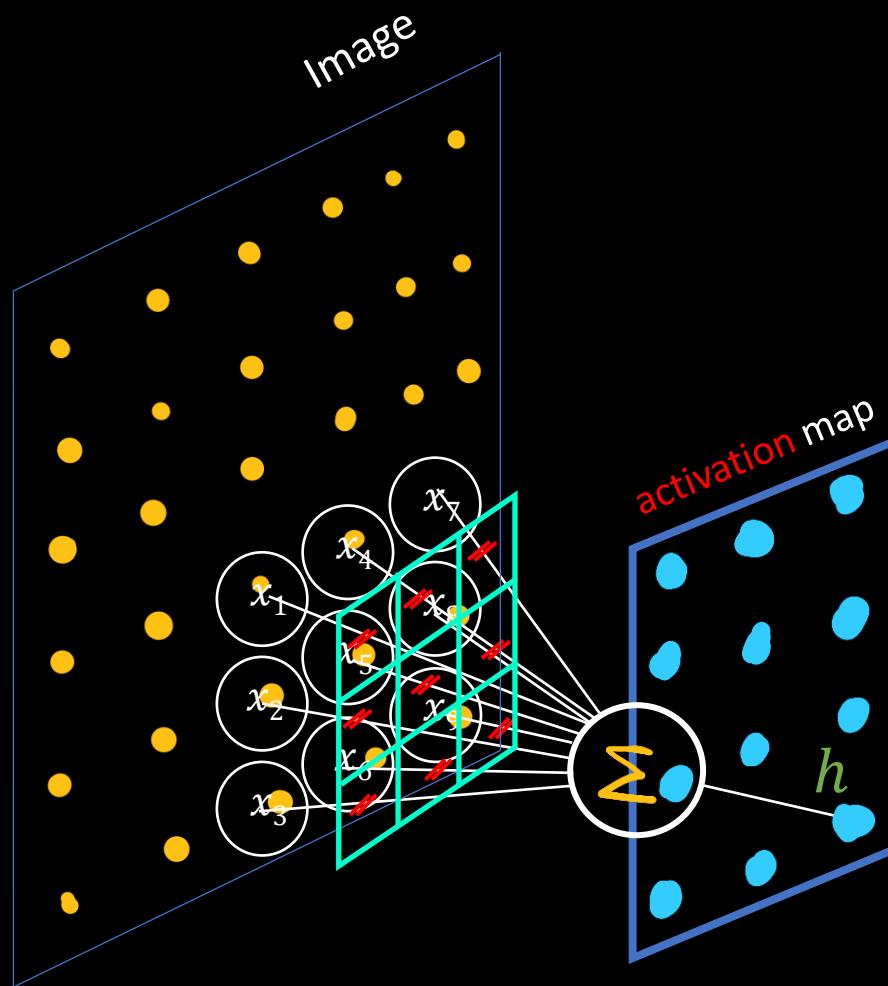
9 connections = 9 synaps (parameters)

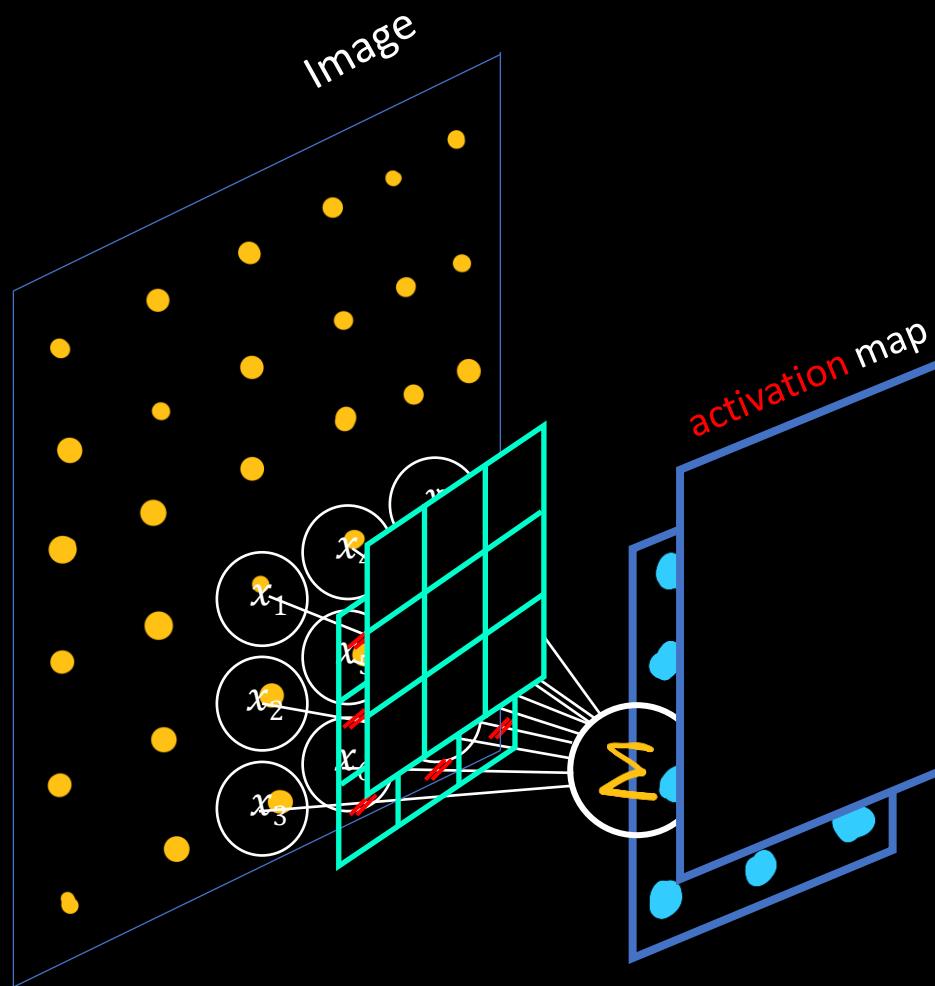
A filter

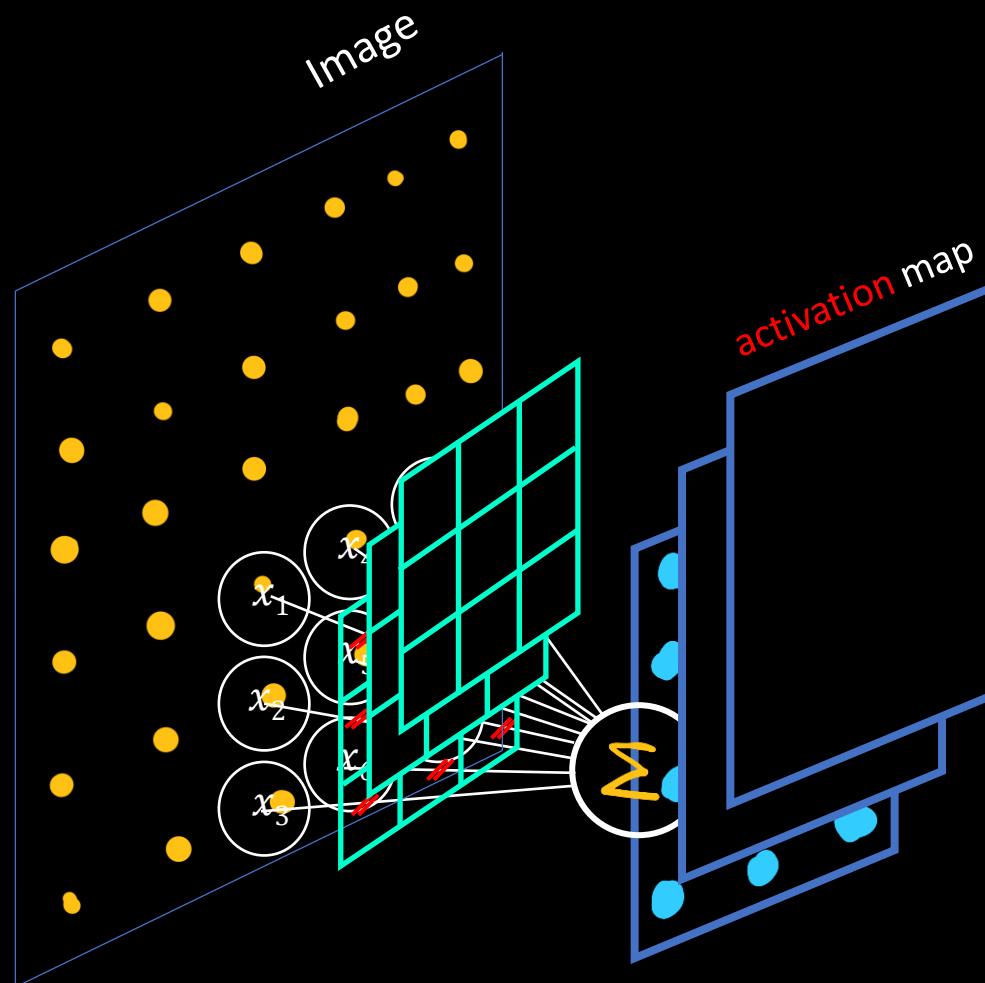
- a group of synapses(**w**)
- parameter tuning with back-propagation → automatic filter setting
- sharing the same synapses while **covolution**(dot-product & shifting)
- 1 filter → 1 activation map

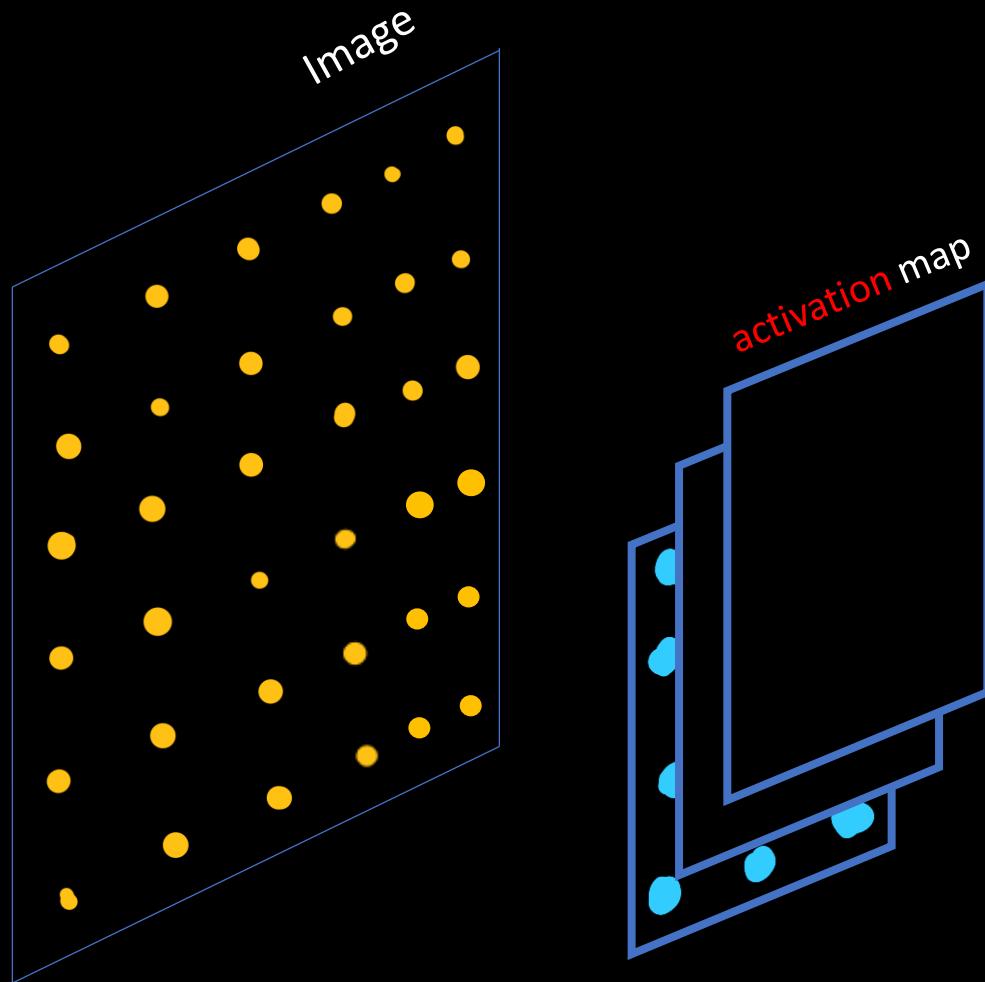
Filter(2D view)

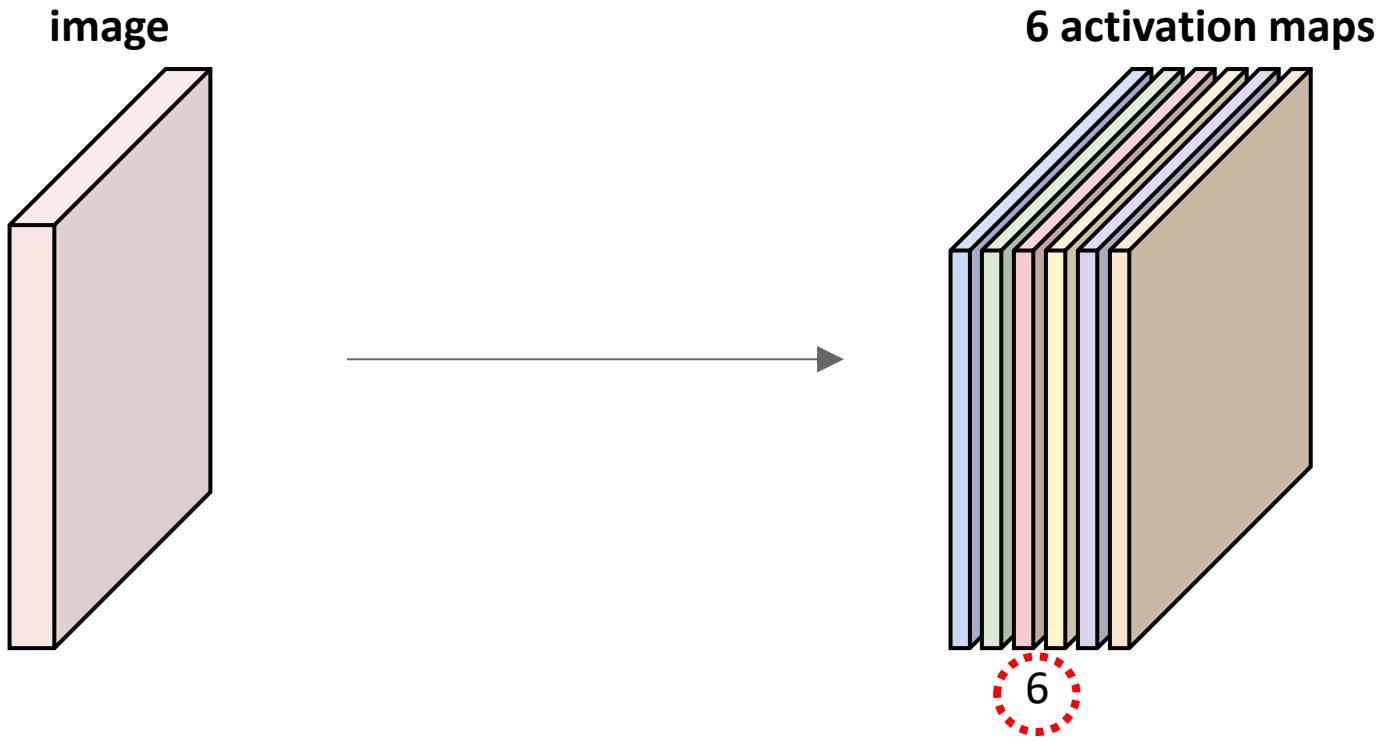




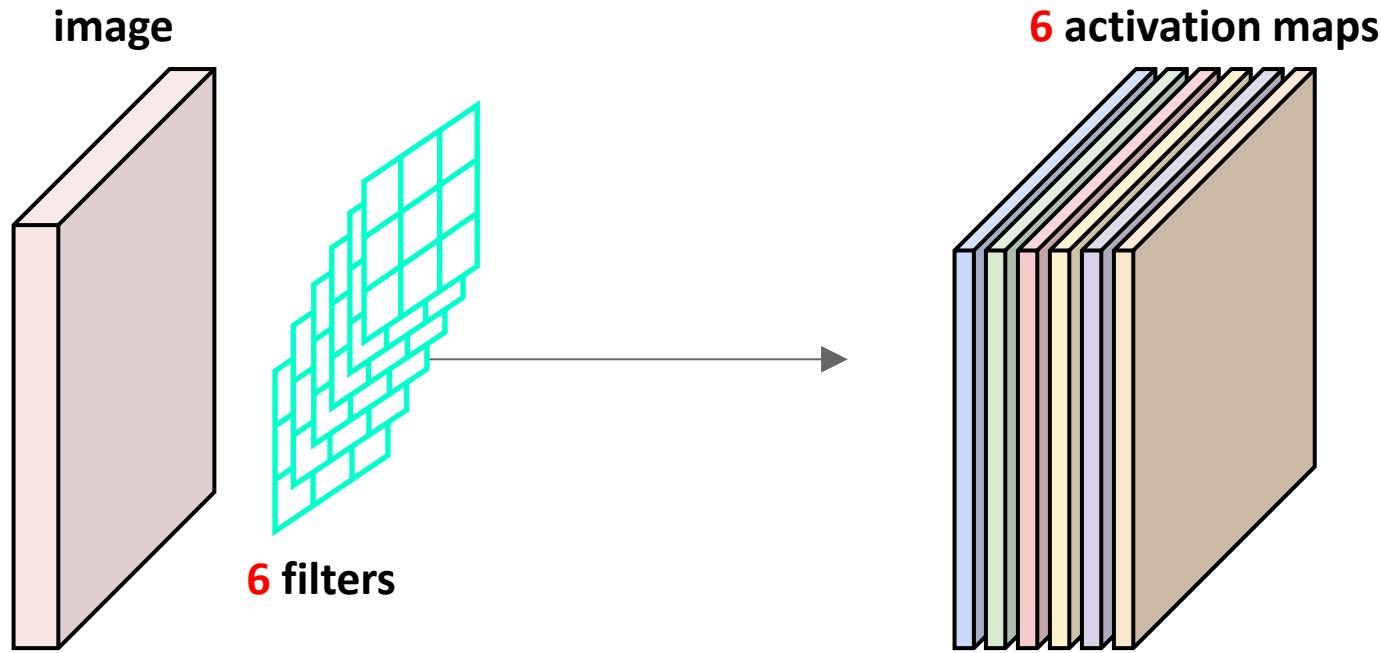








Where are the filter, and how many?



How many filters and where are they?

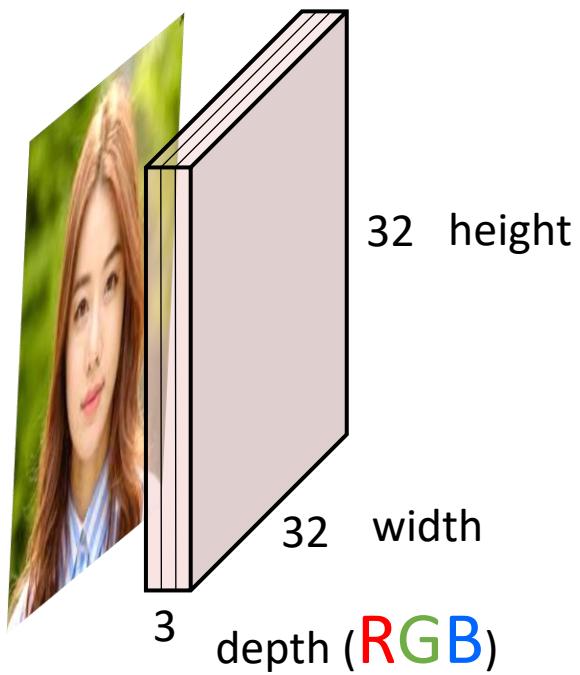
x1.0	x1.5	x1.0	x0.4
x0.7	x1.0	x0.7	x0.6
x-0.4	x-0.4	x-0.4	x0.0

0.73	0.86	0.74	0.85	0.99	0.58	0.96	0.25
0.29	0.92	0.25	0.38	0.67	0.16	0.63	0.73
0.90	0.99	0.95	0.06	0.56	0.85	0.23	0.84
0.00	0.45	0.55	0.28	0.33	0.59	0.69	0.01
0.50	0.23	0.03	0.41	0.49	0.23	0.62	0.69
0.85	0.52	0.11	0.95	0.80	0.24	0.03	0.54
0.99	0.07	0.69	0.35	0.05	0.15	0.61	0.80
0.03	0.84	0.79	0.52	0.17	0.66	0.73	0.01



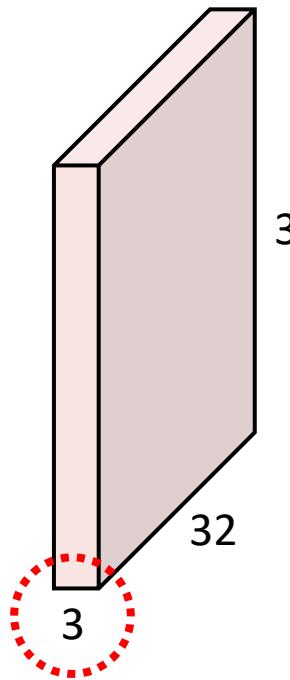
Convolution Layer

32x32x3 pixels image



Convolution Layer

32x32x3 image



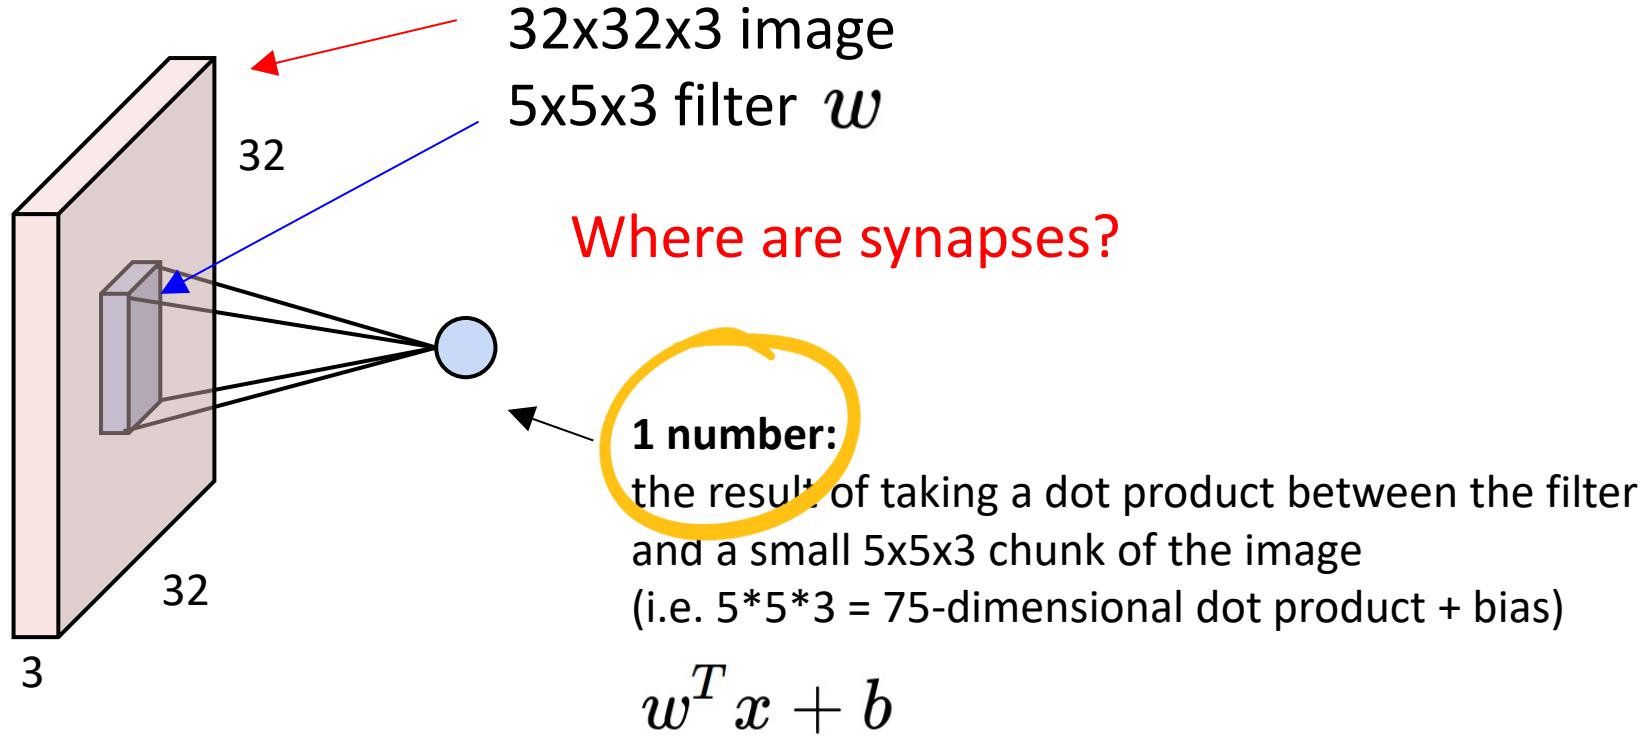
5x5x3 filter



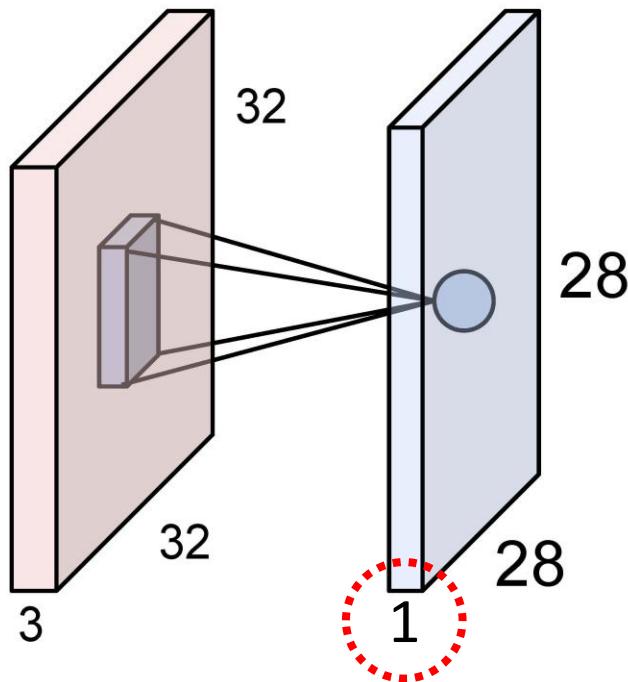
Filters always extend the full depth of the input volume

Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

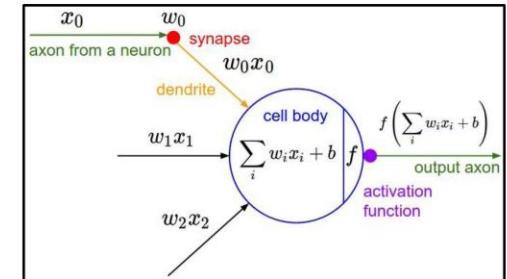
Convolution Layer



The brain/neuron view of CONV Layer



Where are synapses?

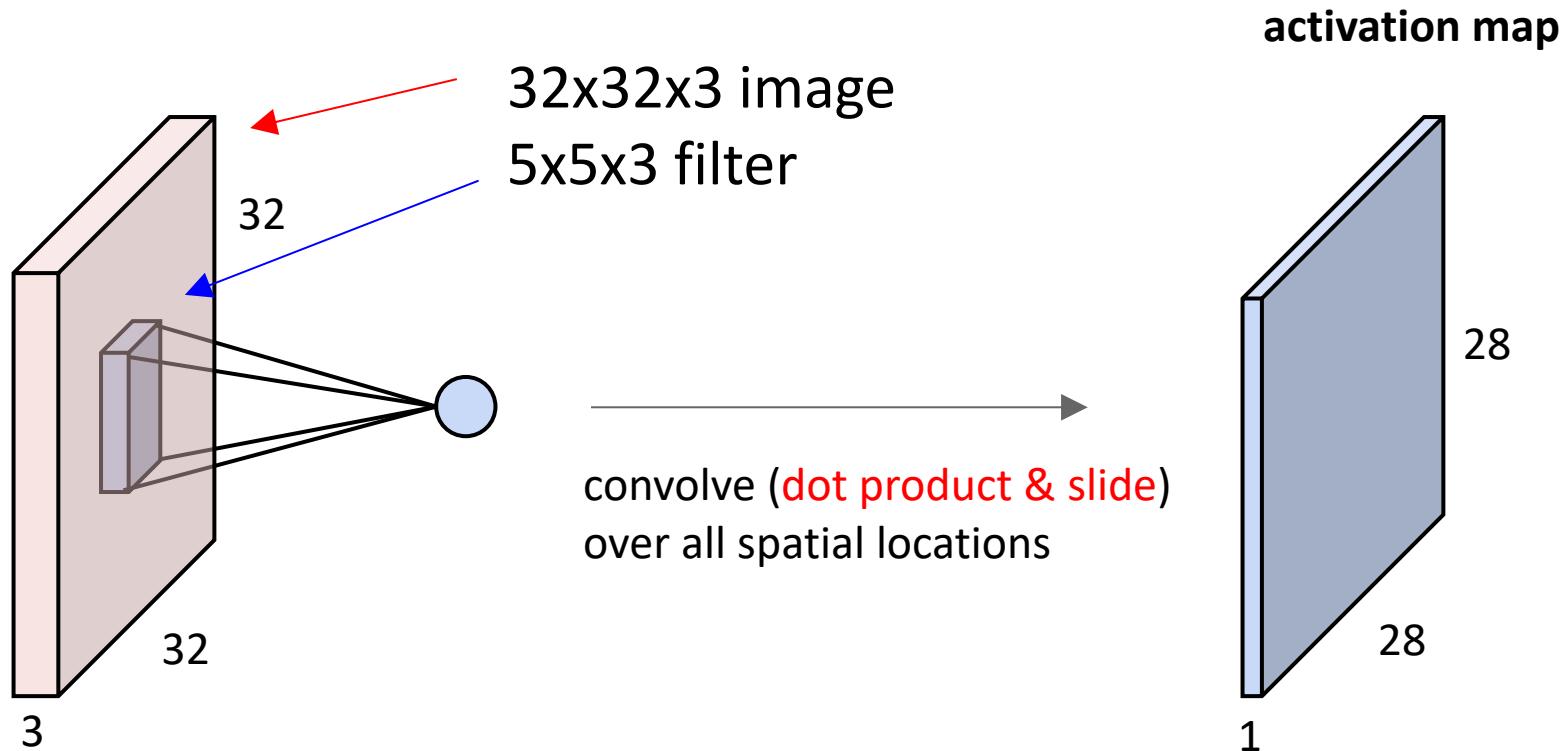


An activation map is a 28x28 sheet of neuron outputs:

1. Each is connected to a small region in the input
2. All of them share parameters

“5x5 filter” -> “5x5 receptive field for each neuron”

Convolution Layer

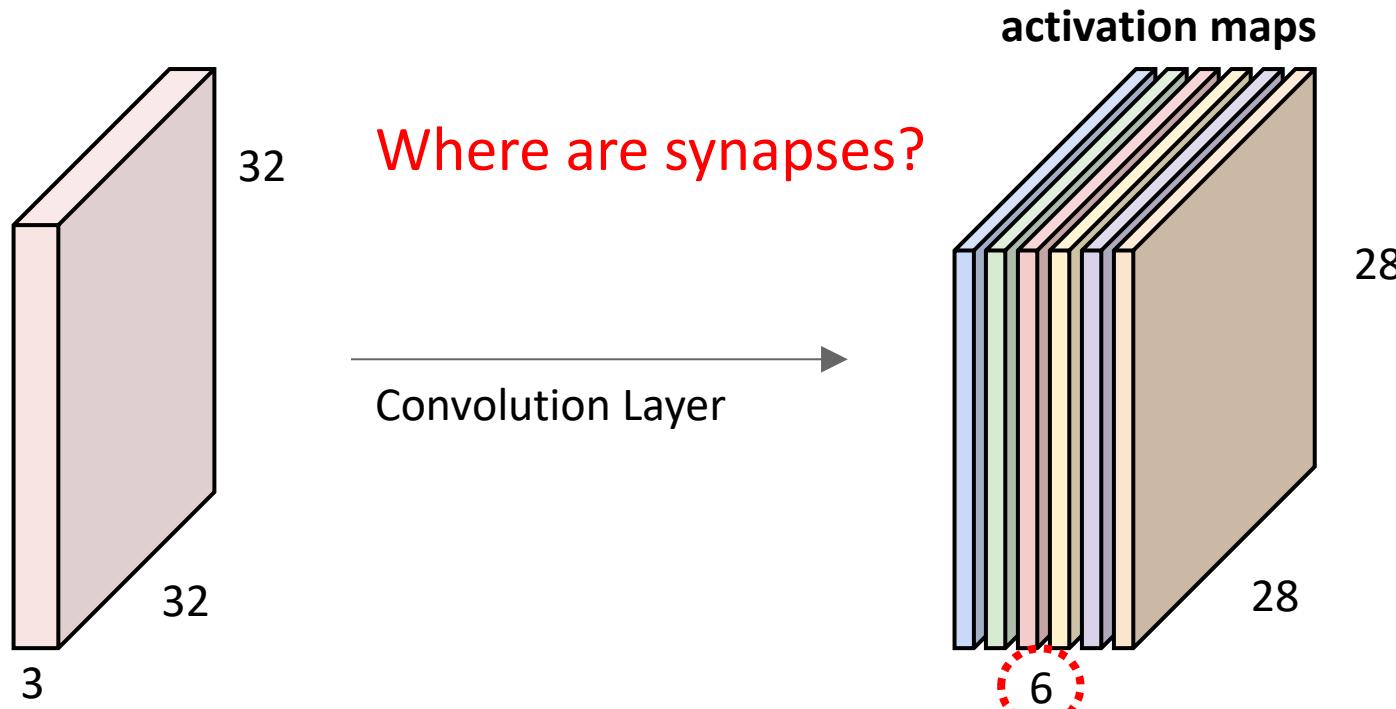


Convolution Layer

consider a second, green filter



For example, if we had **6** **5x5 filters**, we'll get 6 separate activation maps:

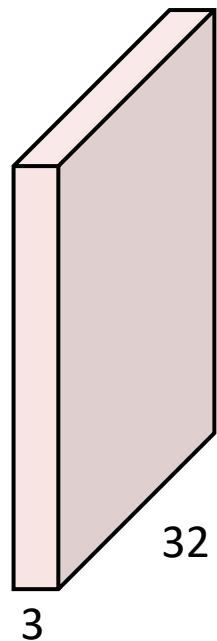


We processed [32x32x3] volume into [28x28x6] volume.

Q: how many parameters would this be if we used a fully connected layer instead?

A: $(32*32*3)*(28*28*6) = 14.5M$ parameters, ~14.5M multiplies

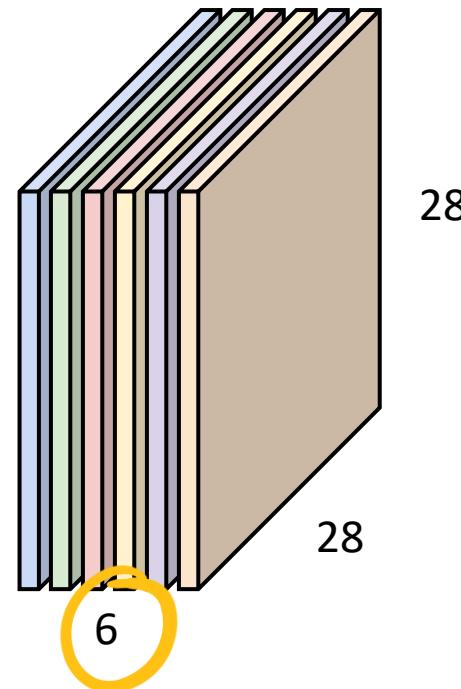
For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



Convolution Layer

1 Filter -> 1 Activation Map
6 Filters -> 6 Activation Maps

activation maps

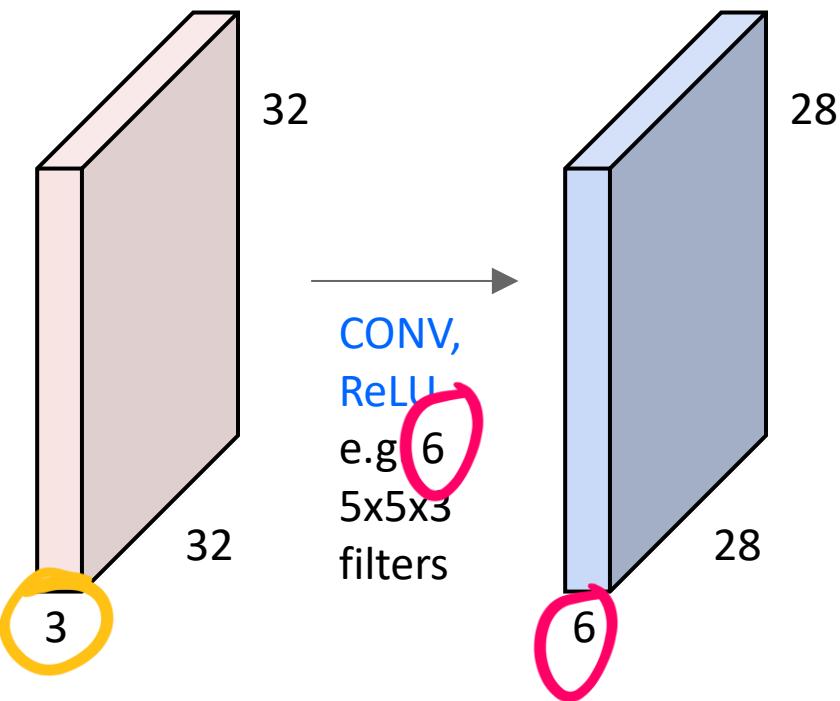


We processed [32x32x3] volume into [28x28x6] volume.

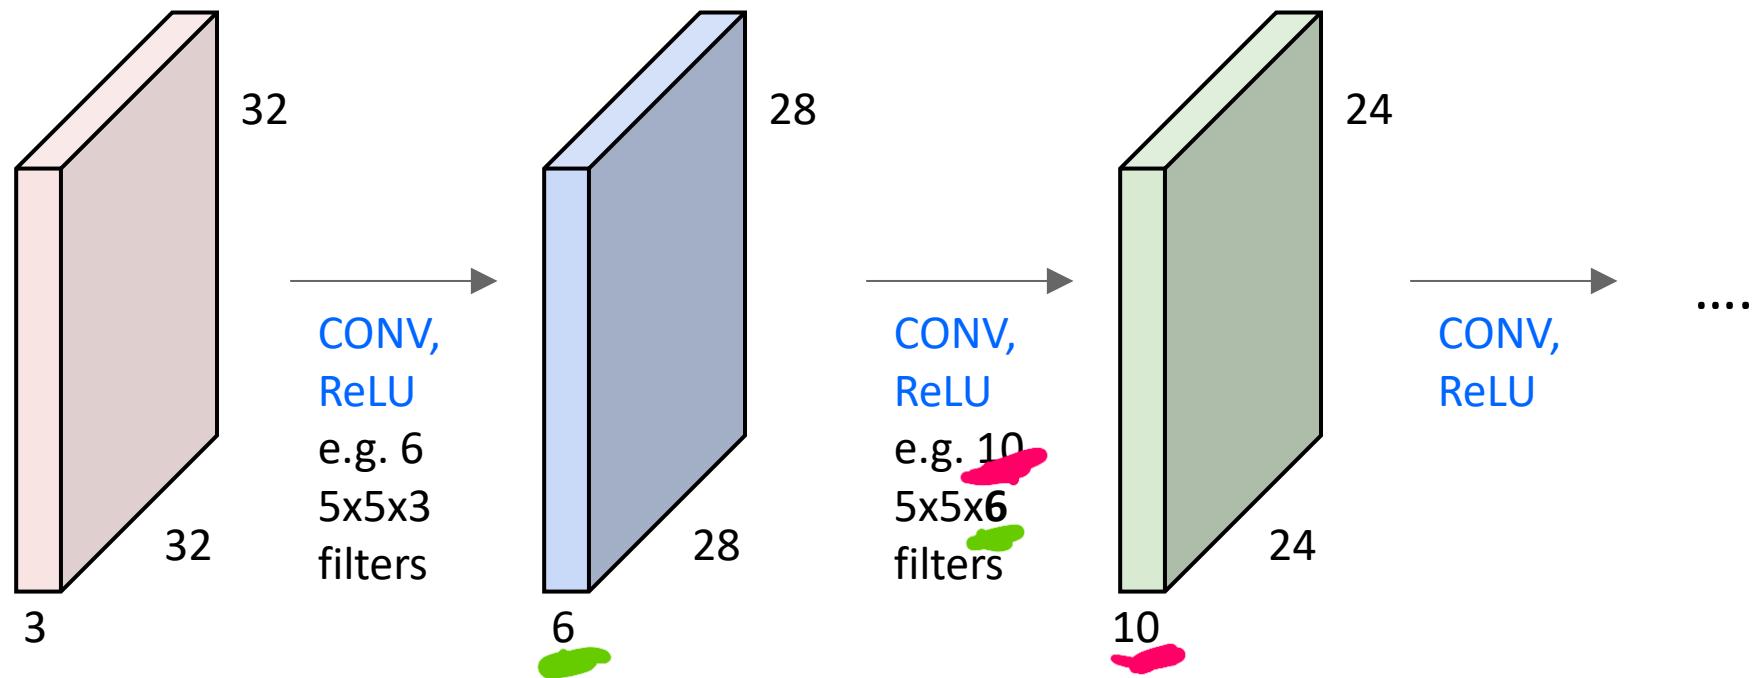
Q: how many parameters are used instead?

A: $(5*5*3)*6 = 450 \text{ parameters}$, $(5*5*3)*(28*28*6) = \sim 350K \text{ multiplies}$

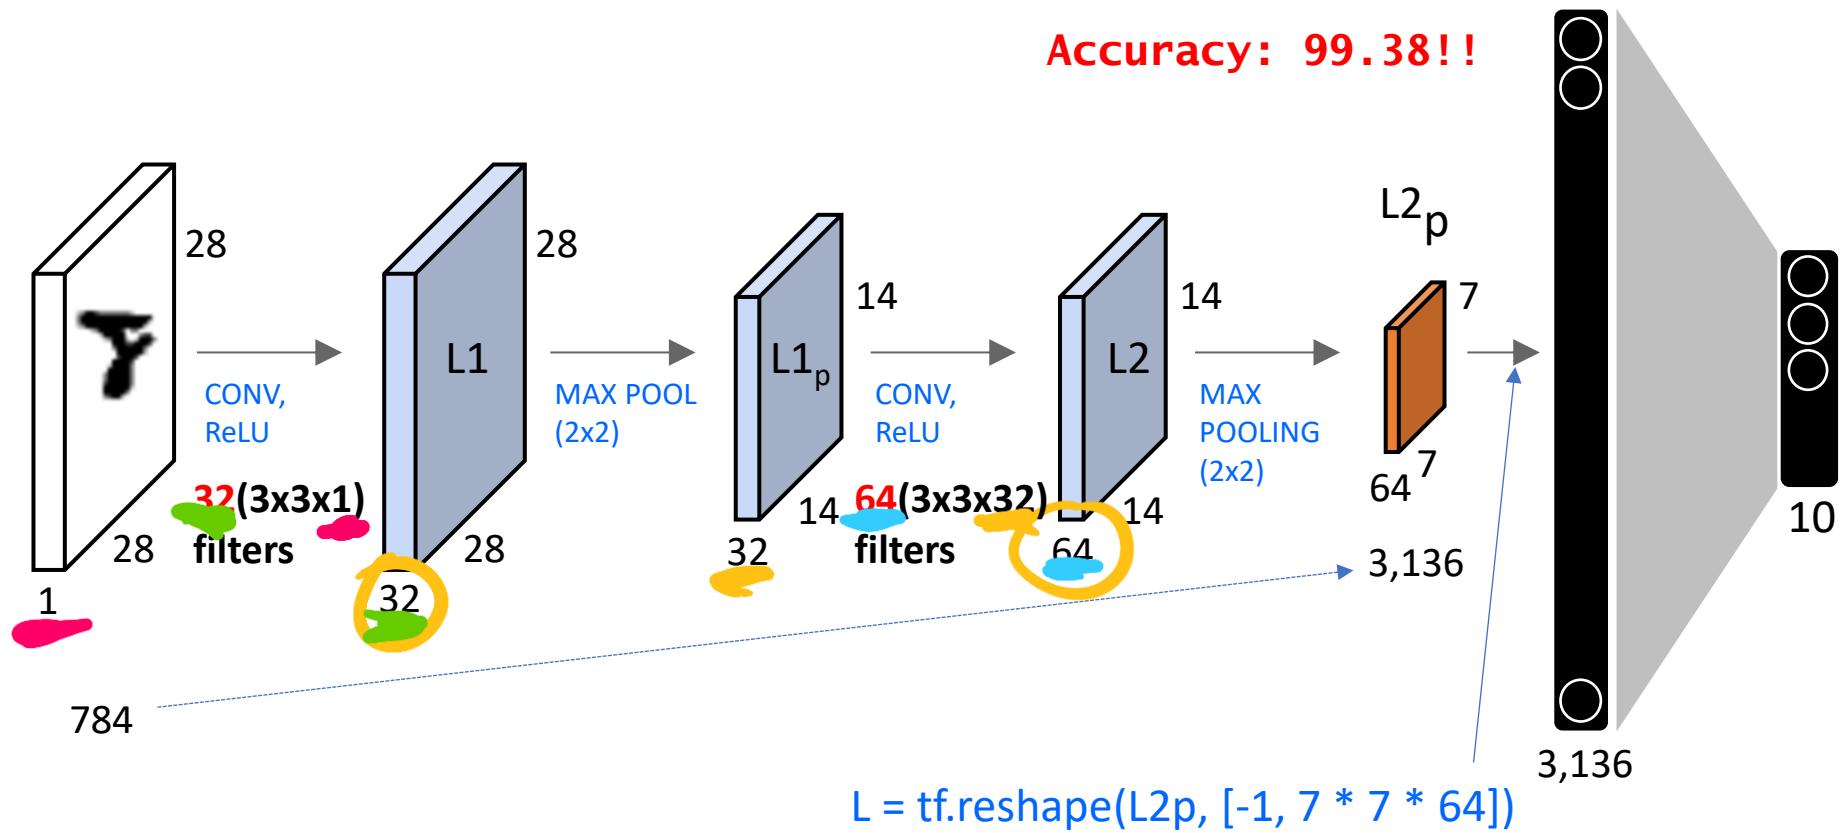
Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions

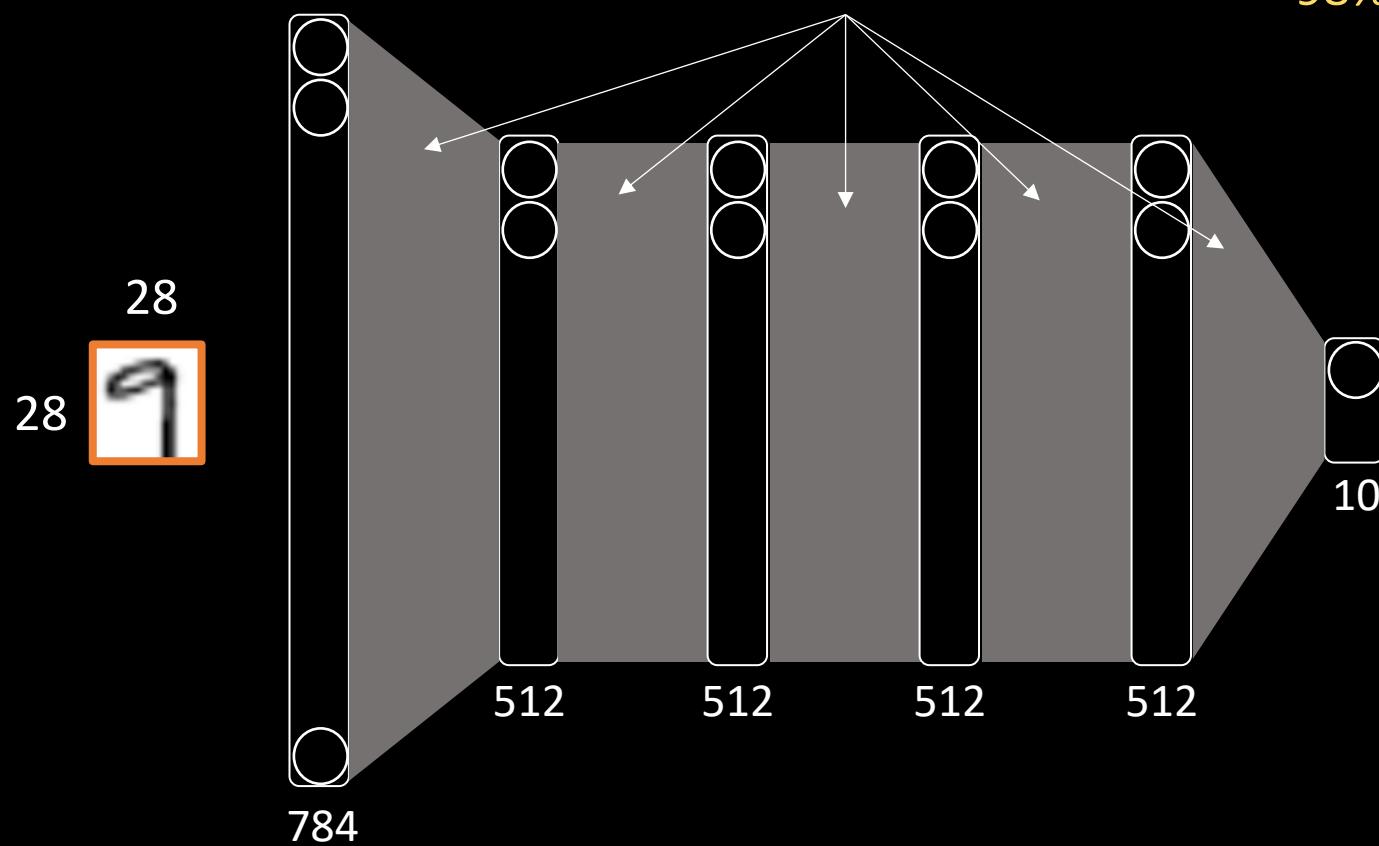


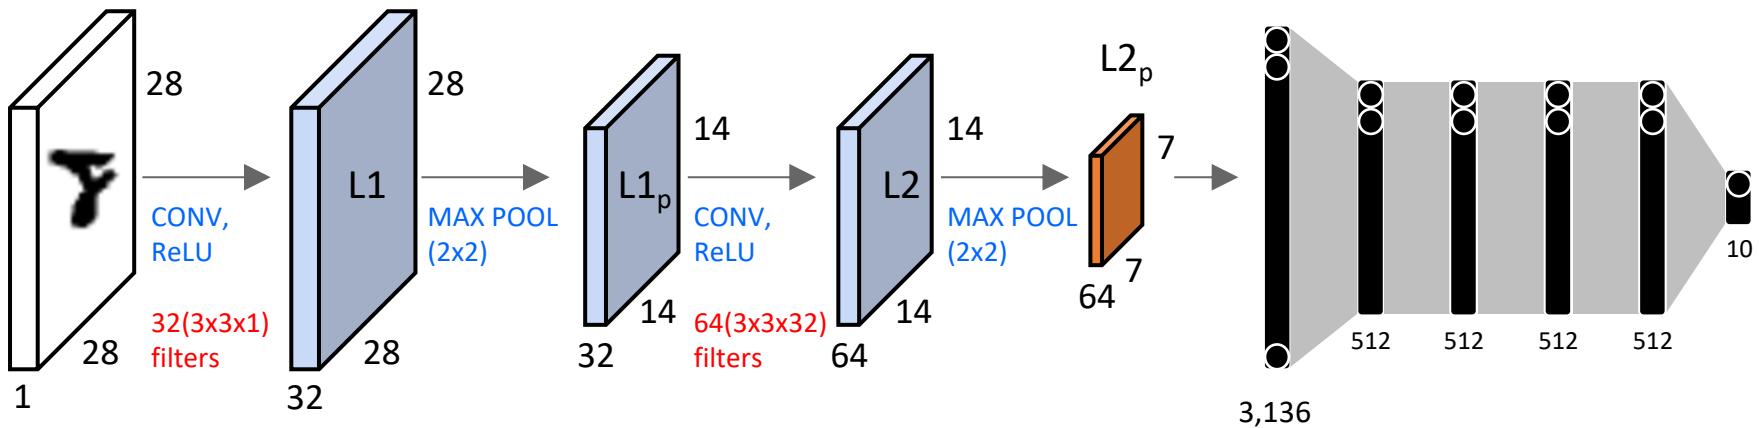
Accuracy: 99.38!!



Where are the **final features** we try to extract?

23_mnist_nn_deep.py
98%

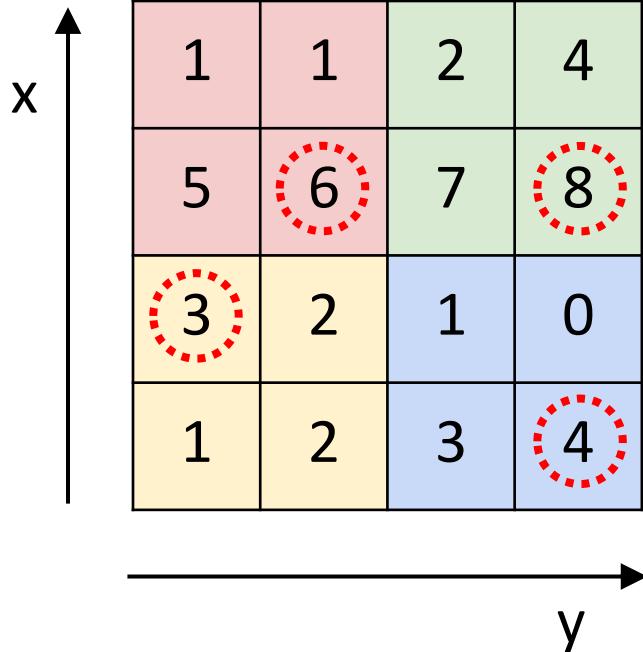




Where are the features we try to extract?

MAX POOLING

Single depth slice



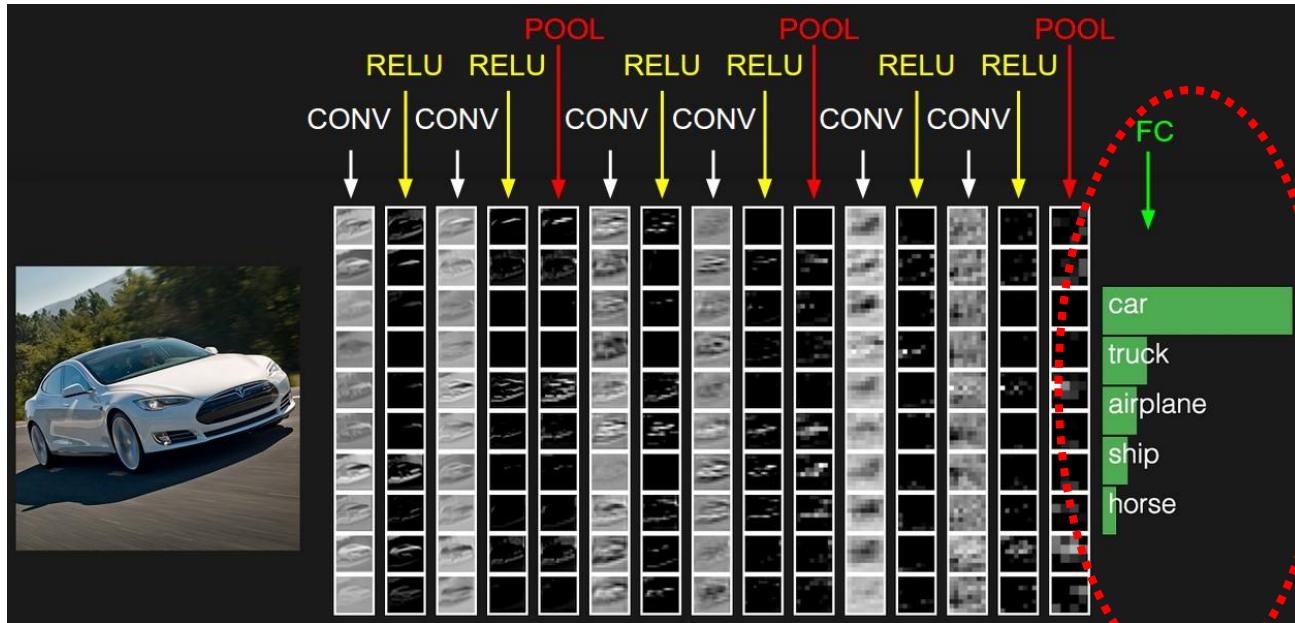
2x2 max pooling
with stride 2

“dimension reduction”

6	8
3	4

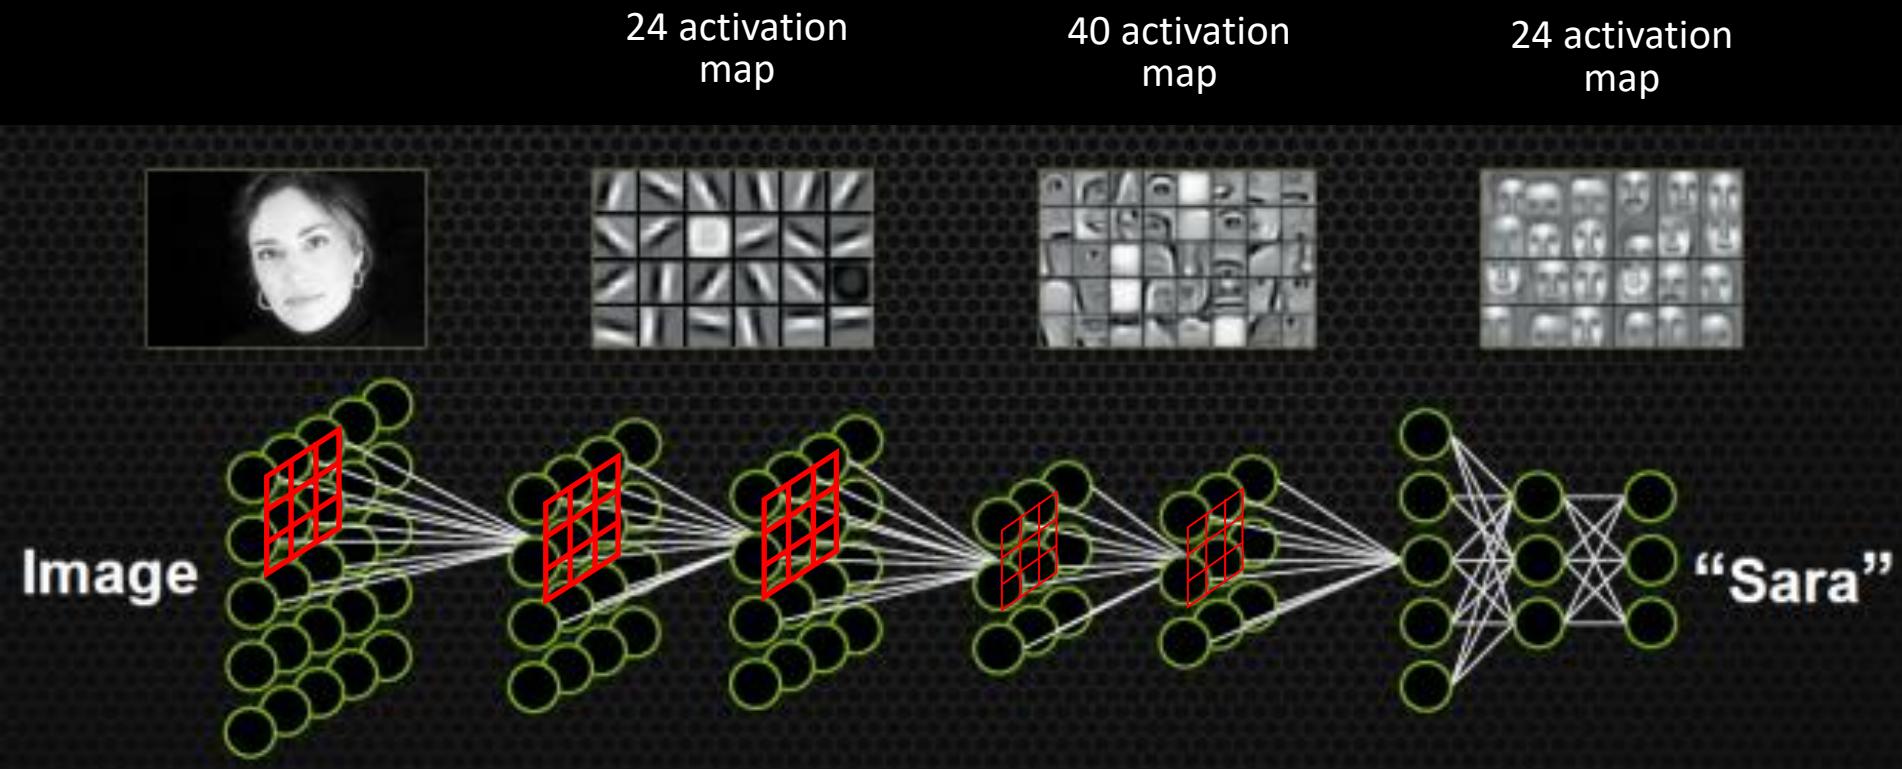
Fully Connected Layer (FC layer)

- Contains neurons that connect to the entire input volume, as in ordinary Neural Networks



10
activation
maps

HOW A DEEP NEURAL NETWORK SEES



HOW A DEEP NEURAL NETWORK SEES

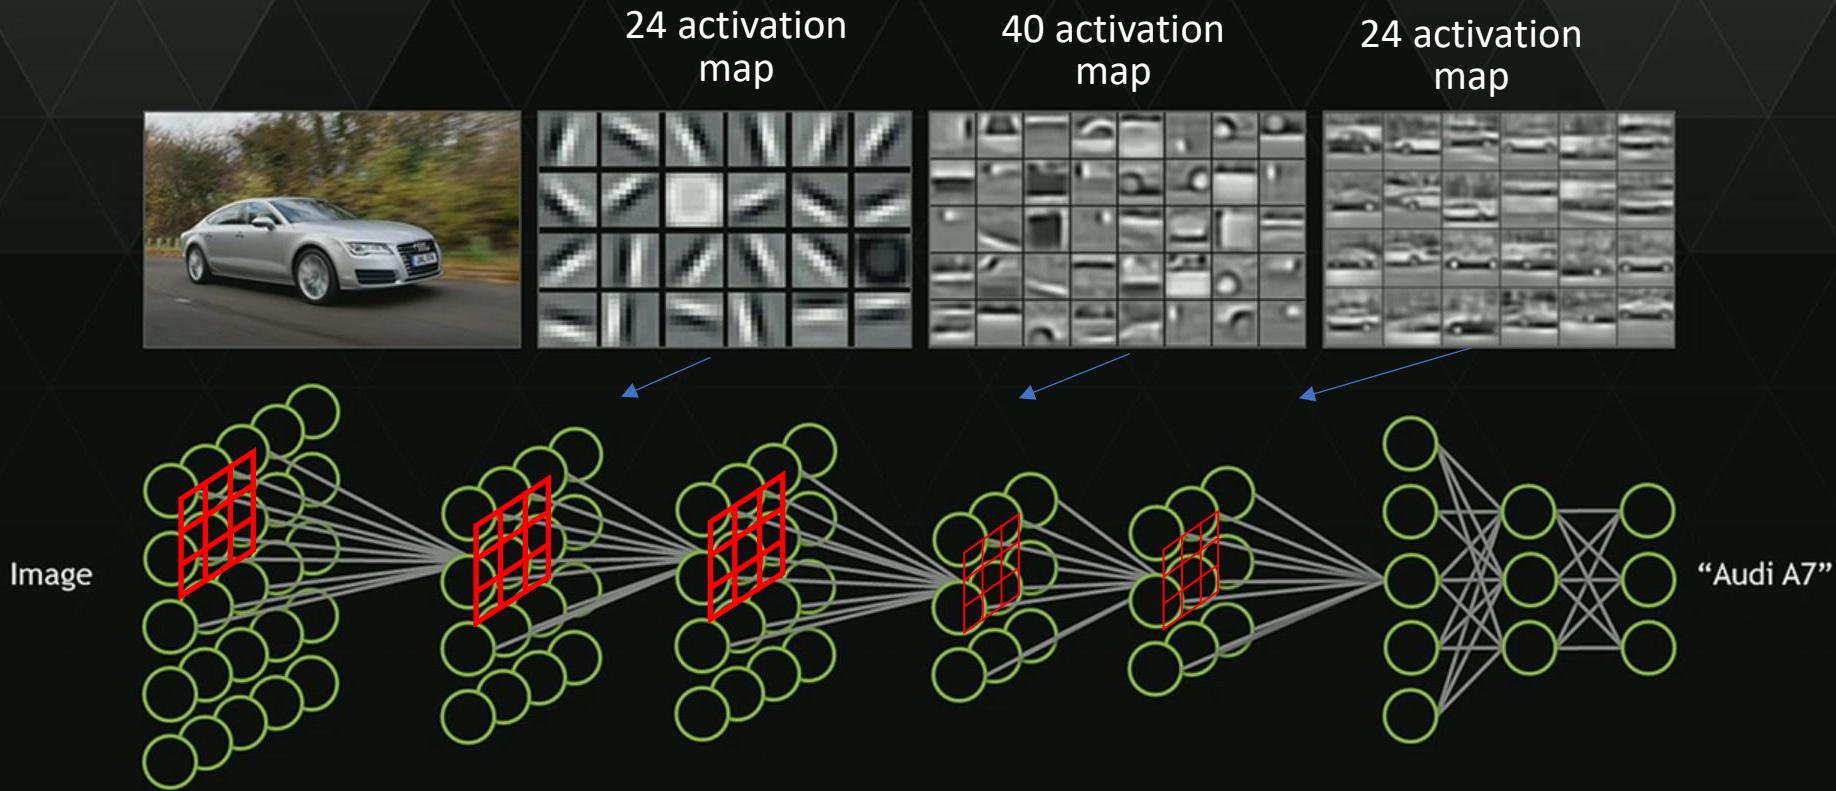
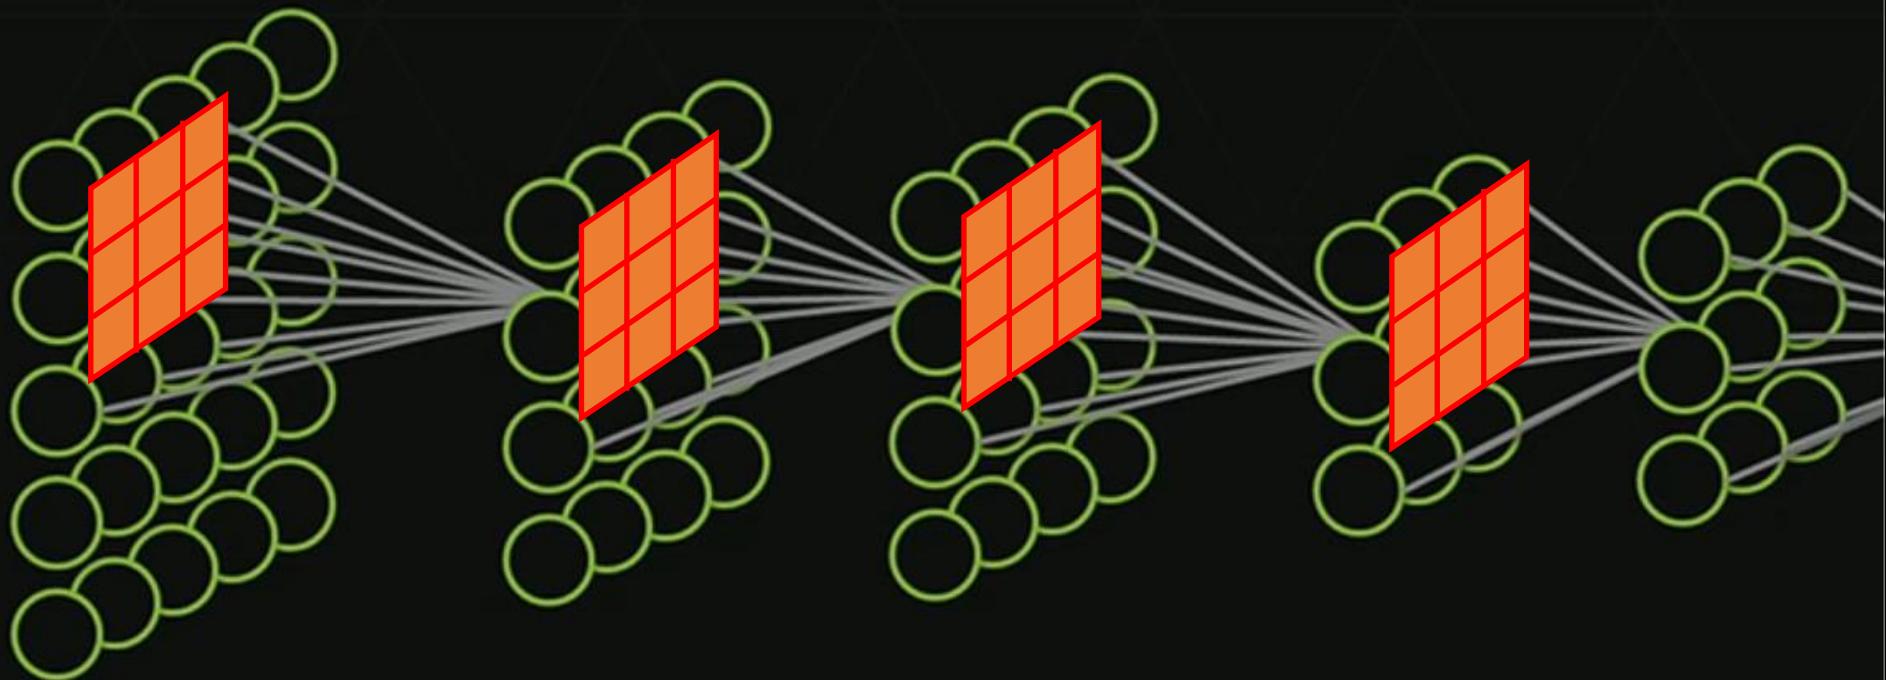


Image source: "Unsupervised Learning of Hierarchical Representations with Convolutional Deep Belief Networks" ICML 2009 & Comm. ACM 2011.
Honglak Lee, Roger Grosse, Rajesh Ranganath, and Andrew Ng.

Filter size : 3x3, 1 filter for each convolution,
then how many parameters to be tuned?

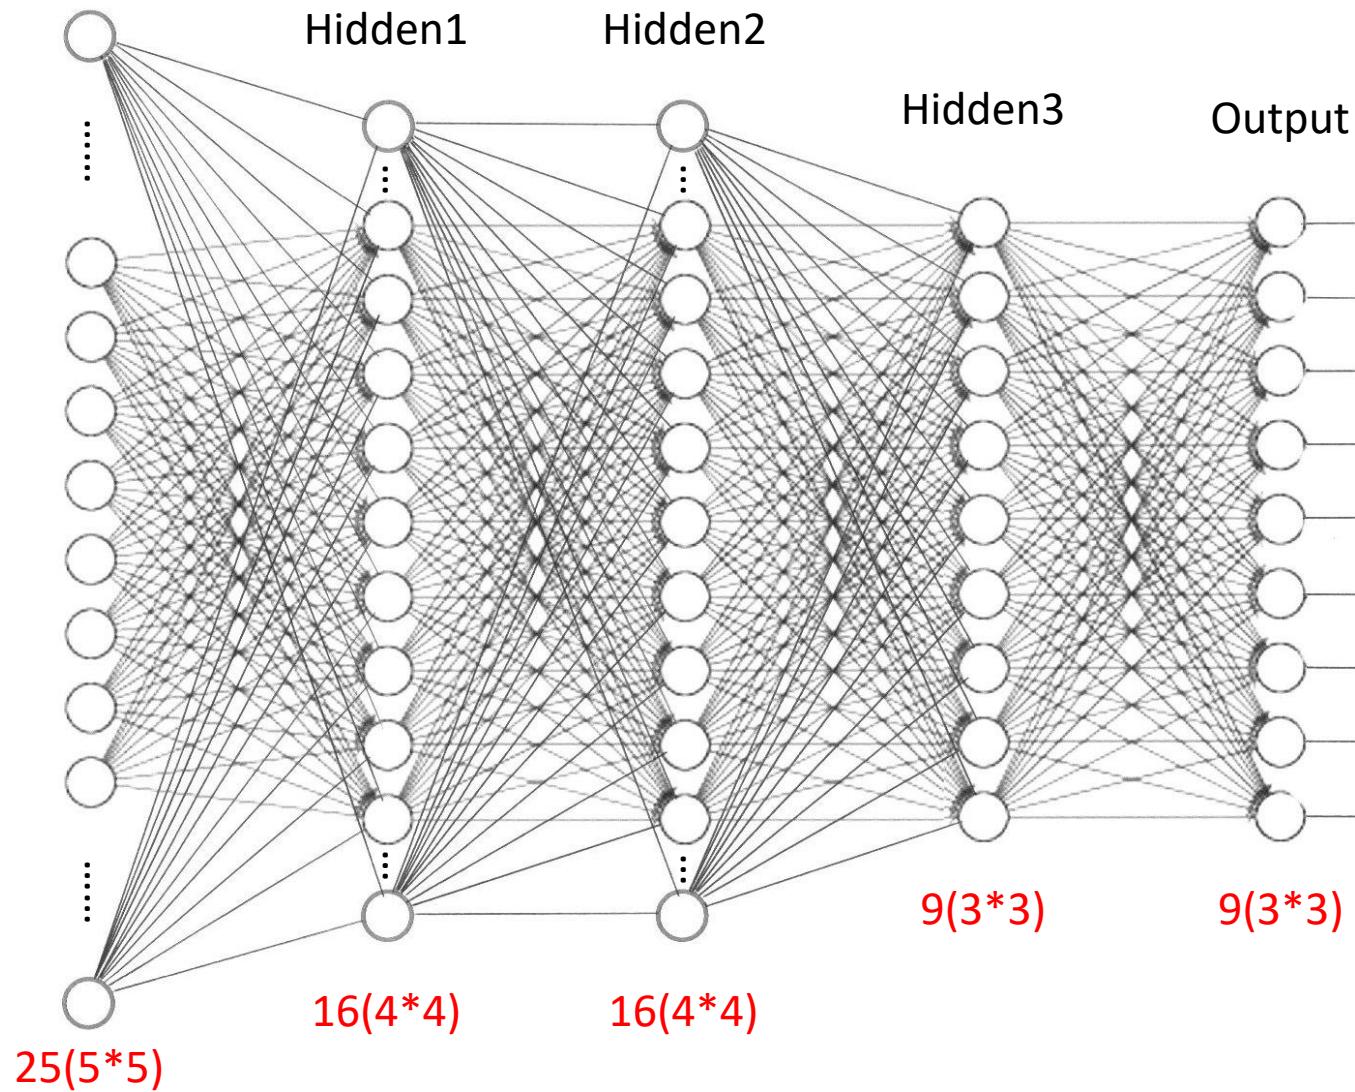
$$9 * 4 = 36$$



If fully connected, $25 * 16 + 16 * 16 + 16 * 9 + 9 * 9 = 881$

$$\frac{881}{36} = 24$$

Input



Some questions about CNN

- Fully connected? (Yes/No)
- Where are the parameters to be tuned?
- We have 2 filters, then?
- Convolution again with the resulting activation map

Learning in CNN

- Parameter tuning in filters
→ filter organizing
- Parameter tuning in fully connected layers

Theoretical backgrounds of CNN

A bit of history:

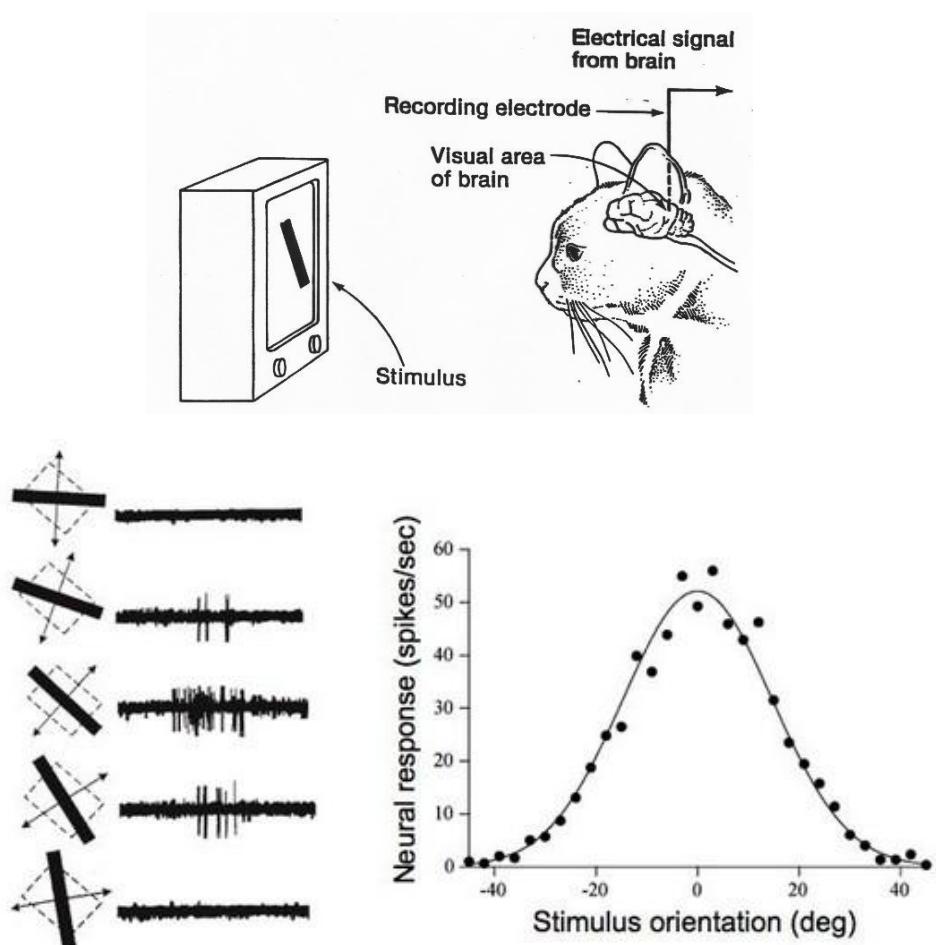
Hubel & Wiesel, 1959

RECEPTIVE FIELDS OF SINGLE
NEURONES IN
THE CAT'S STRIATE CORTEX

1962

RECEPTIVE FIELDS, BINOCULAR
INTERACTION
AND FUNCTIONAL ARCHITECTURE
IN
THE CAT'S VISUAL CORTEX

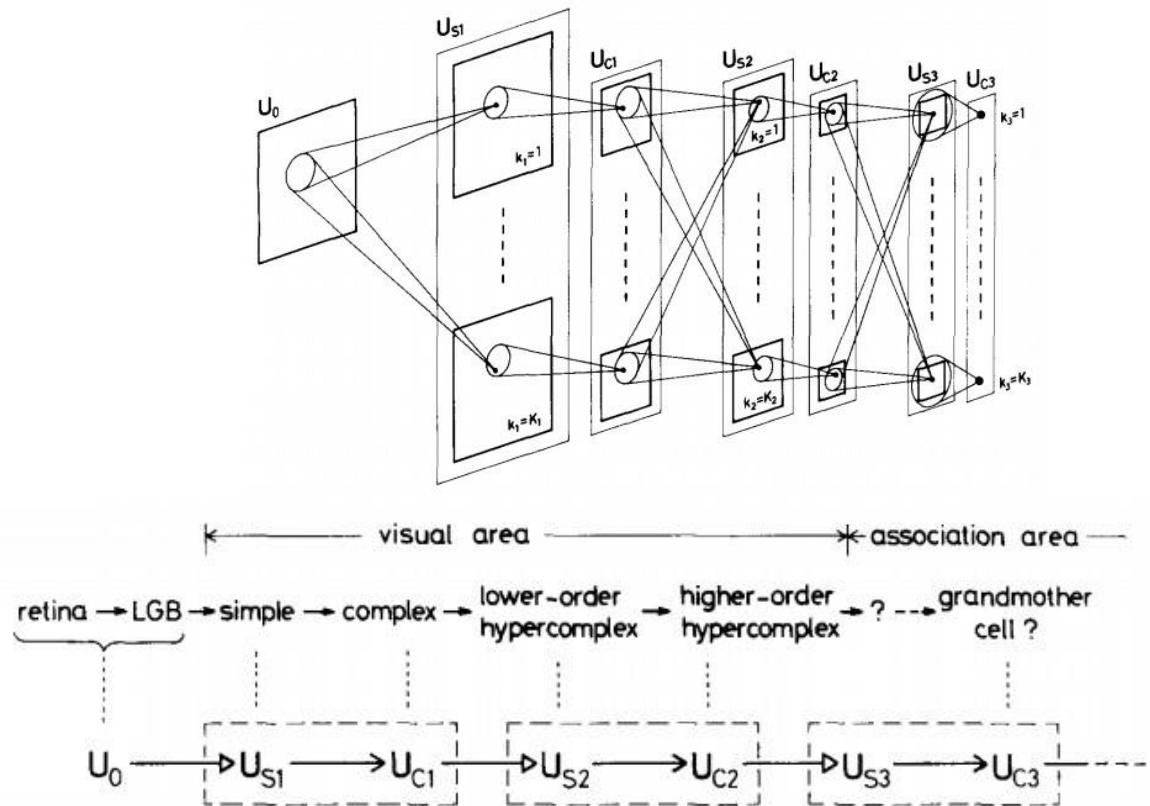
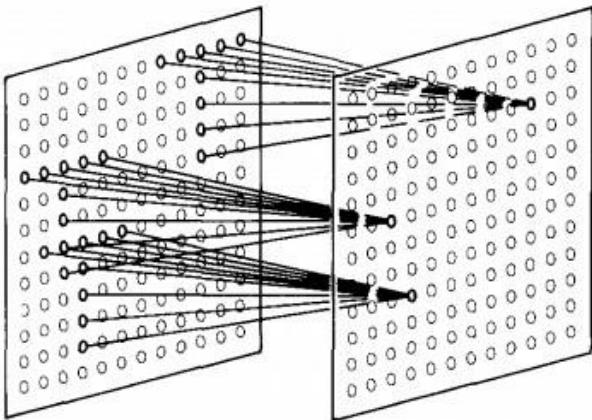
1968...



A bit of history:

“sandwich” architecture (SCSCSC...)
simple cells: modifiable parameters
complex cells: perform pooling

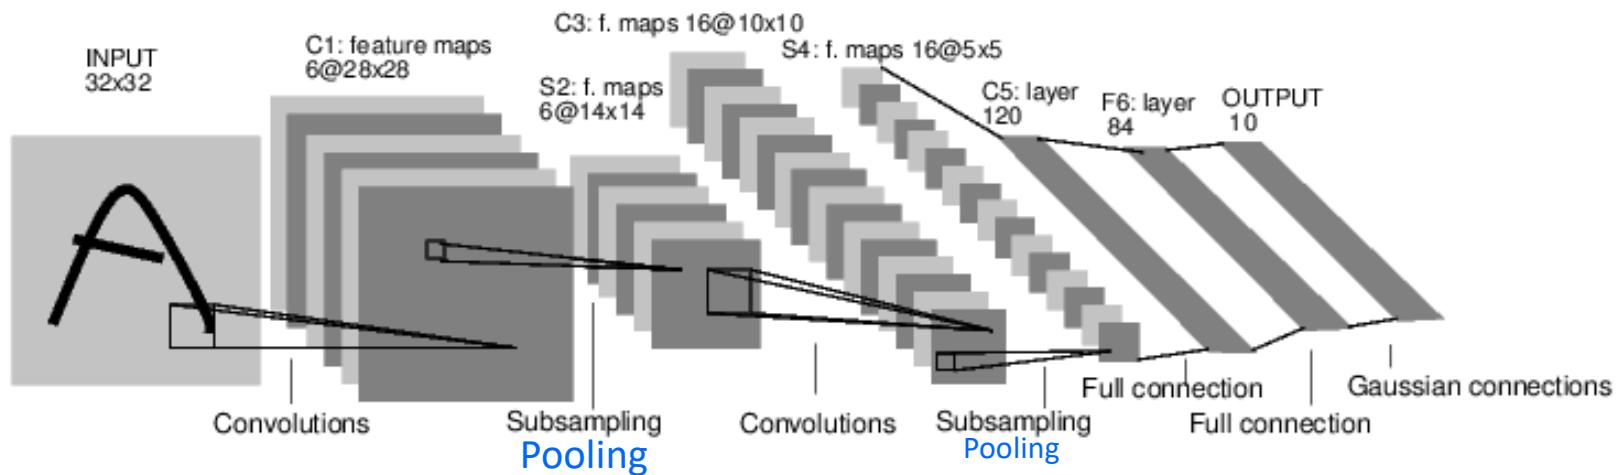
Neuro-cognitron [Fukushima 1980]



LeNet-5

Gradient-based learning applied to document recognition

[LeCun, Bottou, Bengio, Haffner 1998]



Convolution-Pooling-Convolution-Pooling-FC

Where are the features we try to extract?

Case Studies

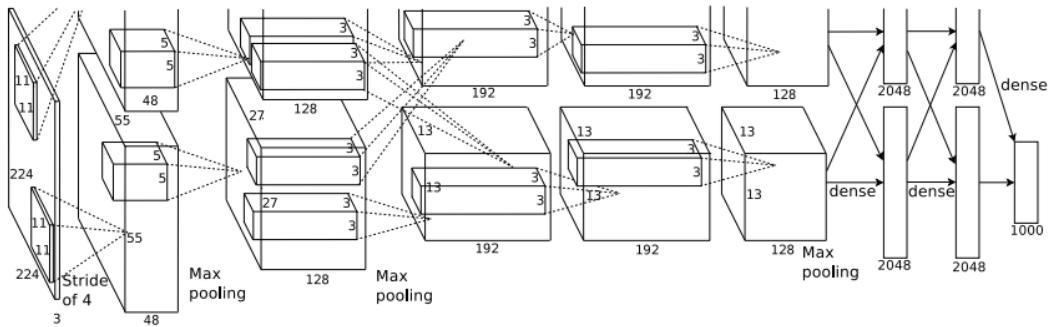


ImageNet and Competition

- **ImageNet**: categorization of huge amount of images according to WordNet schema
- database for ILSVRC (ImageNet Large Scale Visual Recognition Competition)
- **Largest Computer Vision Challenge** showing state-of-the-art performance on the dataset every year
- firstly collected by Deng et al. in 2009
- And it has been organized by Stanford University Vision Lab (prof. Fei-Fei Li) since 2010.

Case Study: AlexNet

[Krizhevsky et al. 2012]



Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] CONV1: 96 11x11 filters at stride 4, pad 0

[27x27x96] MAX POOL1: 3x3 filters at stride 2

[27x27x96] NORM1: Normalization layer

[27x27x256] CONV2: 256 5x5 filters at stride 1, pad 2

[13x13x256] MAX POOL2: 3x3 filters at stride 2

[13x13x256] NORM2: Normalization layer

[13x13x384] CONV3: 384 3x3 filters at stride 1, pad 1

[13x13x384] CONV4: 384 3x3 filters at stride 1, pad 1

[13x13x256] CONV5: 256 3x3 filters at stride 1, pad 1

[6x6x256] MAX POOL3: 3x3 filters at stride 2

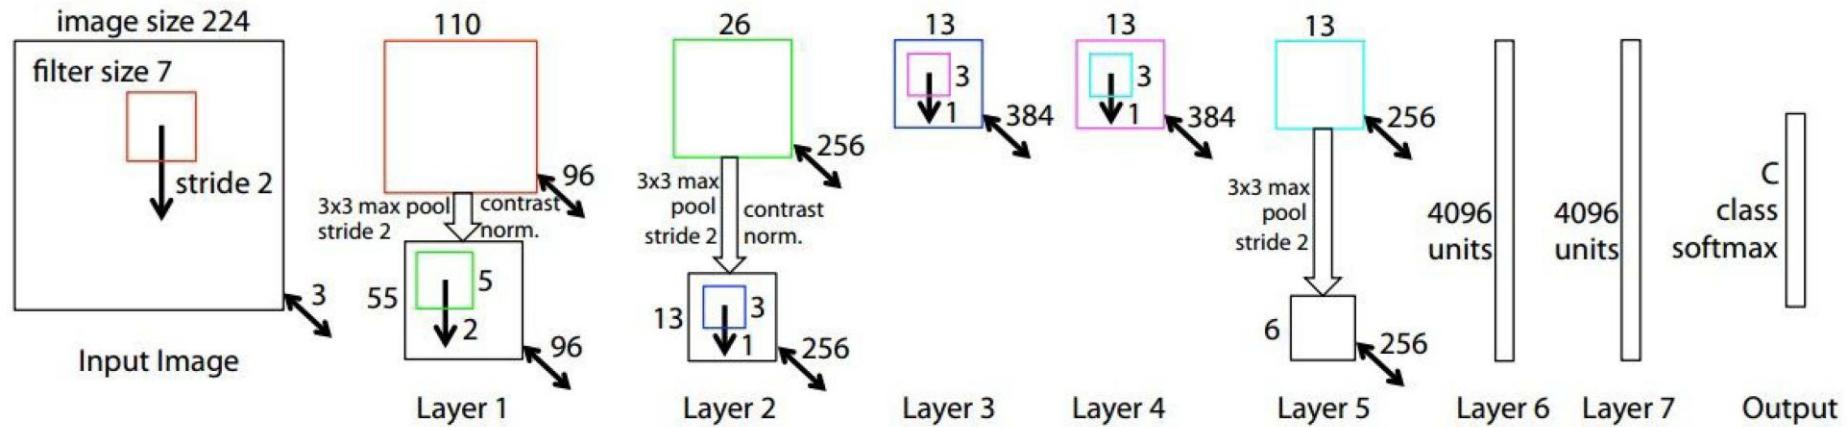
[4096] FC6: 4096 neurons

[4096] FC7: 4096 neurons

[1000] FC8: 1000 neurons (class scores)

Case Study: ZFNet

[Zeiler and Fergus, 2013]



AlexNet but:

CONV1: change from (11x11 stride 4) to (7x7 stride 2)

CONV3,4,5: instead of 384, 384, 256 filters use 512, 1024, 512

ImageNet top 5 error: 15.4% -> 14.8%

Case Study: VGGNet

[Simonyan and Zisserman, 2014]

Only 3x3 CONV stride 1, pad 1
and 2x2 MAX POOL stride 2

best model

11.2% top 5 error in ILSVRC 2013

->

7.3% top 5 error

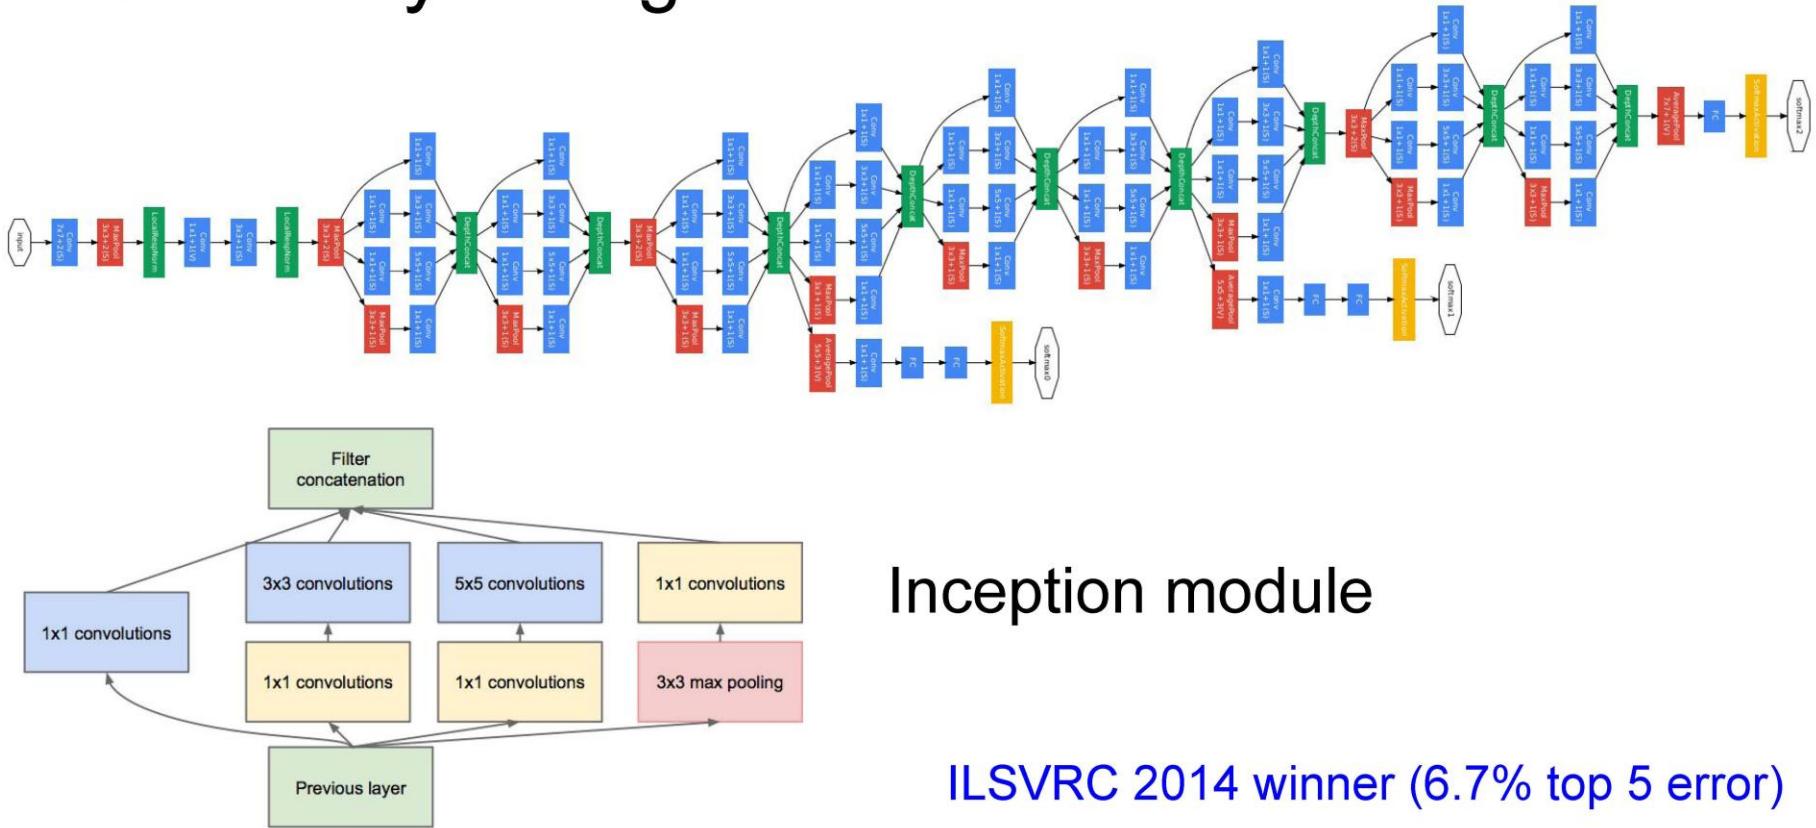
ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: Number of parameters (in millions).

Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

Case Study: GoogLeNet

[Szegedy et al., 2014]



Case Study: ResNet

[He et al., 2015]

ILSVRC 2015 winner (3.6% top 5 error)

Microsoft
Research

MSRA @ ILSVRC & COCO 2015 Competitions

- **1st places in all five main tracks**
 - ImageNet Classification: “Ultra-deep” (quote Yann) **152-layer nets**
 - ImageNet Detection: **16%** better than 2nd
 - ImageNet Localization: **27%** better than 2nd
 - COCO Detection: **11%** better than 2nd
 - COCO Segmentation: **12%** better than 2nd

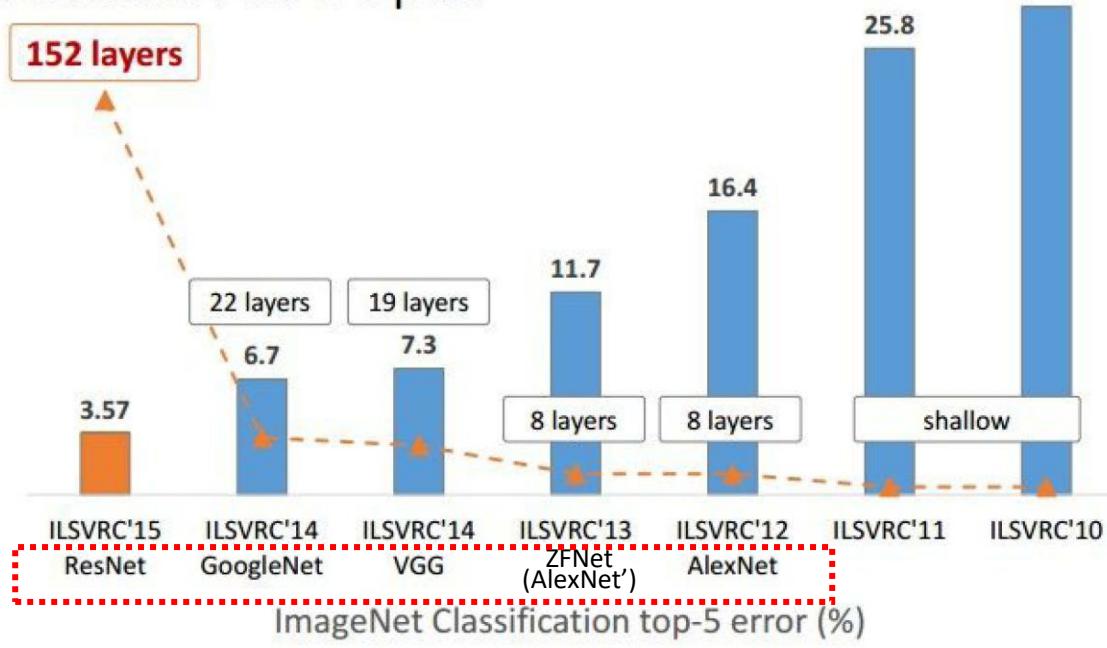
*improvements are relative numbers

ICCV15
International Conference on Computer Vision

Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.

Slide from Kaiming He's recent presentation <https://www.youtube.com/watch?v=1PGLj-uKT1w>

Revolution of Depth



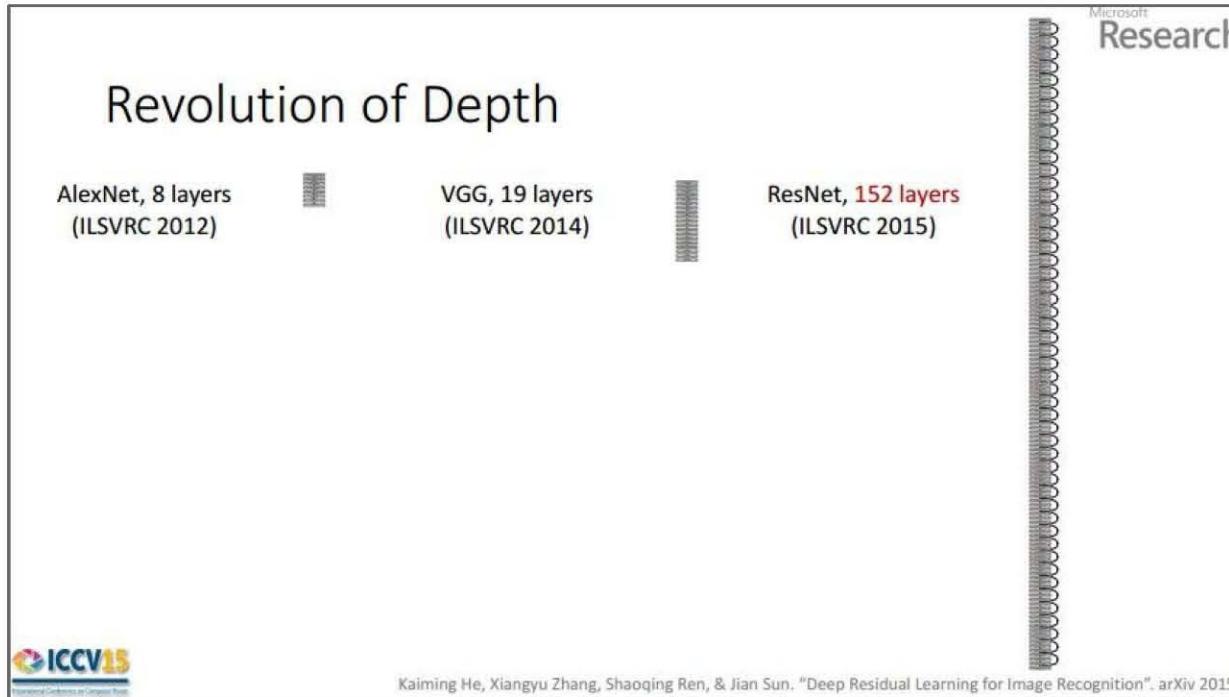
Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". arXiv 2015.

(slide from Kaiming He's recent presentation)

Case Study: ResNet

[He et al., 2015]

ILSVRC 2015 winner (3.6% top 5 error)

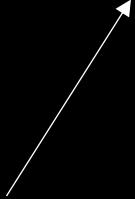


2-3 weeks of training
on 8 GPU machine

at runtime: faster
than a VGGNet!
(even though it has
8x more layers)

(slide from Kaiming He's recent presentation)

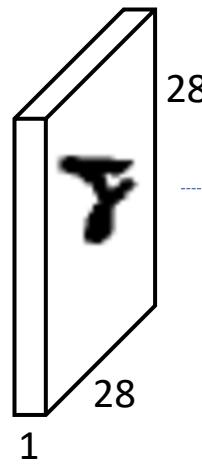
Dot product & shift for feature extraction



CNN, Convolutional Neural Network

20_mnist_softmax.py

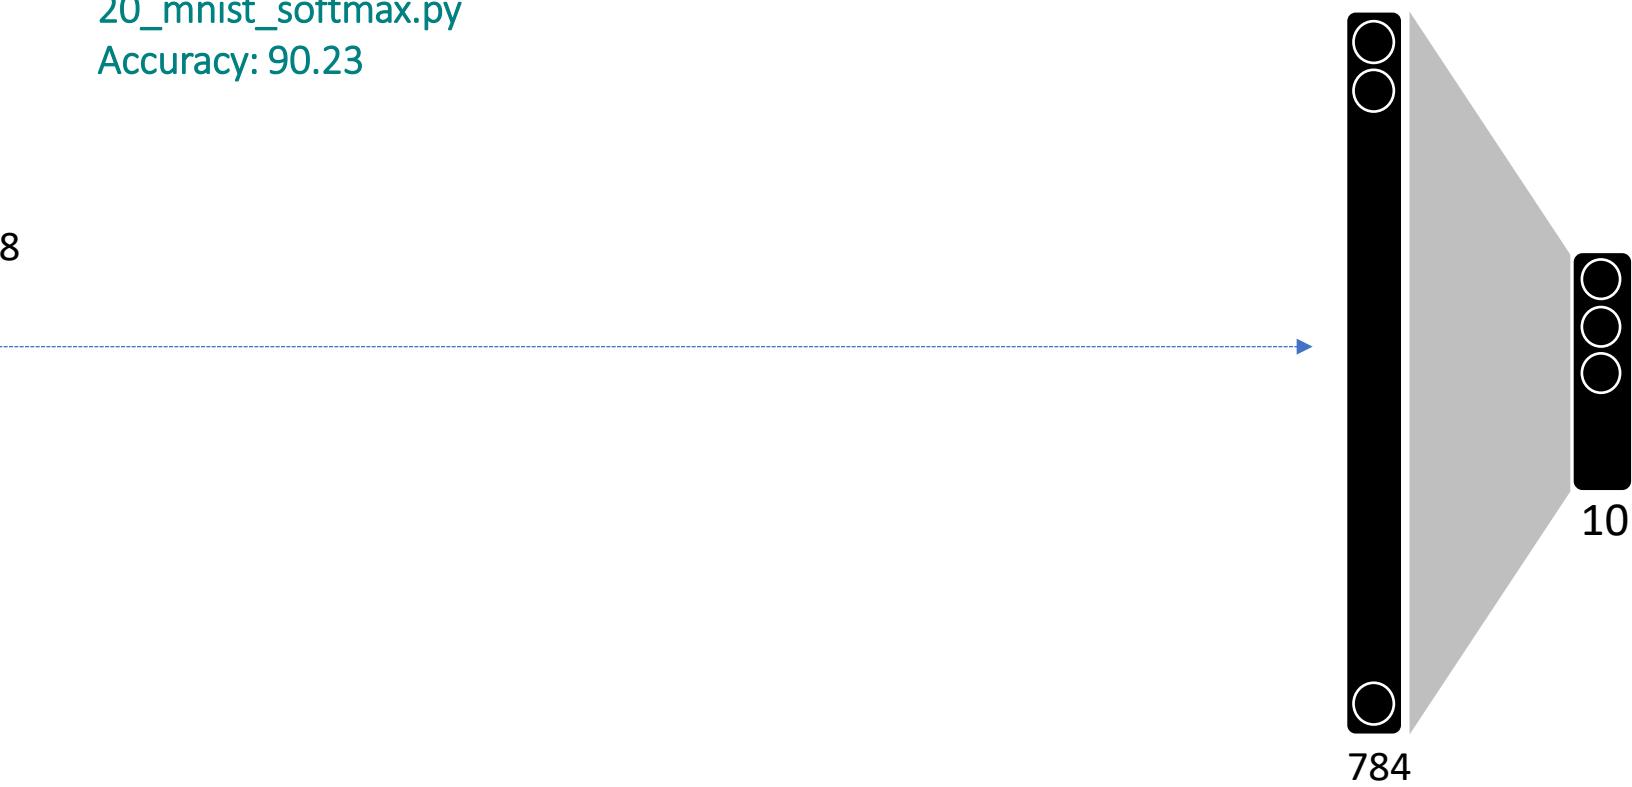
Accuracy: 90.23



28

28

1

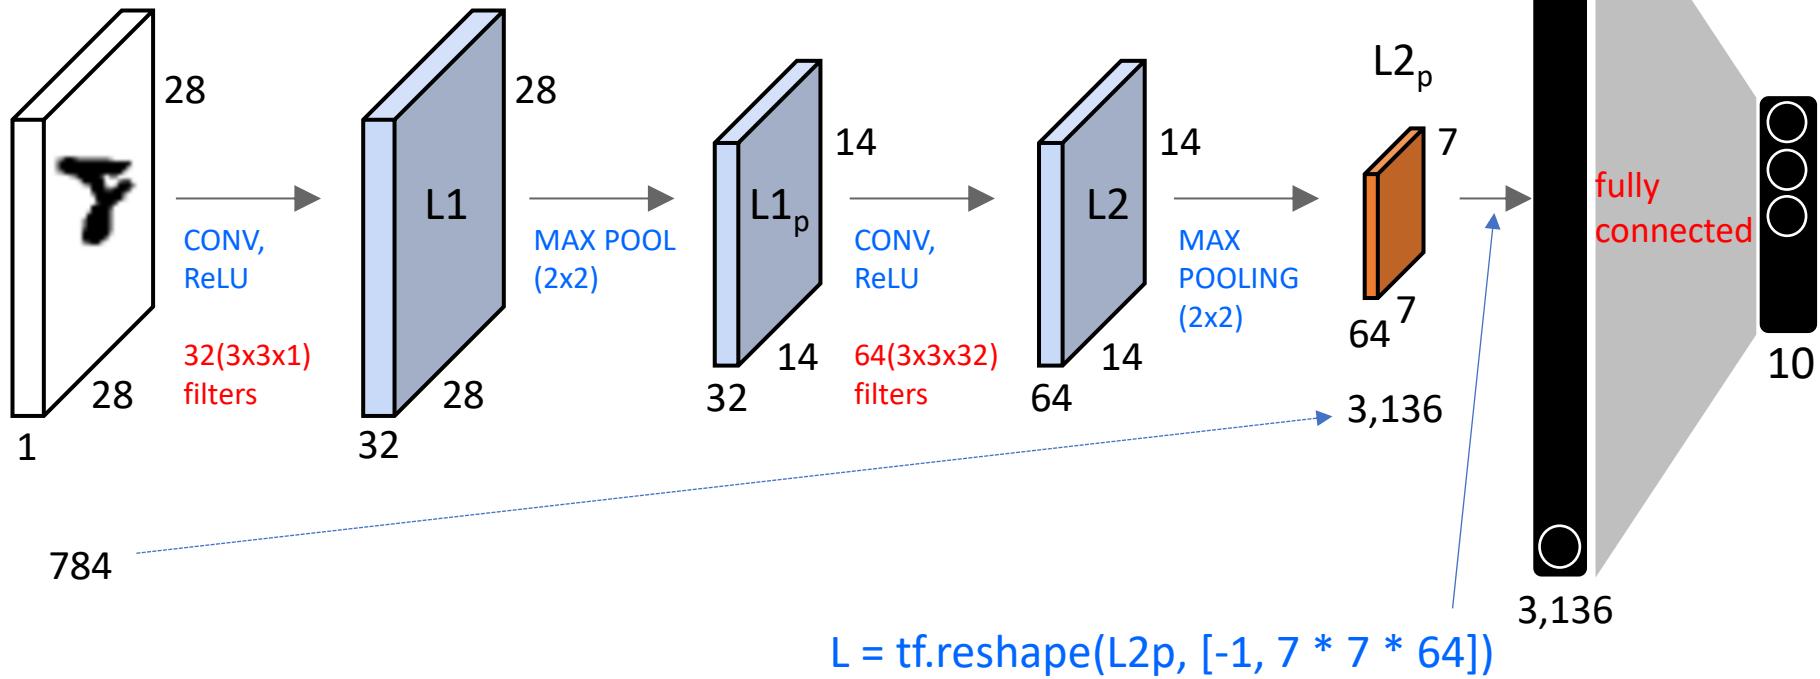


784

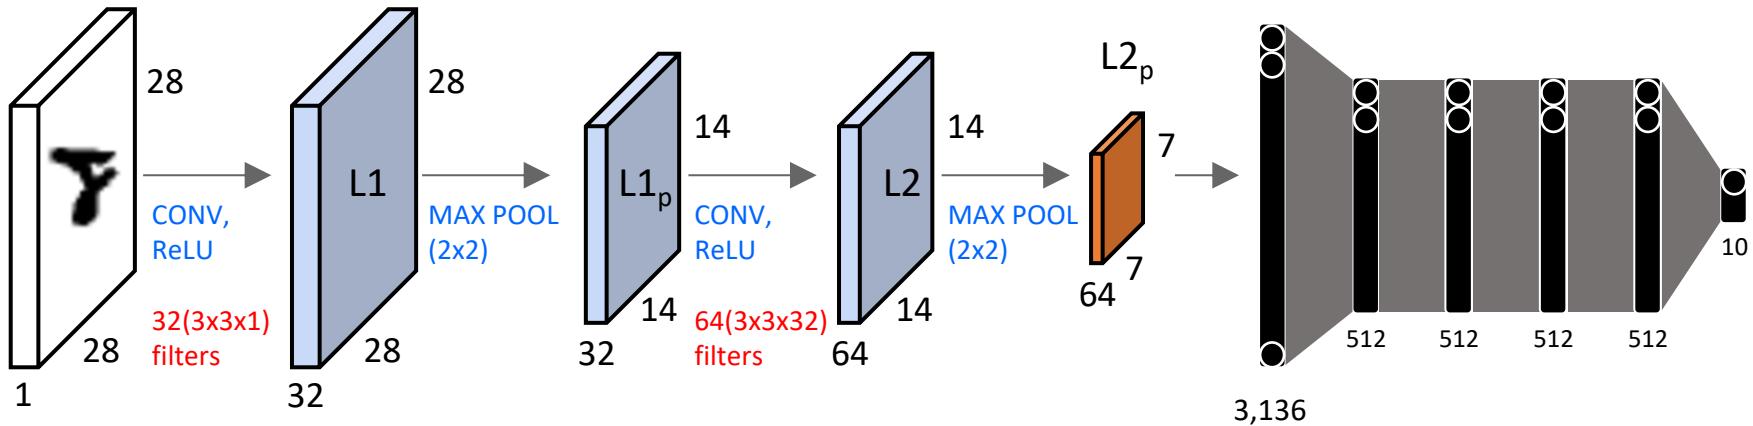
10

26_mnist_softmax.py

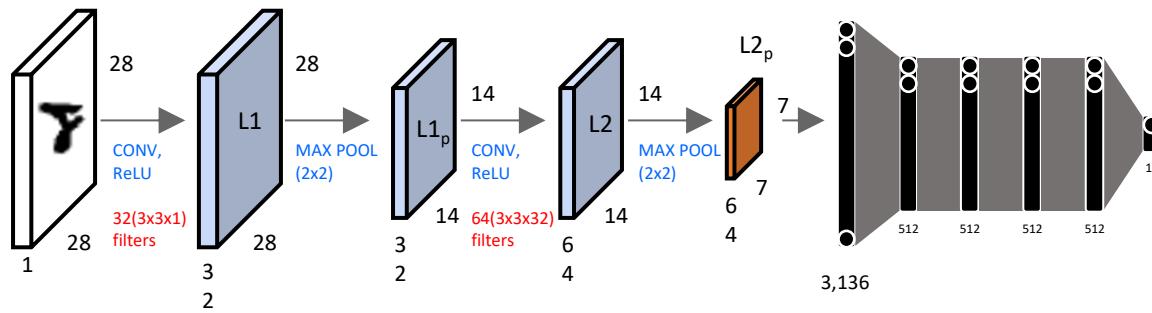
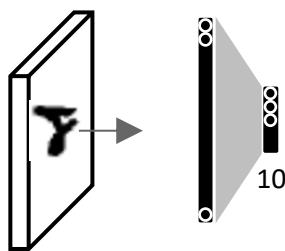
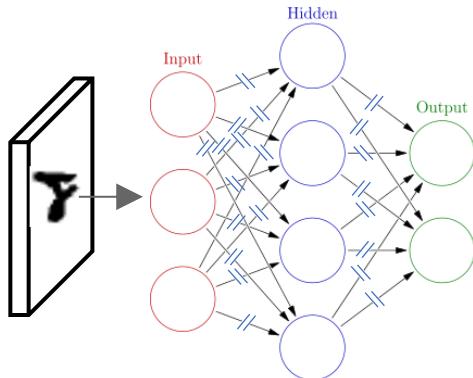
Accuracy: 98.85

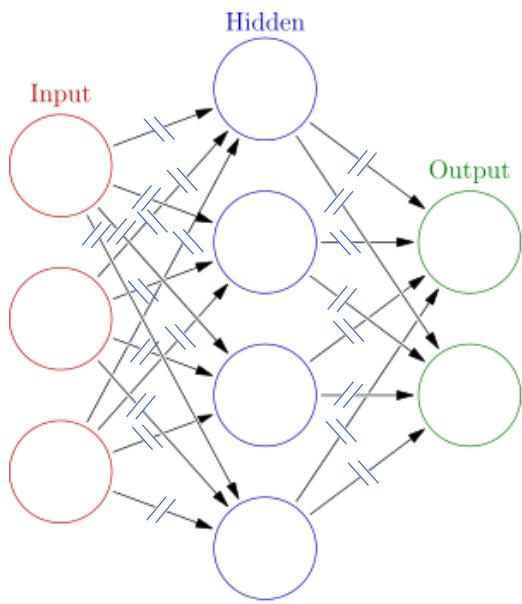


Accuracy: ?????

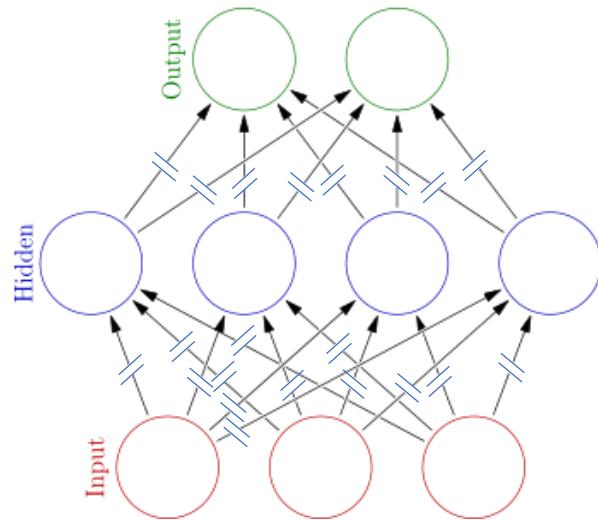


Feedforward NN

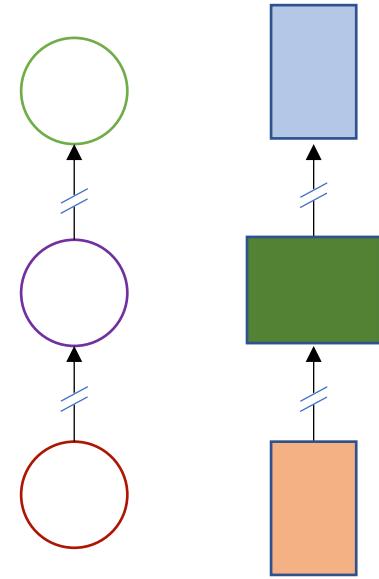




3 layered FCNN

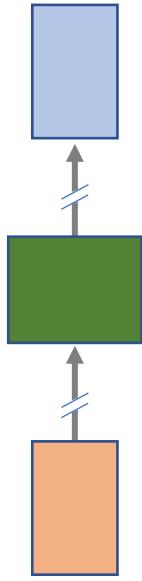


Rotated 90°
clockwise



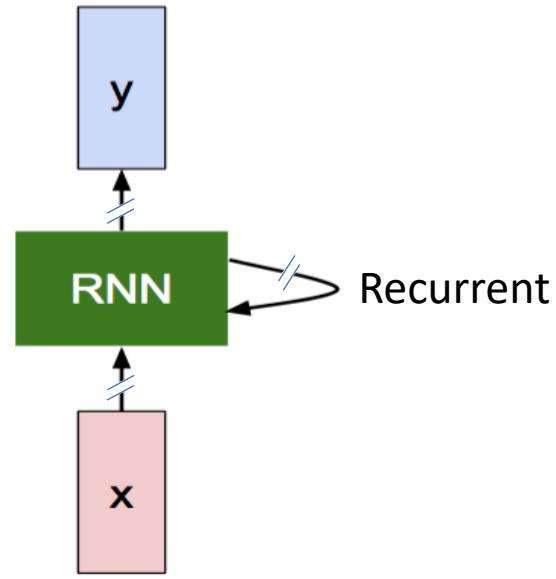
Simplified

Rectangular
representation



Simplified NN

What is the arrow?
Where are synapses?
What is learning?



Recurrent NN

Imagine the inside of RNN

Recurrent Neural Networks
RNN

Seq2Seq

Transformer

LLMs
(ChatGPT, Gemini, Claude, etc)