

---

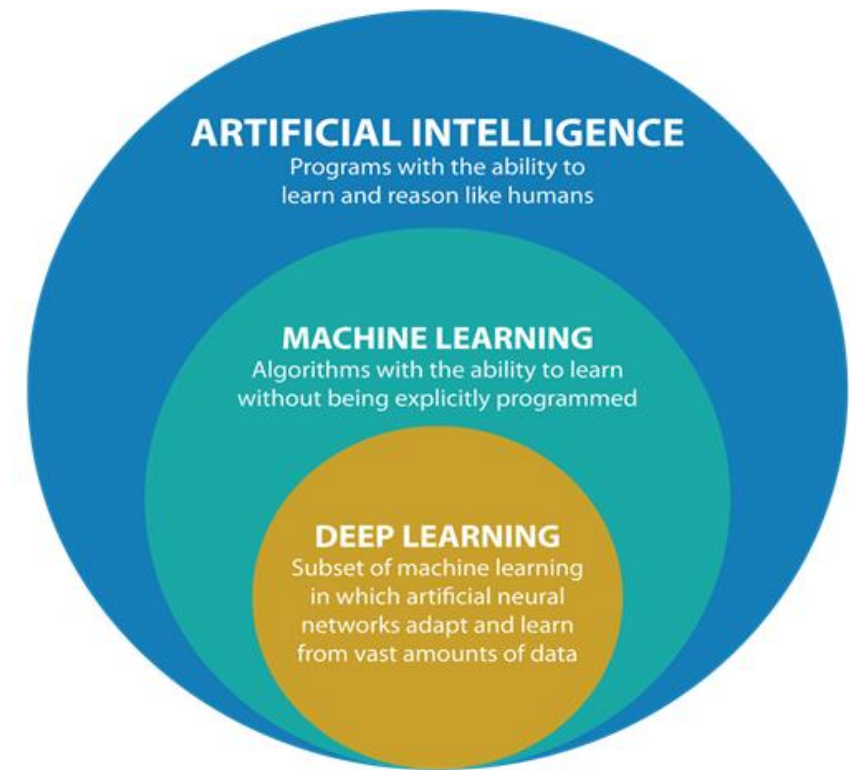
# NEURAL NETWORKS

DEBAPRIYA HAZRA




# HISTORY OF ARTIFICIAL INTELLIGENCE (AI)

- **Dartmouth Conference of 1956** - AI and its mission were defined.
- **1970s** – First AI winter
- The **first robot was constructed in 1972**.
- **1990s** - Second AI winter machine learning
- **1997** - IBM's Deep Blue



# CONVOLUTIONAL NEURAL NETWORKS (CNN)


How a simple ConvNet looks like



## Pooling Layer


- Pooling layer is used to reduce the spatial dimensions of the input data.
- It is commonly used in the first layer of CNNs to reduce the size of the input data.
- Logistic is used for binary classification and softmax for multi-class classification.
- Softmax is used to learn and the amount of computation performed in the network.

**filters** : The number of filters.  
**kernel\_size** : A number specifying both the height and width of the (square) convolution window. There are also some additional optional arguments that you might like to tune.  
**strides** : The stride of the convolution. If you don't specify anything, this is set to one.  
**padding** : This is either **valid** or **same** . If you don't specify anything, the padding is set to **valid** .  
**activation** : This is typically **relu** . If you don't specify anything, no activation is applied. You are strongly encouraged to add a ReLU activation function to every convolutional layer in your networks.




# LENET (2010)

- This network takes a **grayscale 32 x 32 image as input**
- It goes to the **convolution layers (C1)** and then to the **subsampling layer (S2)**
- Nowadays the subsampling layer is replaced by a pooling layer.
- Then, there is another **sequence of convolution layers (C3) followed by a pooling** (that is, subsampling) layer (S4)
- Finally, there are **three fully connected layers**, including the **OUTPUT layer** at the end
- This network was used for **zip code recognition in post offices**



# ALEXNET (2012)

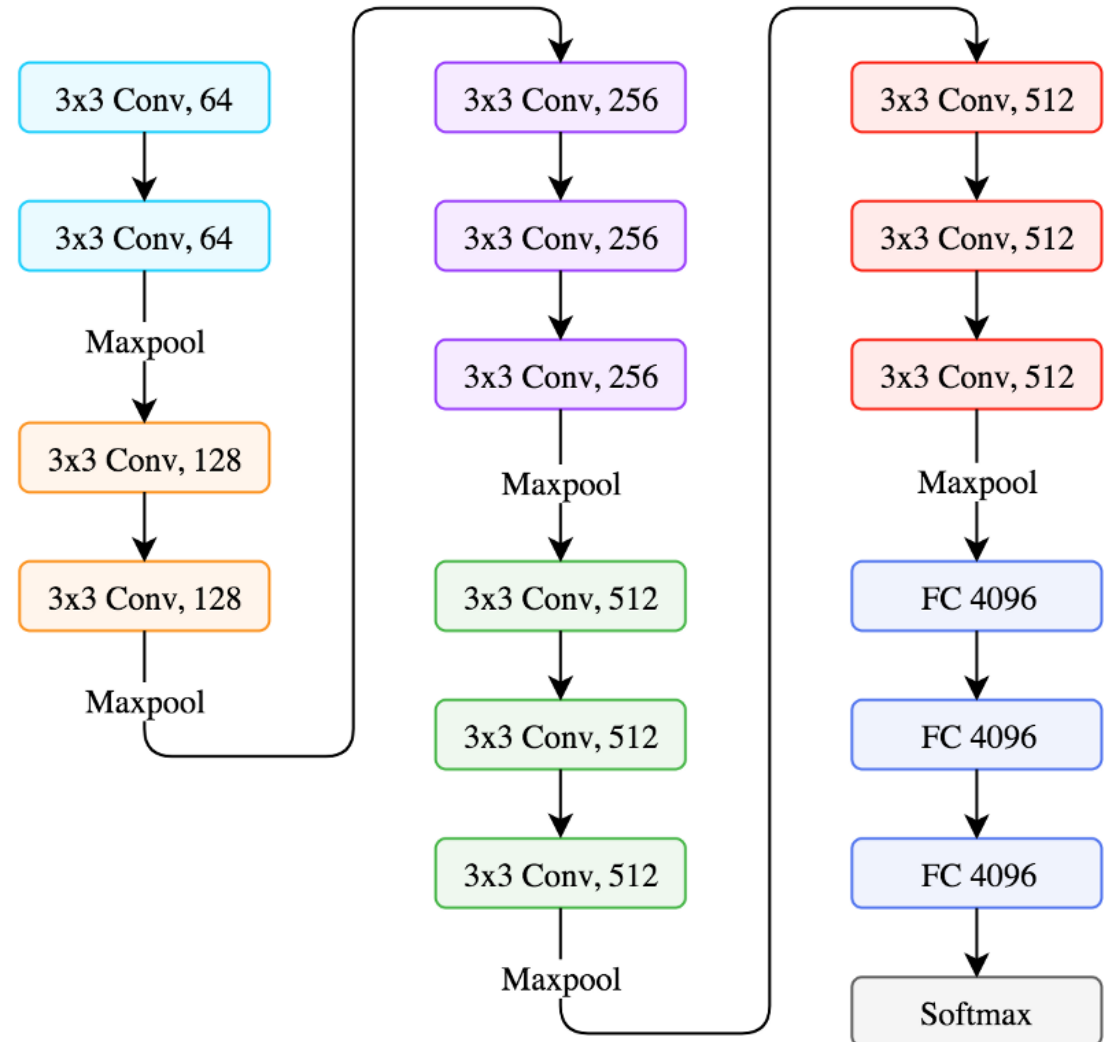
- A **ReLU activation function** and a **dropout of 0.5** were used in this network to **fight overfitting**
- There is a **normalization layer** used in the architecture, but this is **not used in practice anymore** as it used heavy **data augmentation**
- **Simple structure**
- AlexNet is trained on the ImageNet database using **two separate GPUs**, possibly due to processing limitations with inter-GPU connections at the time



When AlexNet is processing an image, this is what is happening at each layer.

# VGGNET (2014)

- The **default input size** of an image for this model is **224 x 224 x 3**.
- The image is passed through a stack of convolution layers with a **stride of 1 pixel** and **padding of 1**.
- It uses **3 x 3 convolution** throughout the network.
- **Max pooling** is done over a **2 x 2 pixel** window with a **stride of 2**
- The first two **fully connected layers have 4,096 neurons each**
- The third **fully connected layers** are responsible for classification with **1,000 neurons**
- All **hidden layers** are built with the **ReLU activation function**



# VGGNET

- It is painfully slow to train.
- The network architecture weights themselves are quite large (in terms of disk/bandwidth).

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input ( $224 \times 224$ RGB image)					
conv3-64	conv3-64 <b>LRN</b>	conv3-64 <b>conv3-64</b>	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 <b>conv3-128</b>	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 <b>conv1-256</b>	conv3-256 conv3-256 <b>conv3-256</b>	conv3-256 conv3-256 conv3-256 <b>conv3-256</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 <b>conv1-512</b>	conv3-512 conv3-512 <b>conv3-512</b>	conv3-512 conv3-512 conv3-512 <b>conv3-512</b>
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

# VGGNET

## Keras

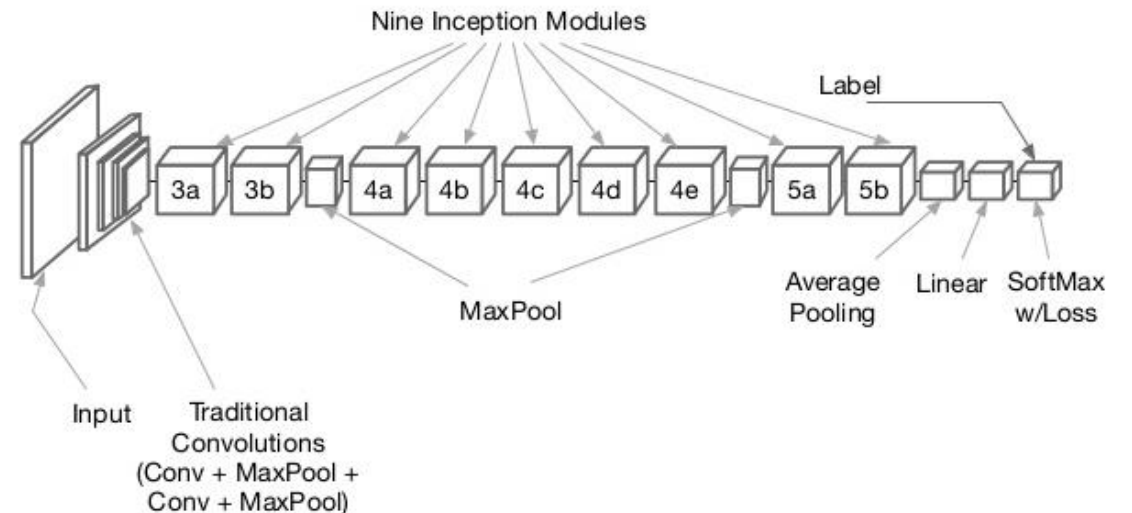
- The Keras Applications module has **pre-trained neural network models**, along with its **pre-trained weights trained on ImageNet**.
- These models can be used directly for **prediction**, **feature extraction**, and **fine-tuning**.
- The first time it executes the preceding script, Keras will automatically download and **cache the architecture weights to disk in the ~/.keras/models directory**
- From next time it will be faster.

```
#import VGG16 network model and other necessary libraries
from keras.applications.vgg16 import VGG16
from keras.preprocessing import image
from keras.applications.vgg16 import preprocess_input
import numpy as np
#Instantiate VGG16 and returns a vgg16 model instance
vgg16_model = VGG16(weights='imagenet', include_top=False)
#include_top: whether to include the 3 fully-connected layers at the top of the network.
#This has to be True for classification and False for feature extraction. Returns a model
#instance weights:'imagenet' means model is pre-training on ImageNet data.
model = VGG16(weights='imagenet', include_top=True)
model.summary()
#image file name to classify
image_path = 'jumping_dolphin.jpg'
#load the input image with keras helper utilities and resize the image.
#Default input size for this model is 224x224 pixels.
img = image.load_img(image_path, target_size=(224, 224))
#convert PIL (Python Image Library??) image to numpy array
x = image.img_to_array(img)
print (x.shape)
#image is now represented by a NumPy array of shape (224, 224, 3),
# but we need to expand the dimensions to be (1, 224, 224, 3) so we can
# pass it through the network -- we'll also preprocess the image by
# subtracting the mean RGB pixel intensity from the ImageNet dataset
#Finally, we can load our Keras network and classify the image:
x = np.expand_dims(x, axis=0)
print (x.shape)
preprocessed_image = preprocess_input(x)
preds = model.predict(preprocessed_image)
print('Prediction:', decode_predictions(preds, top=2)[0])
```



# GOOGLENET (2014)

- The main attractive feature of GoogLeNet is that it runs very fast due to the introduction of a new concept called **inception module**.
- GoogLeNet has **22 layers**
- The main advantage of GoogLeNet is that it is **more accurate than VGG**, while using **much fewer parameters** and **less compute power**.



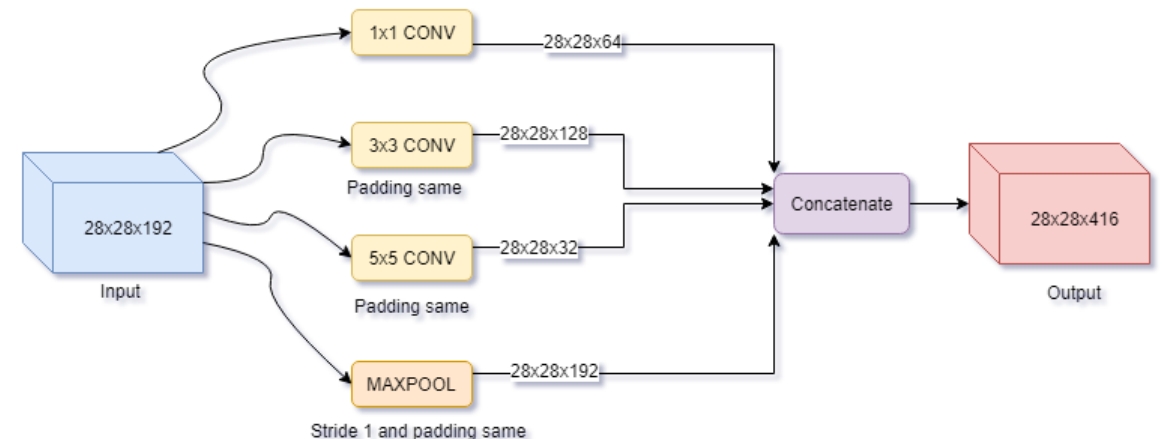
# GOOGLNET

## GoogleNet Architecture

- **Vanishing gradient** problem is avoided by adding **auxiliary classifiers to intermediate layers**
- Since **fully connected layers** are prone to overfitting, it is replaced with a global average pooling (**GAP layer**)
- GoogLeNet added a **linear layer** for the ease of transfer learning

## Inception Module

- In addition to creating deeper networks, the inception block introduces the idea of **parallel convolutions**.
- In-parallel convolutions of different sizes are performed on the output of the previous layer.



# GOOGLNET

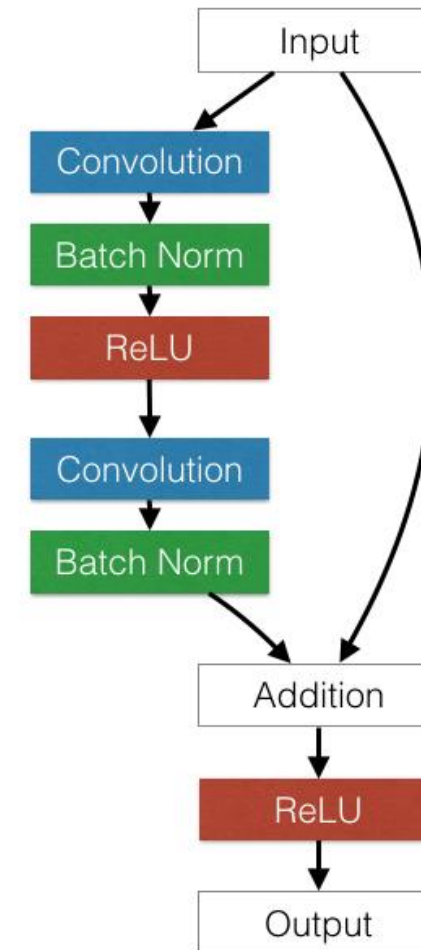
- Googlenet will be just **many inception blocks in cascade**
- The main disadvantage is still the **gradient vanishing** that would occur if we begin stacking lots and lots of inception layers
- **Multiple branches and losses**

## Inception Block

```
import tensorflow as tf
def inception_block_a(x, name='inception_a' ):
    # num of channels: 384 = 96*4
    with tf.variable_scope(name):
        # Pooling part
        b1 = tf.layers.average_pooling2d(x, [3,3], 1, padding='SAME' )
        b1 = tf.layers.conv2d(inputs=b1, filters=96, kernel_size=[1, 1], padding="same", activation=tf.nn.relu)
        # 1x1 part
        b2 = tf.layers.conv2d(inputs=x, filters=96, kernel_size=[1, 1], padding="same", activation=tf.nn.relu)
        # 3x3 part
        b3 = tf.layers.conv2d(inputs=x, filters=64, kernel_size=[1, 1], padding="same", activation=tf.nn.relu)
        b3 = tf.layers.conv2d(inputs=b3, filters=96, kernel_size=[3, 3], padding="same", activation=tf.nn.relu)
        # 5x5 part
        b4 = tf.layers.conv2d(inputs=x, filters=64, kernel_size=[1, 1], padding="same", activation=tf.nn.relu)
        # 2 3x3 in cascade with same depth is the same as 5x5 but with less parameters
        # b4 = tf.layers.conv2d(inputs=b4, filters=96, kernel_size=[5, 5], padding="same", activation=tf.nn.relu)
        b4 = tf.layers.conv2d(inputs=b4, filters=96, kernel_size=[3, 3], padding="same", activation=tf.nn.relu)
        b4 = tf.layers.conv2d(inputs=b4, filters=96, kernel_size=[3, 3], padding="same", activation=tf.nn.relu)
        concat = tf.concat([b1, b2, b3, b4], axis=-1)
        return concat
```

# RESIDUAL NETWORKS (2015)

- After a certain depth, adding additional layers to feed-forward convNets results in a **higher training error** and **higher validation error**.
- When adding layers, performance increases only up to a certain depth, and then it rapidly decreases.
- The **ResNet** team added connections that **can skip layers**
- The **residual block**, which adds the **output of the previous layer to the output of the next layer**




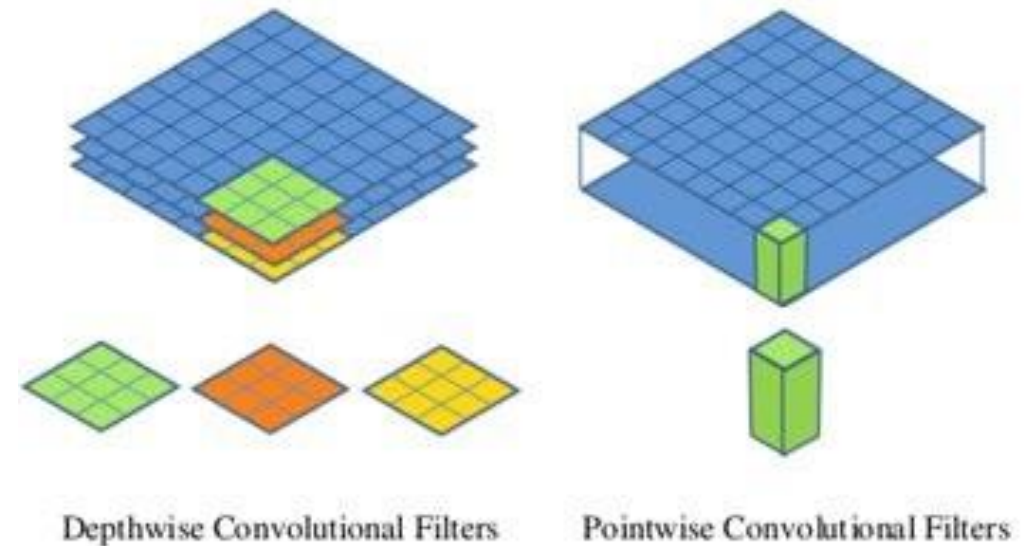
# RESIDUAL NETWORKS

- It has different depth variations, such as **34, 50, 101,** and **152 layers**
- The ResNet architecture is a stack of residual blocks.
- Every residual block has **3 x 3 convolution layers**
- After the last conv layer, a **GAP layer** is added.
- There is only **one fully connected** layer to classify 1,000 classes
- For a deeper network, say more than 50 layers, it uses the **bottleneck features** concept to improve efficiency.
- **No dropout** is used in this network.

```
import tensorflow as tf
from collections import namedtuple
# Configurations for each bottleneck group.
BottleneckGroup = namedtuple('BottleneckGroup',
['num_blocks', 'num_filters', 'bottleneck_size' ])
groups = [
    BottleneckGroup(3, 128, 32), BottleneckGroup(3, 256, 64), BottleneckGroup(3, 512, 128),
    BottleneckGroup(3, 1024, 256)
]
# Create the bottleneck groups, each of which contains `num_blocks`
# bottleneck groups.
for group_i, group in enumerate(groups):
    for block_i in range(group.num_blocks):
        name = 'group_%d/block_%d' % (group_i, block_i)
        # 1x1 convolution responsible for reducing dimension
        with tf.variable_scope(name + '/conv_in' ):
            conv = tf.layers.conv2d(
                net,
                filters=group.num_filters,
                kernel_size=1,
                padding='valid' ,
                activation=tf.nn.relu)
            conv = tf.layers.batch_normalization(conv, training=training) with
            tf.variable_scope(name + '/conv_bottleneck' ):
                conv = tf.layers.conv2d(
                    conv,
                    filters=group.bottleneck_size,
                    kernel_size=3,
                    padding='same' ,
                    activation=tf.nn.relu)
                conv = tf.layers.batch_normalization(conv, training=training)
        # 1x1 convolution responsible for restoring dimension
        with tf.variable_scope(name + '/conv_out' ):
            input_dim = net.get_shape()[-1].value
            conv = tf.layers.conv2d(
                conv,
                filters=input_dim,
                kernel_size=1,
                padding='valid' ,
                activation=tf.nn.relu)
            conv = tf.layers.batch_normalization(conv, training=training)
        # shortcut connections that turn the network into its counterpart
        # residual function (identity shortcut)
        net = conv + net
```

# MOBILENETS (2017)

- Works **faster on mobile devices**.
- Created by Google, MobileNet's key feature is that it uses a **different "sandwich" form of convolution block**.
- Instead of the usual (CONV, BATCH\_NORM, RELU), it **splits 3x3 convolutions up into a 3x3 depthwise convolution, followed by a 1x1 Pointwise CONV**.
- They call this block a **depthwise separable convolution**.



Depthwise Separable Convolution

# MOBILENETS

## Depthwise Separable Convolution

- The new convolution block (`tf.layers.separable_conv2d`) consists of two main parts:
  - a depthwise convolution layer
  - followed by a `1x1` pointwise convolution layer.
- However, it only filters input channels, and it does not combine them to create new features.


## Control Parameters

- MobileNets uses **two hyperparameters** to help control the trade-off between **accuracy** and **speed**
- **Width Multiplier**: Controls the Depthwise CONVs accuracy by uniformly reducing the number filters used throughout the network
- **Resolution Multiplier**: Simply scales down the input image to different sizes

# WHY GENERATIVE MODELS ?

- With discriminative models, we generally try to find ways of separating or "discriminating" between different classes in our data
- However, with generative models, we try to find out the probability distribution of our data


- Pretrain a model with unlabeled data
- Augment your dataset (in theory, if you capture the probability distribution of your data, you can generate more data)
- Compress your data (lossy)
- Create some sort of simulator (for example, a quadcopter can be controlled with four inputs; if you capture this data and train a generative model on it, you can learn the dynamics of the drone)





# AUTOENCODERS

- An autoencoder is a **regular neural network**, an **unsupervised learning model** that takes an input and produces the same input in the output layer.
- The idea is that the **encoder part will compress** your input into a smaller dimension.
- From this smaller dimension, it then tries to **reconstruct the input using the decoder** part of the model.
- An encoder can be a **fully connected neural network** or a **convolutional neural network (CNN)**.
- A decoder also uses the **same kind of network** as an encoder.



# AUTOENCODERS

## Uses


- The use of autoencoders can be **good for pretraining**
- Once trained, you can use the weights of the encoder and then fine-tune them to your intended task.
- Another use is as a form of **compression** for your data if it isn't too complicated.
- You can use the autoencoder to **reduce the dimensionality down to two or three dimensions** and then try to visualize your inputs in the latent space to see whether it shows you anything useful.

## Limitations

- They cannot be used to generate more data for us.
- This is because we don't know how to create new latent vectors to feed to the decoder

# VARIATIONAL AUTOENCODERS (VAE)

- The first generative model that could **create more data that resembles the training data**
- VAE has a new constraint
- When we want to use our VAE to generate new data, we can just **create sample vectors that come from a unit Gaussian distribution and give them to the trained decoder.**



# VARIATIONAL AUTOENCODERS (VAE)

- We need **two parameters** to keep track and to enforce our VAE model to produce a **normal distribution in the latent space**:
  - **Mean** (should be zero)
  - **Standard deviation** (should be one)

## VAE loss function

- **Generative loss**: This loss compares the model output with the model input.
- **Latent loss**: The loss we use here will be the **KL divergence loss**. This loss term penalizes the VAE if it starts to produce latent vectors that are not from the desired distribution.

```
generation_loss = mean(square(generated_image - real_image))  
latent_loss = KL-Divergence(latent_variable, unit_gaussian)  
loss = generation_loss + latent_loss
```

# VARIATIONAL AUTOENCODERS (VAE)

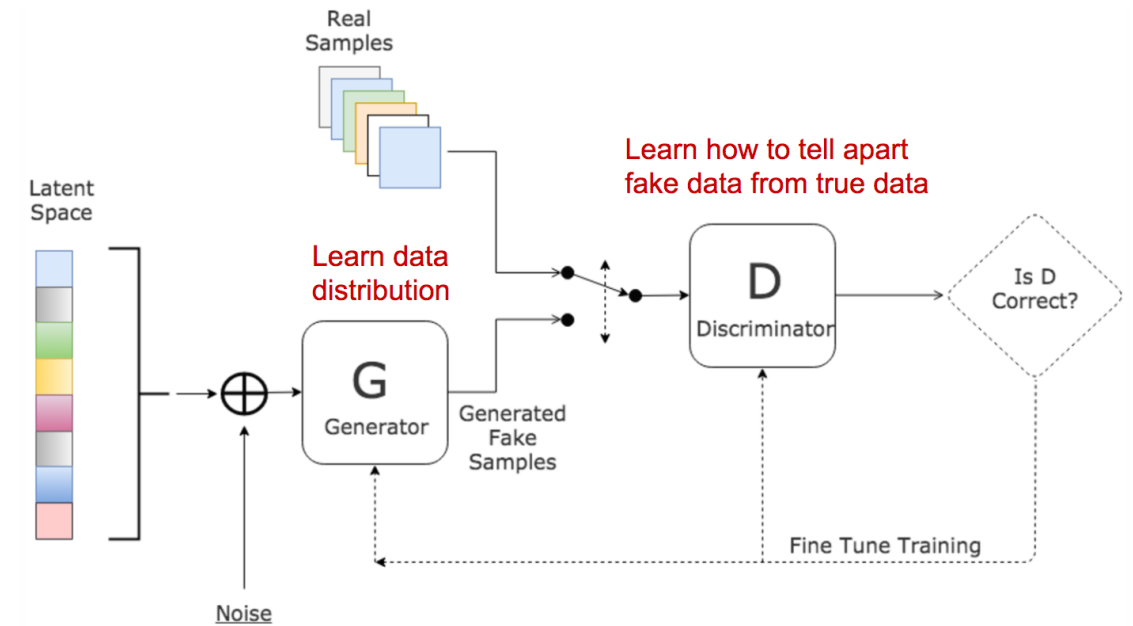


## Kullback-Leibler divergence

- The blue distribution is trying to model the green distribution
- As the blue distribution comes closer and closer to the green one, the KL divergence loss will get closer to zero.

# GENERATIVE ADVERSARIAL NETWORKS (GAN) (2014)

- **Generator:** Create images similar to the **real images** dataset using a size N, 1-D vector as input (Choice of N is up to us)
- **Discriminator:** Verify that the image given to it is real or a fake generated one



# GENERATIVE ADVERSARIAL NETWORKS (GAN) (2014)

## Practical usage of GANs

- Use the discriminator network weights as the initialization for a different task **similar to what we can do with autoencoders**
- Use the generator network to **create new images**, possibly to augment your dataset, like we can do with the trained decoder of the VAE
- Use the discriminator as a **loss function** (potentially, better than L1/L2 for images) and can also be used back in the **VAE**
- **Semi-supervised learning** by mixing generated data with labeled data

# GENERATIVE ADVERSARIAL NETWORKS (GAN) (2014)

## The Discriminator

- The discriminator takes as input a batch of **784 length vectors**, which is our **28x28 MNIST images** flattened.
- The output will be just a single number for each image
- We use Leaky ReLu as the activation function to prevent ReLu units from dying out.

```
def discriminator(x):  
    with tf.variable_scope("discriminator"):  
        fc1 = tf.layers.dense(inputs=x, units=256, activation=tf.nn.leaky_relu)  
        fc2 = tf.layers.dense(inputs=fc1, units=256, activation=tf.nn.leaky_relu)  
        logits = tf.layers.dense(inputs=fc2, units=1)  
        return logits
```

## Discriminator Loss

$$\frac{1}{m} \sum_{i=1}^m [\log D(x^i) + \log(1 - D(G(z^i)))]$$

It wants to output 1 for real image and 0 for generated images.



# GENERATIVE ADVERSARIAL NETWORKS (GAN) (2014)

## The Generator

- The job of the generator is to **take a vector of random noise as input** and from this, **produce an output image**
- We use the **tanh activation** on the output to restrict generated images to be in **the range -1 to 1**.

```
def generator(z):  
    with tf.variable_scope("generator"):  
        fc1 = tf.layers.dense(inputs=z, units=1024, activation=tf.nn.relu)  
        fc2 = tf.layers.dense(inputs=fc1, units=1024, activation=tf.nn.relu)  
        img = tf.layers.dense(inputs=fc2, units=784, activation=tf.nn.tanh)  
        return img
```

## Generator Loss

$$\frac{1}{m} \sum_{i=1}^m \log(D(G(z^i)))$$

- m: Batch size
- D: Discriminator
- G: Generator
- z: Random noise vector

We want to **maximize this loss function** when training our GAN. When the loss is maximized, it means the generator is capable of generating images that can fool the discriminator, and the discriminator is outputting 1 for generated images.

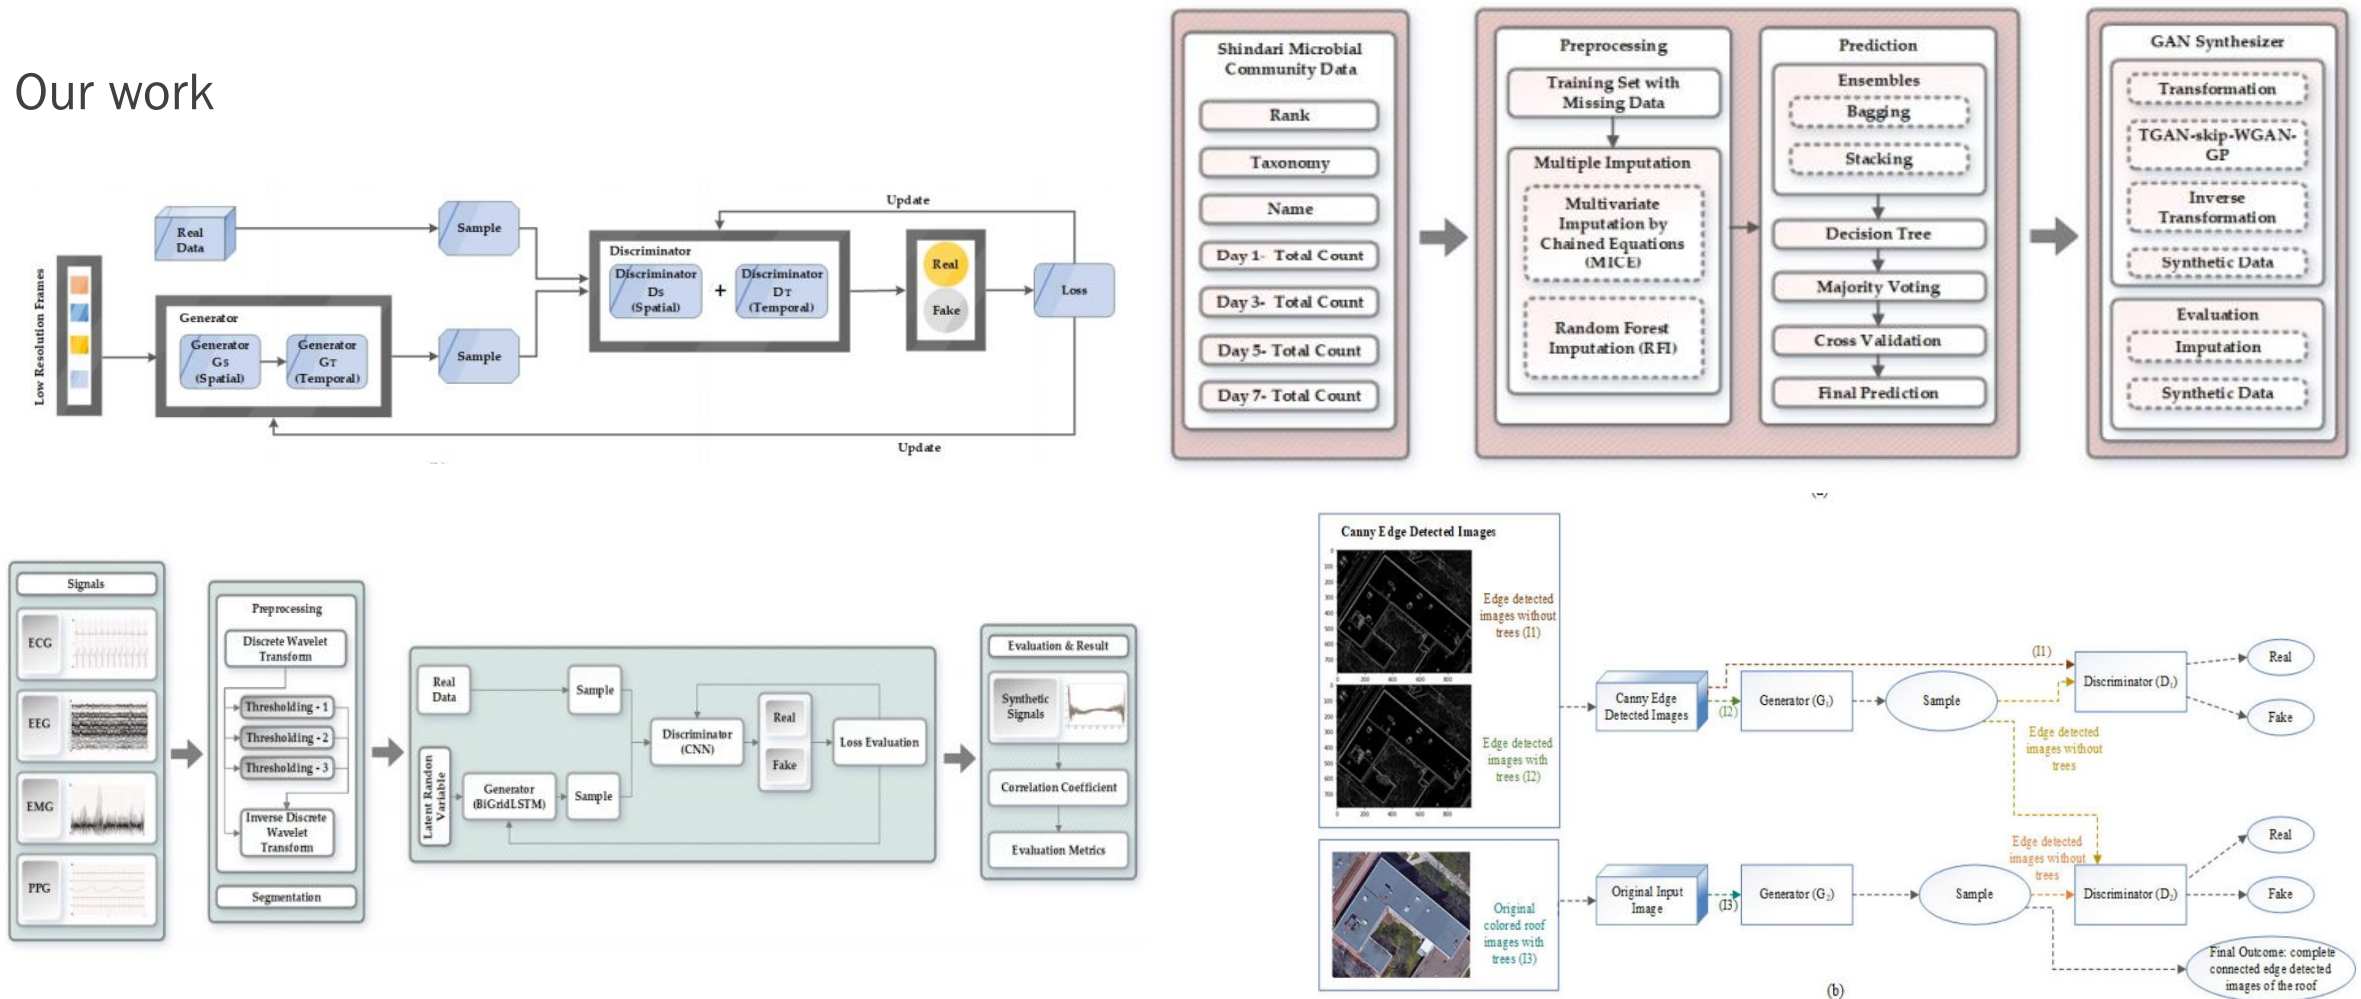


# PROBLEMS WITH GAN

- **Non-convergence:** the model parameters oscillate, destabilize and never converge,
- **Mode collapse:** the generator collapses which produces limited varieties of samples,
- **Diminished gradient:** the discriminator gets too successful that the generator gradient vanishes and learns nothing,
- Unbalance between the generator and discriminator causing overfitting, &
- Highly sensitive to the hyperparameter selections.

# GENERATIVE ADVERSARIAL NETWORKS (GAN) (2014)

Our work



# TRANSFER LEARNING

- The idea is to **transfer something learned from one task and apply it to another.**

When should we use transfer learning? Transfer learning can be applied in the following situations, depending on:

- The size of the new (target) dataset
- Similarity between the original and target datasets

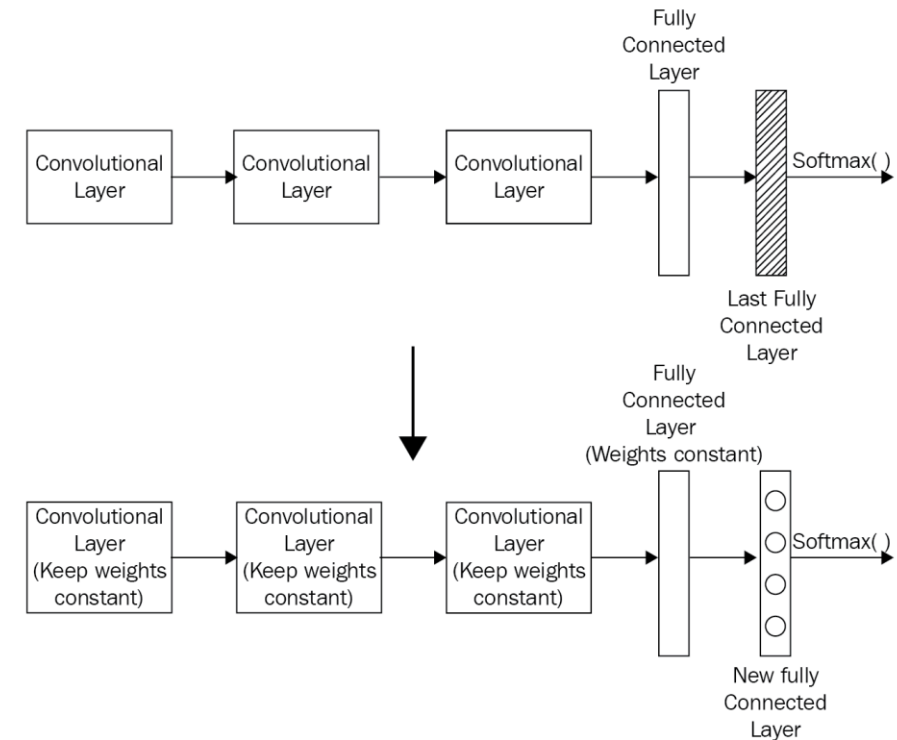
There are four main use cases:

- **Case 1:** New (target) dataset is small and is similar to the original training dataset
- **Case 2:** New (target) dataset is small but is different from the original training dataset
- **Case 3:** New (target) dataset is large and is similar to the original training dataset
- **Case 4:** New (target) dataset is large and is different from the original training dataset

# TRANSFER LEARNING

Target dataset is small and is similar to the original training dataset

- In this case, replace the last fully connected layer with a **new fully connected layer that matches with the number of classes of the target dataset**
- Initialize **old weights** with **randomized weights**
- Train the network to update the weights of the new, fully connected layer:

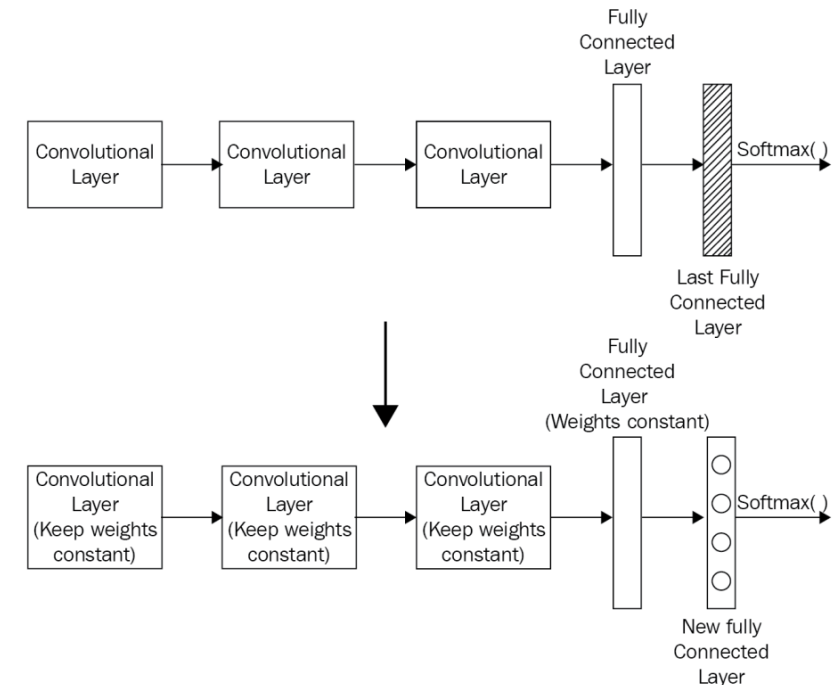


(Transfer Learning : Small data set and Similar data)

# TRANSFER LEARNING

Target dataset is small but different from the original training dataset

- Slice most of the initial layers of the network
- Add to the remaining pre-trained layers a **new fully connected layer** that matches the number of classes of the target dataset
- **Randomize the weights** of the new fully connected layer and freeze all the weights from the pre-trained network
- Train the network to **update the weights** of the new fully connected layer

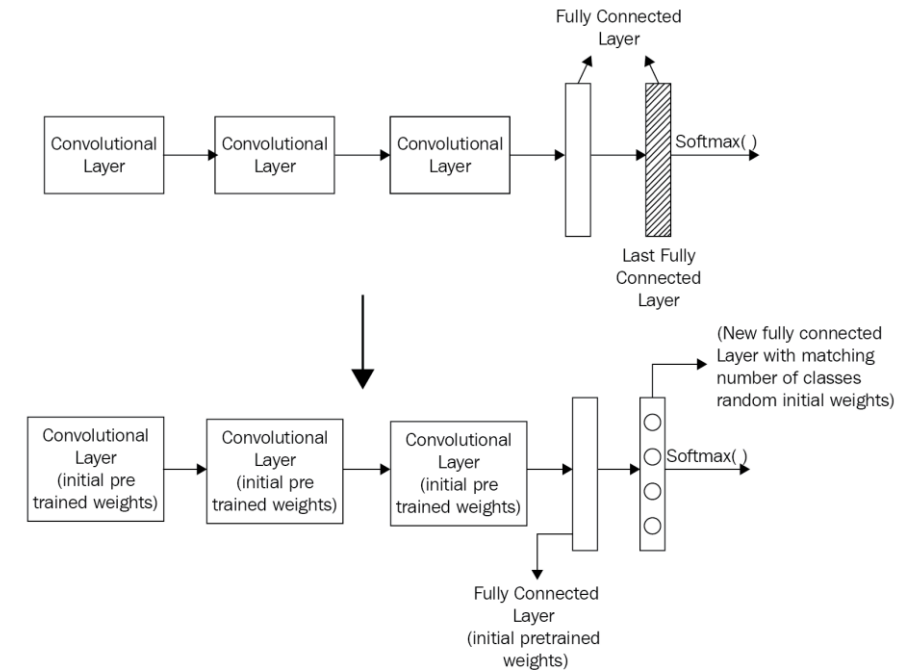


(Transfer Learning : Small data set and Similar data)

# TRANSFER LEARNING

Target dataset is large and similar to the original training dataset

- Remove the last fully connected layer and replace it with a fully connected layer that matches the number of classes in the target dataset
- Randomly initialize the weights of this newly added, fully connected layer
- Initialize the rest of the weights with pre-trained weights
- Train the entire network:

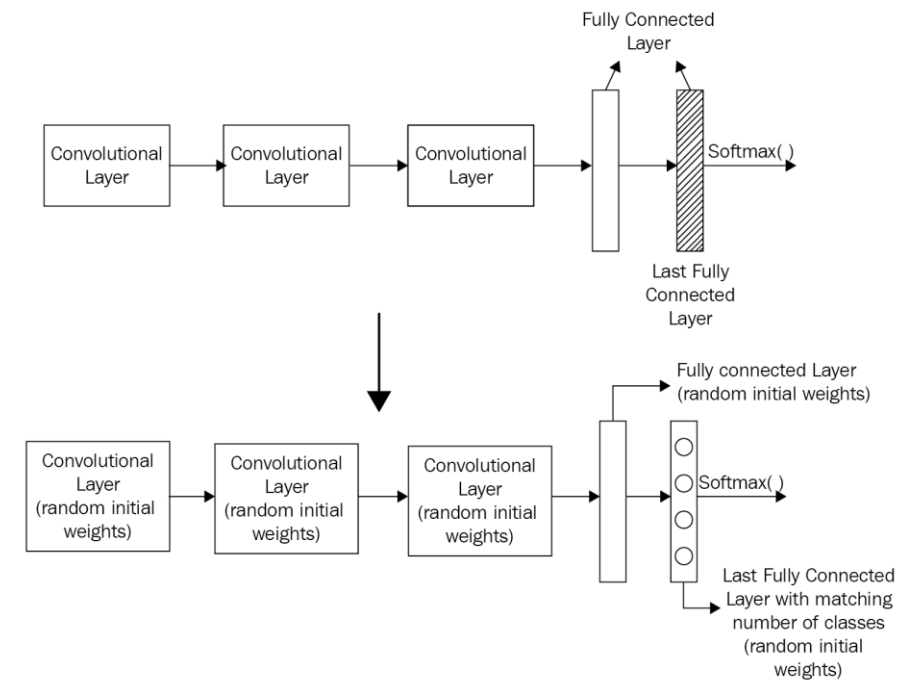


(Transfer Learning : Large data set and Similar data)

# TRANSFER LEARNING

Target dataset is large and different from the original training dataset

- Remove the last fully connected layer and replace it with a **fully connected layer** that matches the number of classes in the target dataset
- Train the entire network from scratch with **randomly initialized weights**:



(Transfer Learning : Large data set and Different data)





# Thank you !

**ANY QUESTIONS ?**